

18-447 Lecture 7: Pipelined Implementation

James C. Hoe


Department of ECE

Carnegie Mellon University

Housekeeping



- Your goal today
 - getting started on pipelined implementations
- Notices
 - Lab 1, Part B, **due Friday midnight**
 - HW1, **past due**
 - Handout #5: HW 2
 - Handout #6: HW 1 solutions
- Readings
 - P&H Ch 4

Doing laundry more quickly: in theory



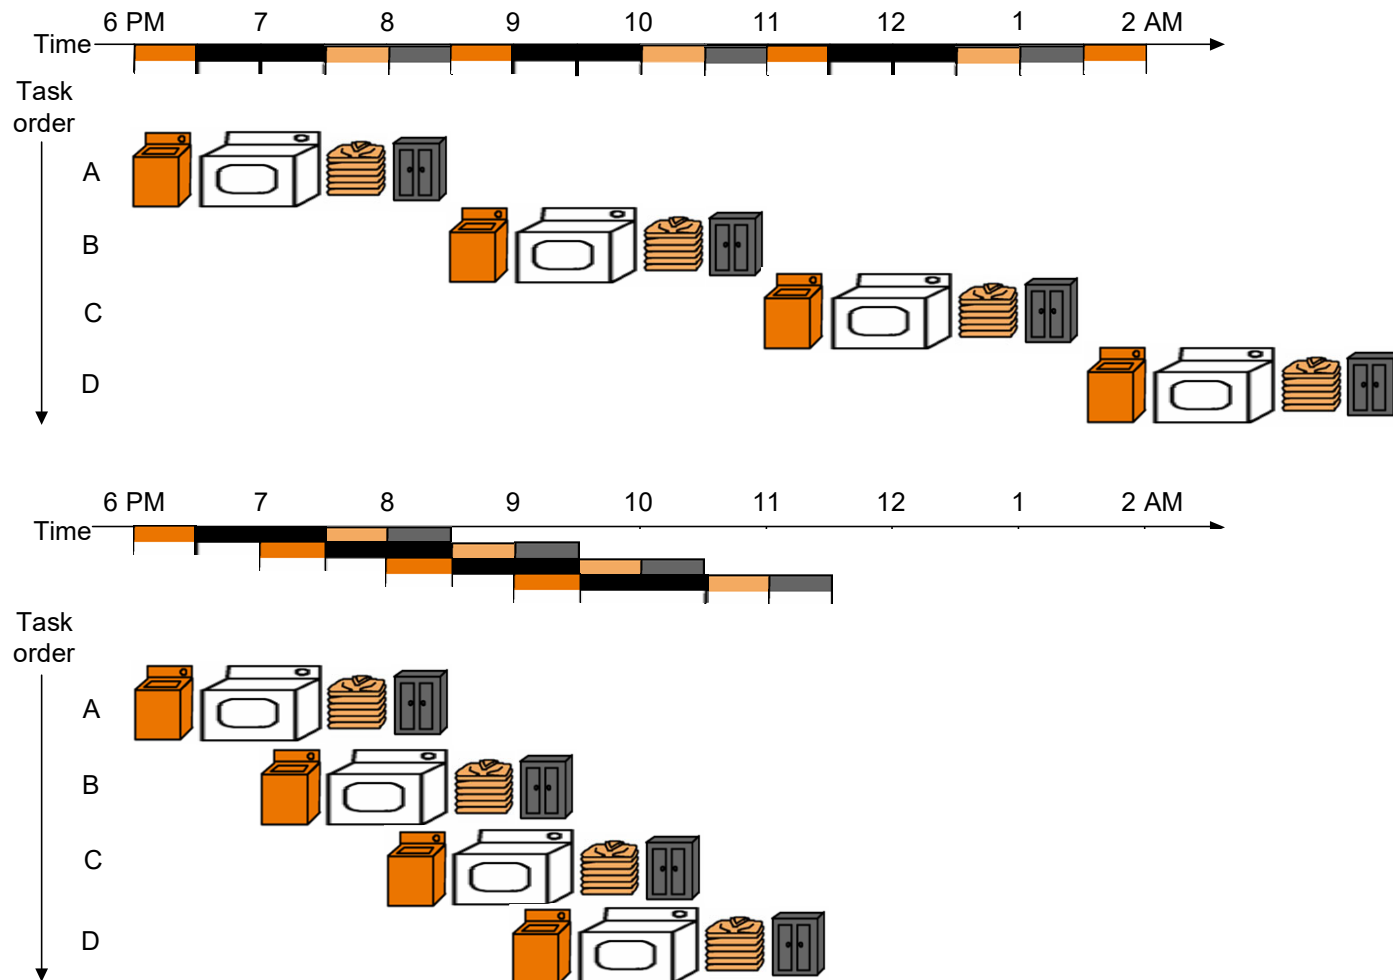
1. “place one dirty load of clothes in **washer**”
2. “when washer is finished, place wet clothes in **dryer**”
3. “when dryer is finished, **you** fold dried clothes”
4. “when folding is finished, ask **friend** to put clothes away”
 - steps to do a load are sequentially dependent
 - no dependence between different loads
 - different steps do not share **resources**

Doing laundry more quickly: in theory



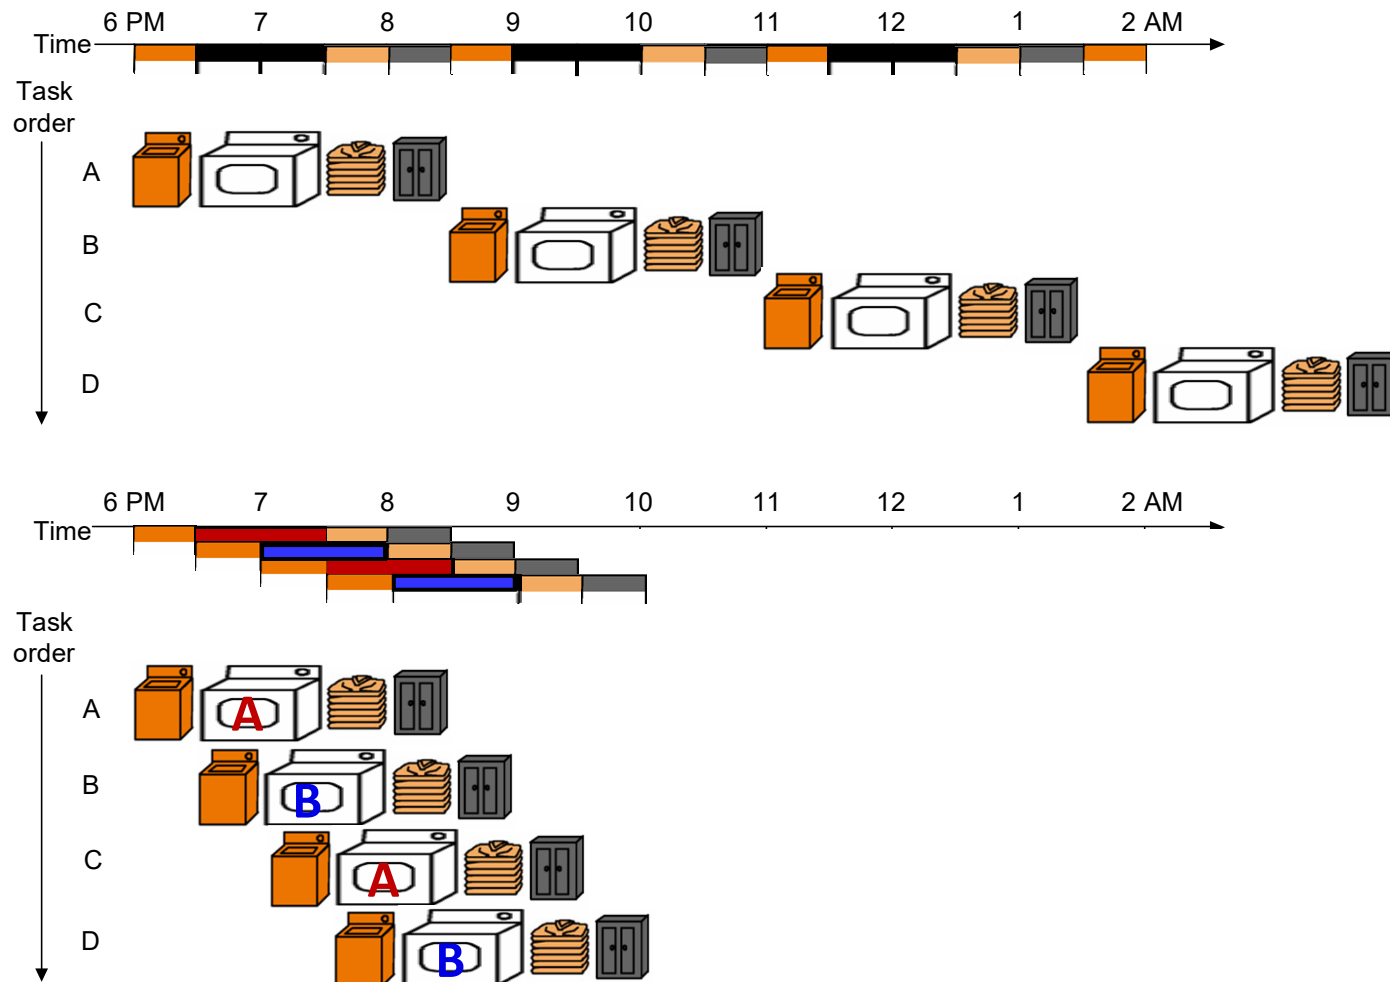
- 4-loads of laundry in parallel
- no additional resources
(all resources always busy!)
- **throughput** increased by 4
- **latency** for a load is the same

Doing laundry more quickly: in practice






the slowest step decides **throughput**

Doing laundry more quickly: in practice



Throughput restored (2 loads per hour) using 2 dryers

(Ideal) HW Pipelining




Notice: evenly divisible; no feedback wires

Performance Model

- Nonpipelined version with delay T


$\text{throughput} = 1/(T+S)$ where S = latch delay



- k -stage pipelined version

$\text{throughput}_{k\text{-stage}} = 1 / (T/k + S)$

$\text{throughput}_{\text{max}} = 1 / (1 \text{ gate delay} + S)$



per-task latency became longer: $T+kS$

Cost Model


- Nonpipelined version with combinational cost G

$\text{Cost} = G + L$ where L = latch cost



- k -stage pipelined version

$\text{Cost}_{k\text{-stage}} = G + Lk$



Pipeline Idealism

Motivation: Increase throughput without adding hardware cost

- Repetition of identical tasks

same task repeated for many different inputs

- Repetition of independent tasks

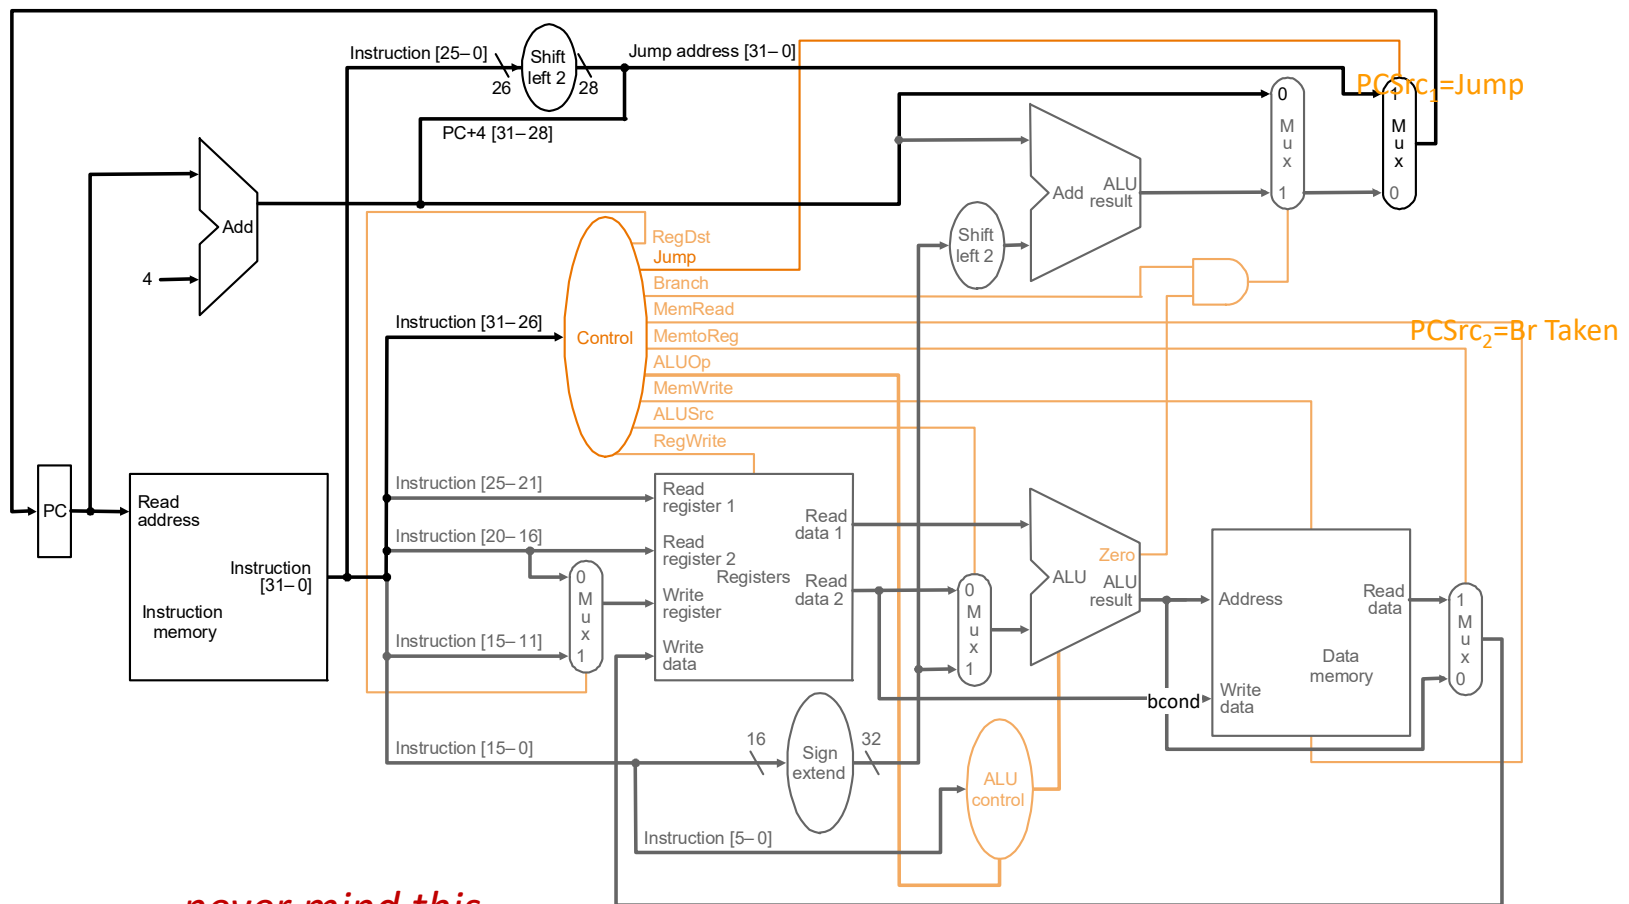
no ordering dependencies between repeated tasks

- Uniformly partitionable suboperations

arbitrary number and placement of boundaries

Good examples: automobile assembly line, doing laundry, but instruction execution???

Reality of Instruction Pipelining

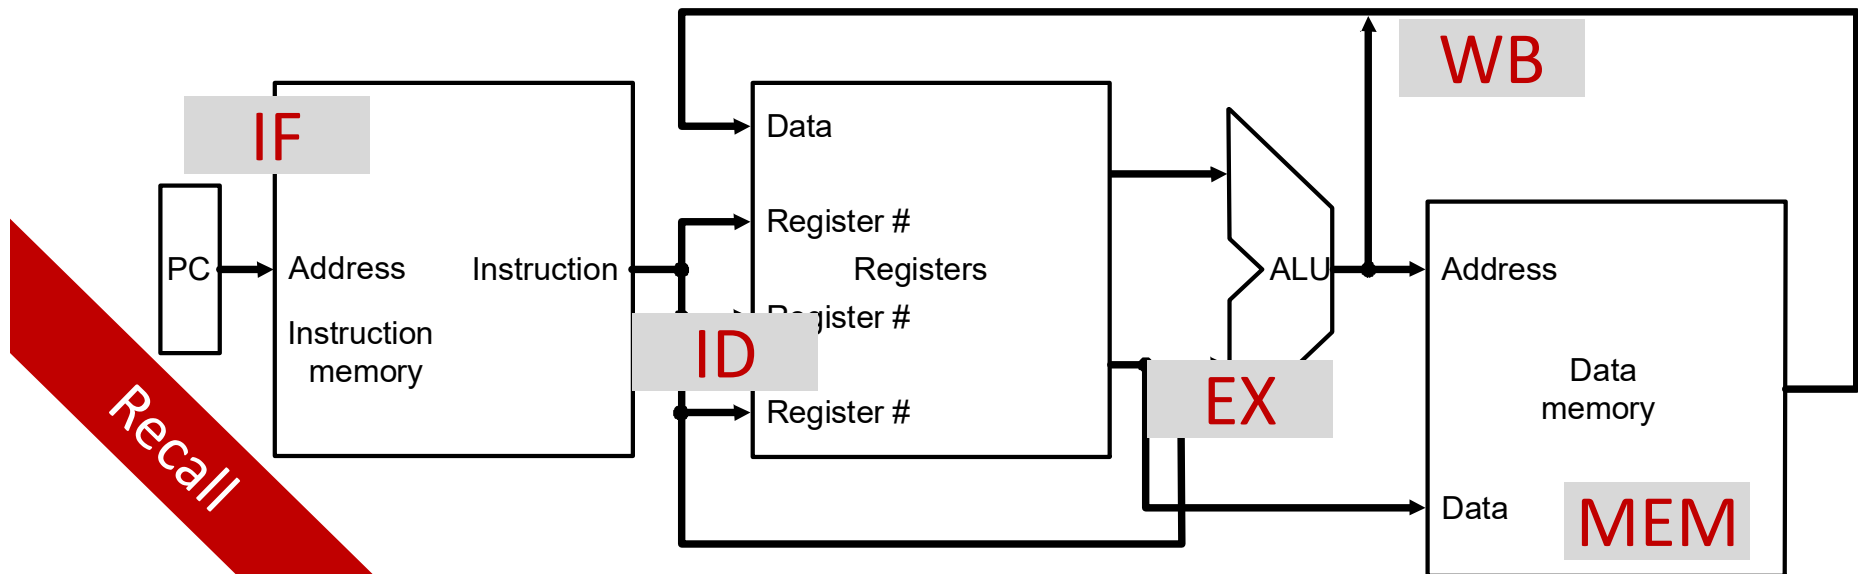


*never mind this
complication today*



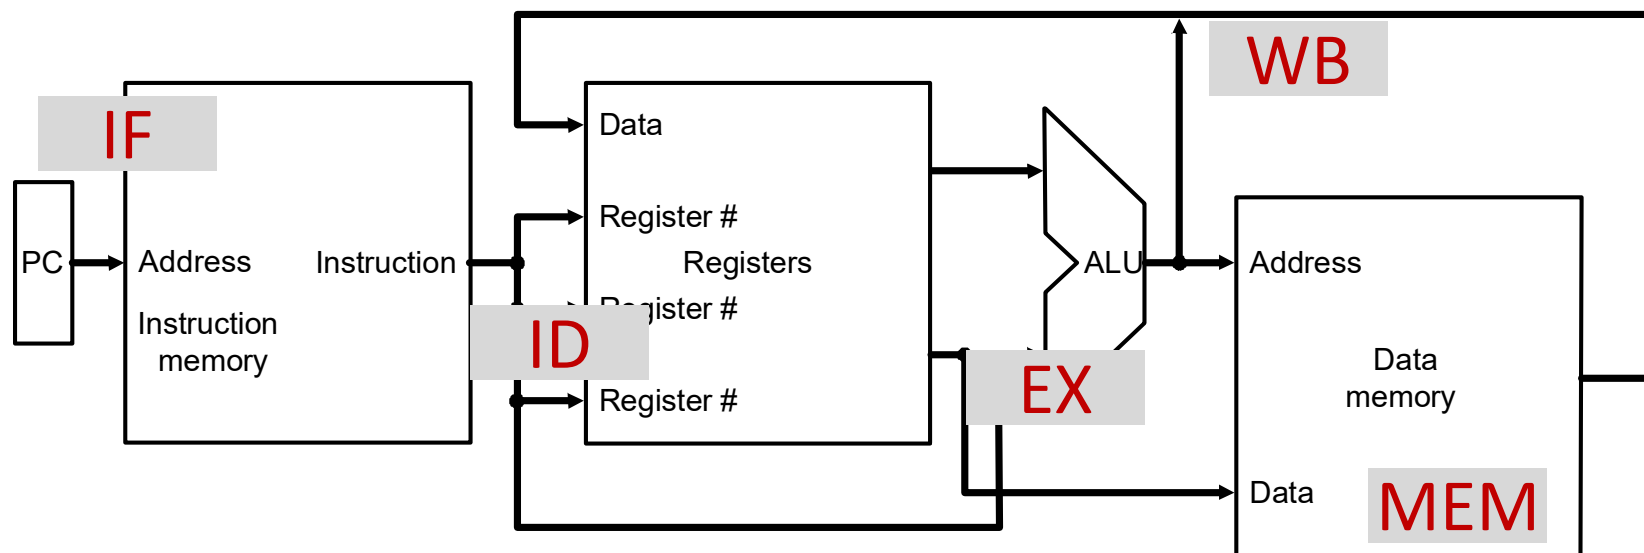
RISC Instruction Processing

- 5 generic steps
 - instruction fetch
 - instruction decode and operand fetch
 - ALU/execute
 - memory access
 - write-back




Coalescing and “External Fragmentation”

steps	IF	ID	EX	MEM	WB
R-type	✓	✓	✓		✓
I-type	✓	✓	✓		✓
LW	✓	✓	✓	✓	✓
SW	✓	✓	✓	✓	
Bxx/JALR	✓	✓	✓		--/✓
JAL	✓		✓		✓



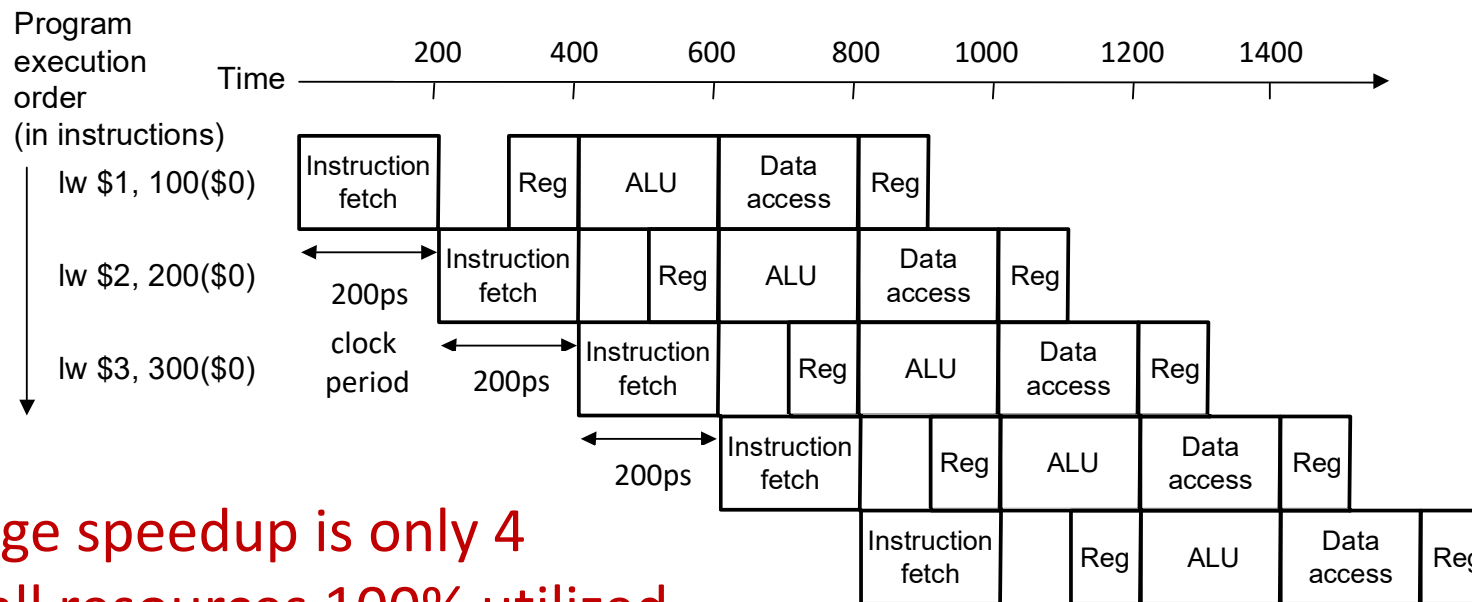
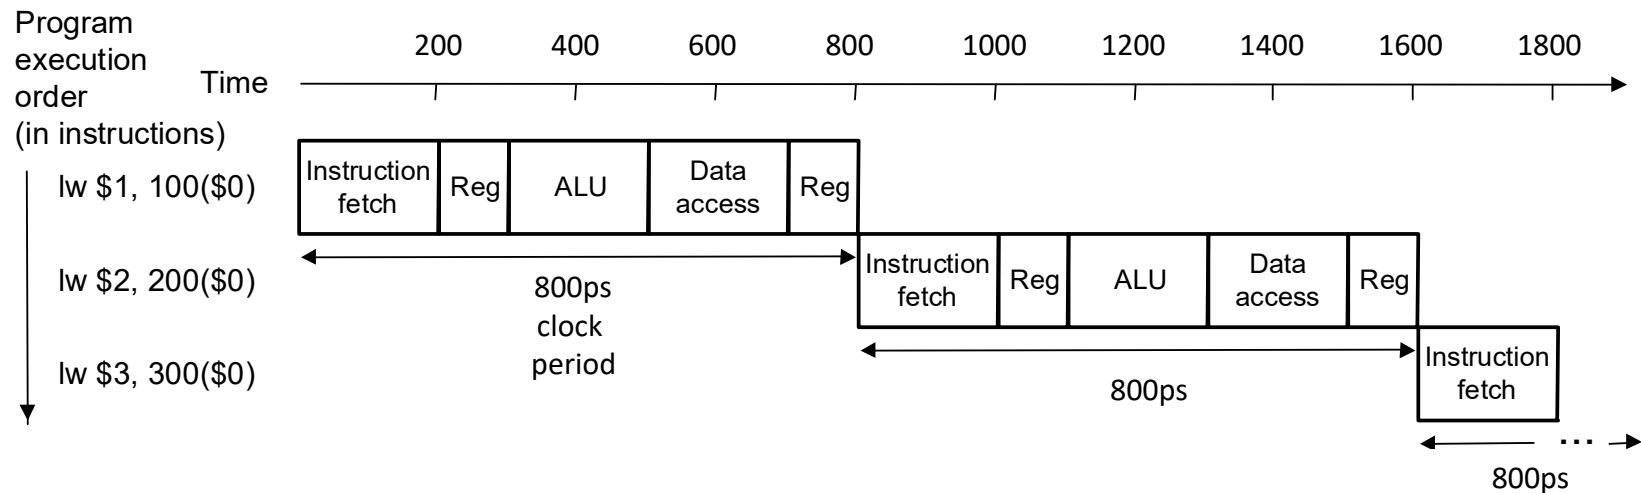
Dividing into Stages



Is this the correct partitioning?

Why not 4 or 6 stages? Why not different boundaries


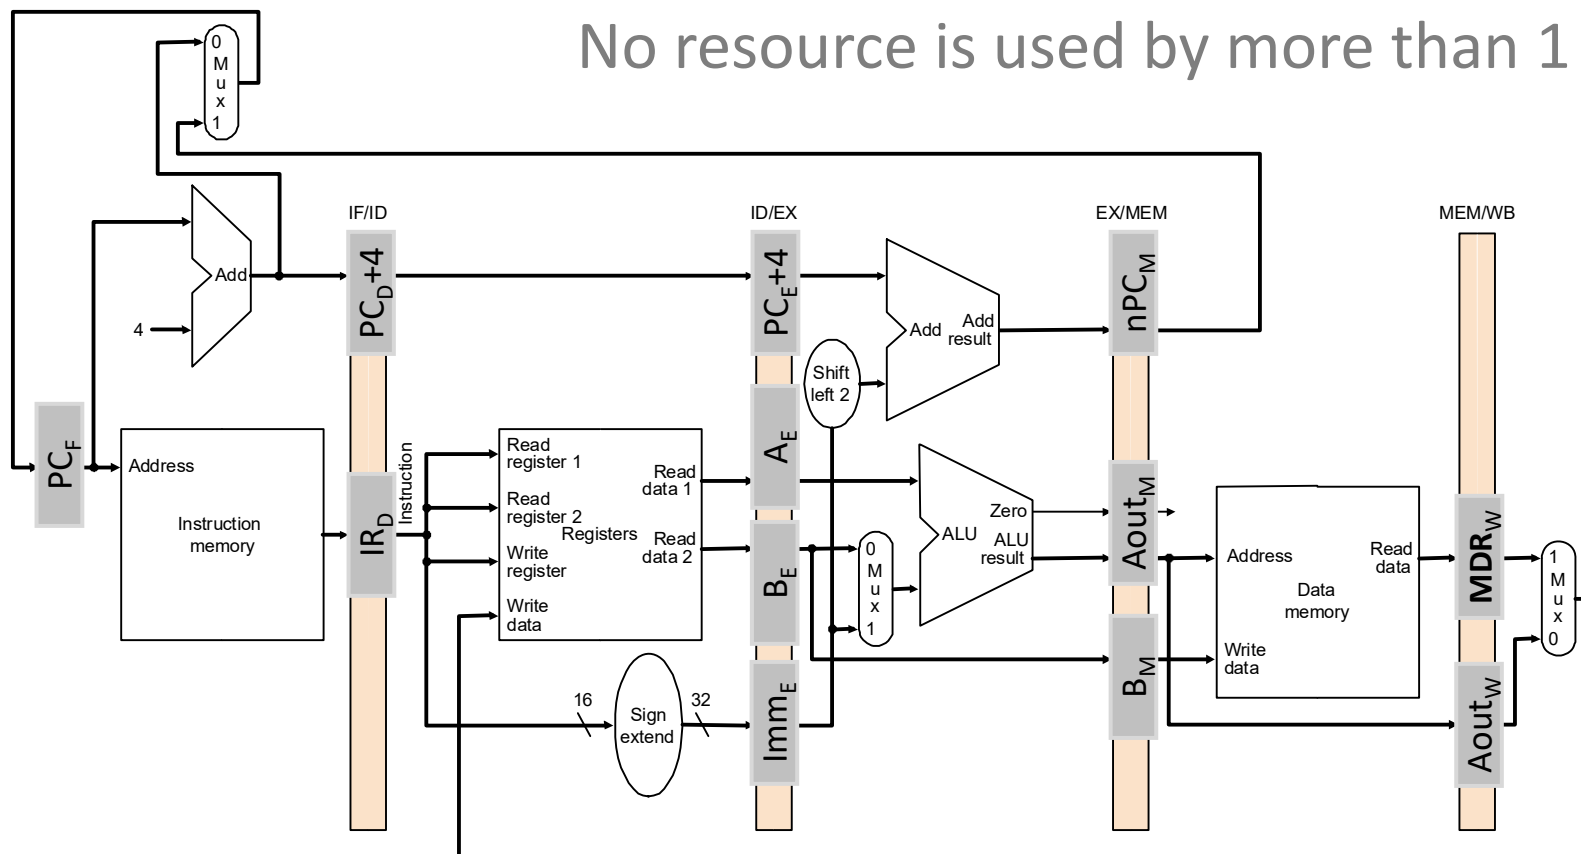
Internal and External Fragmentation



- 5-stage speedup is only 4
- Not all resources 100% utilized

Pipeline Registers

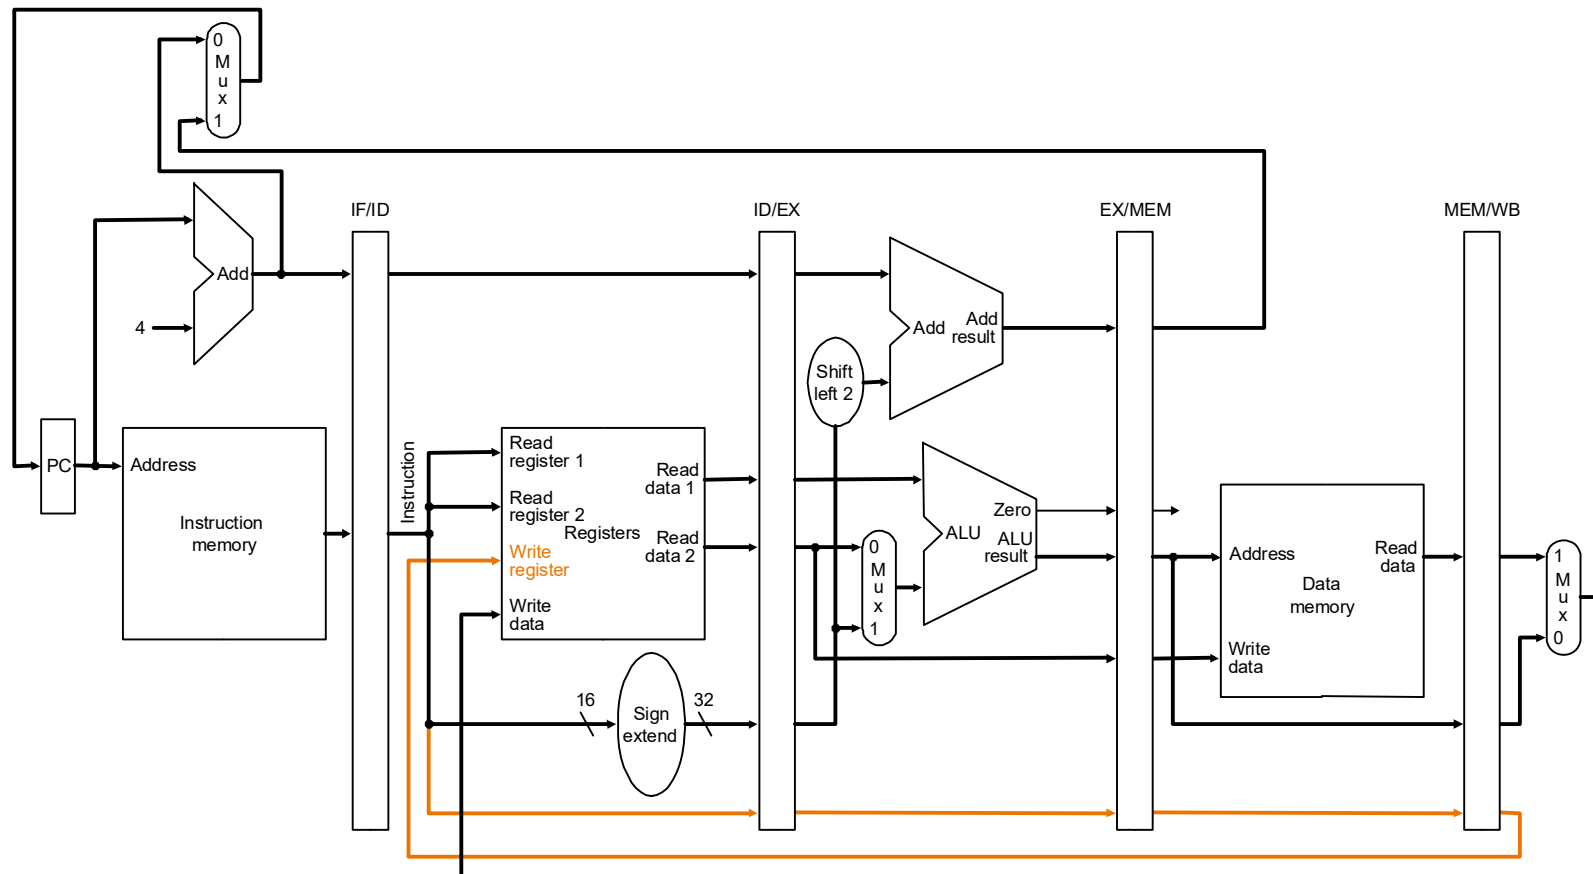
No resource is used by more than 1 stage!



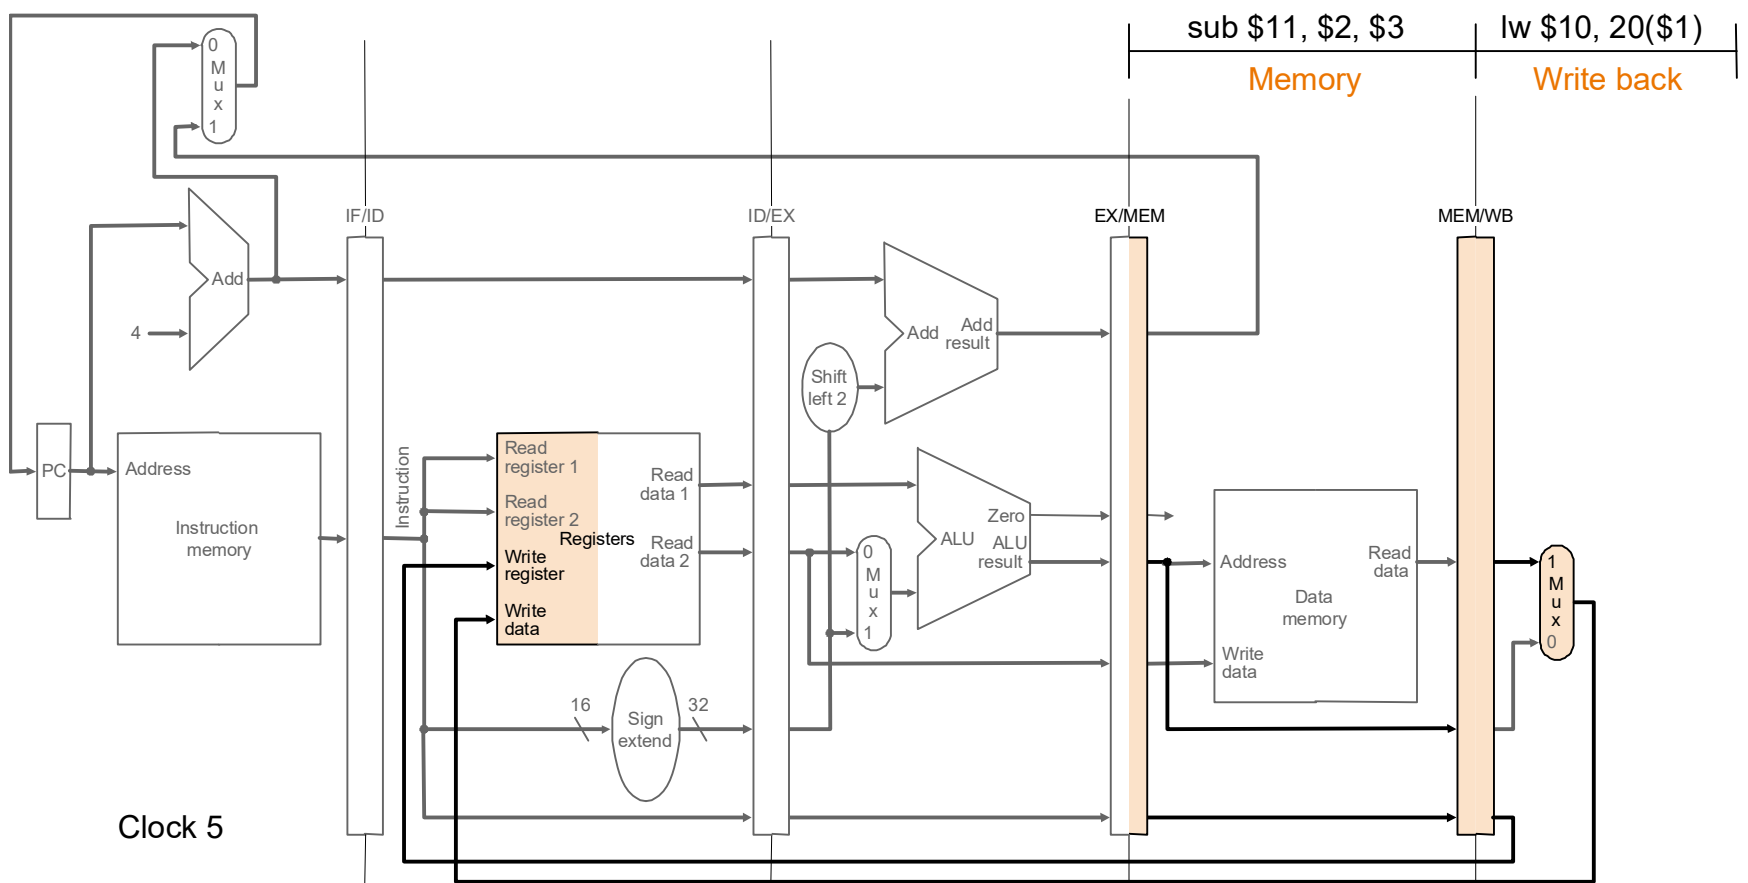
Based on original figure from [P&H CO&D, COPYRIGHT 2004 Elsevier. ALL RIGHTS RESERVED.]

18-447-S21-L07-S16, James C. Hoe, CMU/ECE/CALCM, ©2021

Pipelined Operation



Pipelined Operation



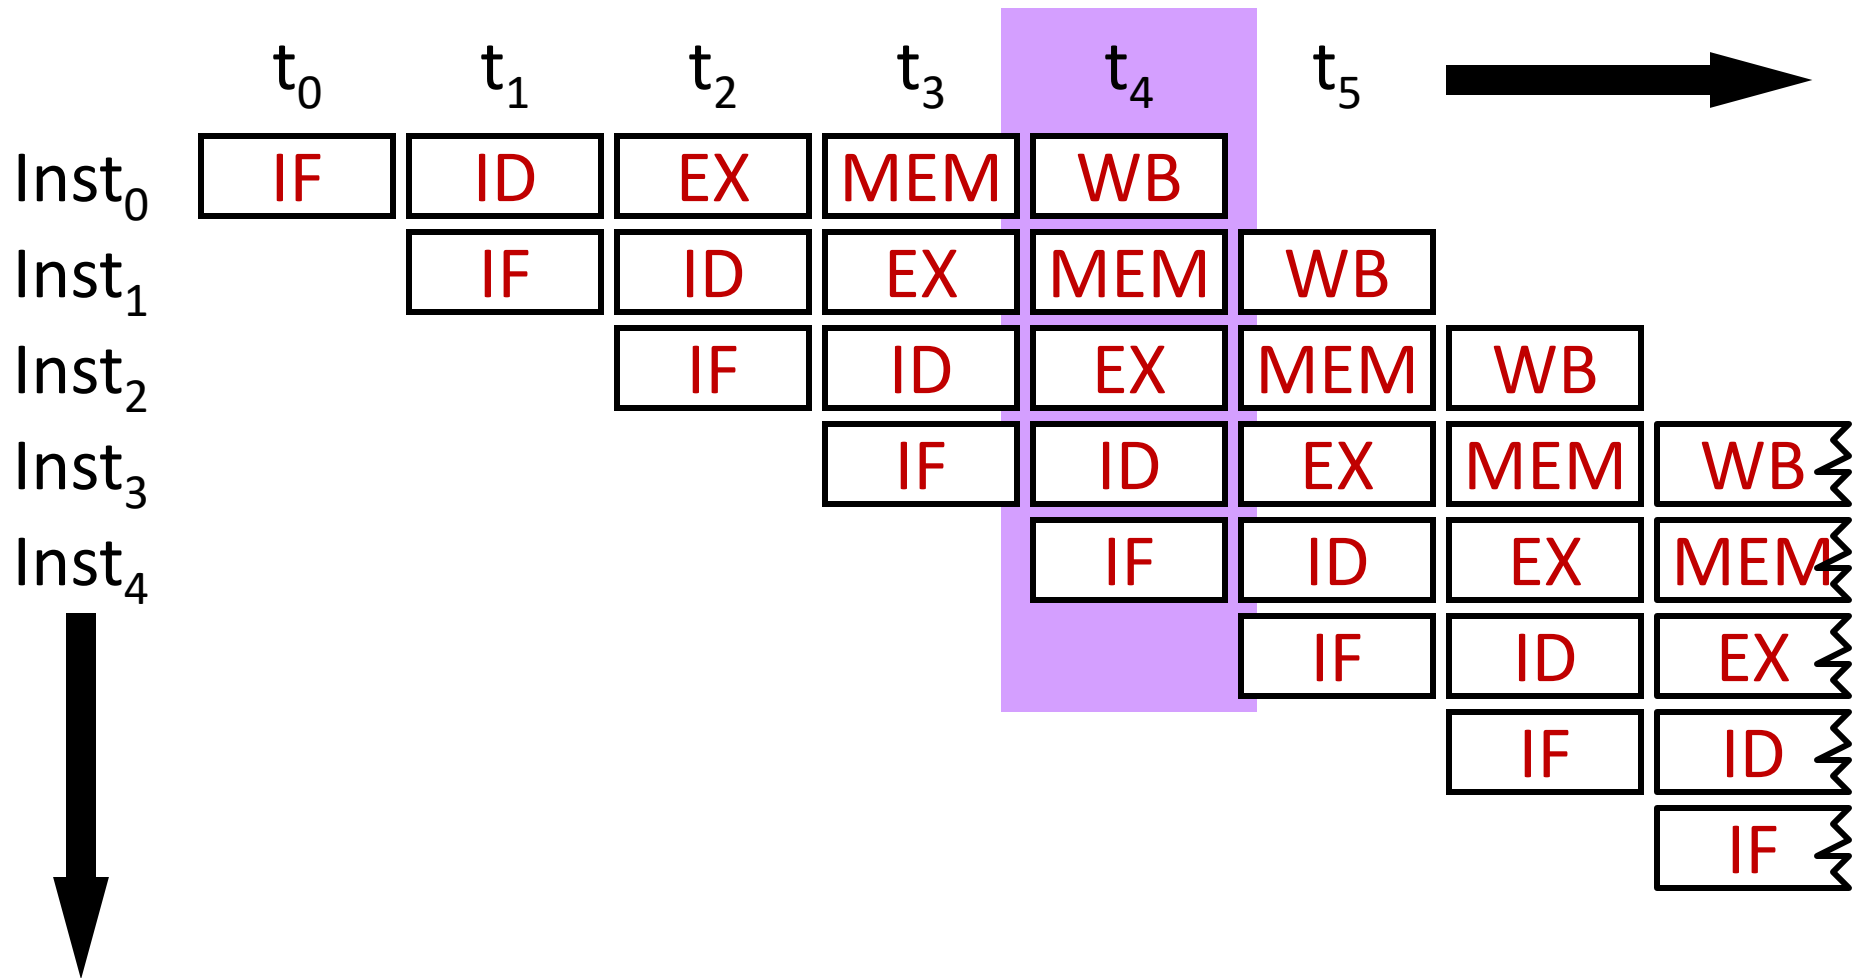
What if LW dest is \$2?

Illustrating Pipeline Operation:

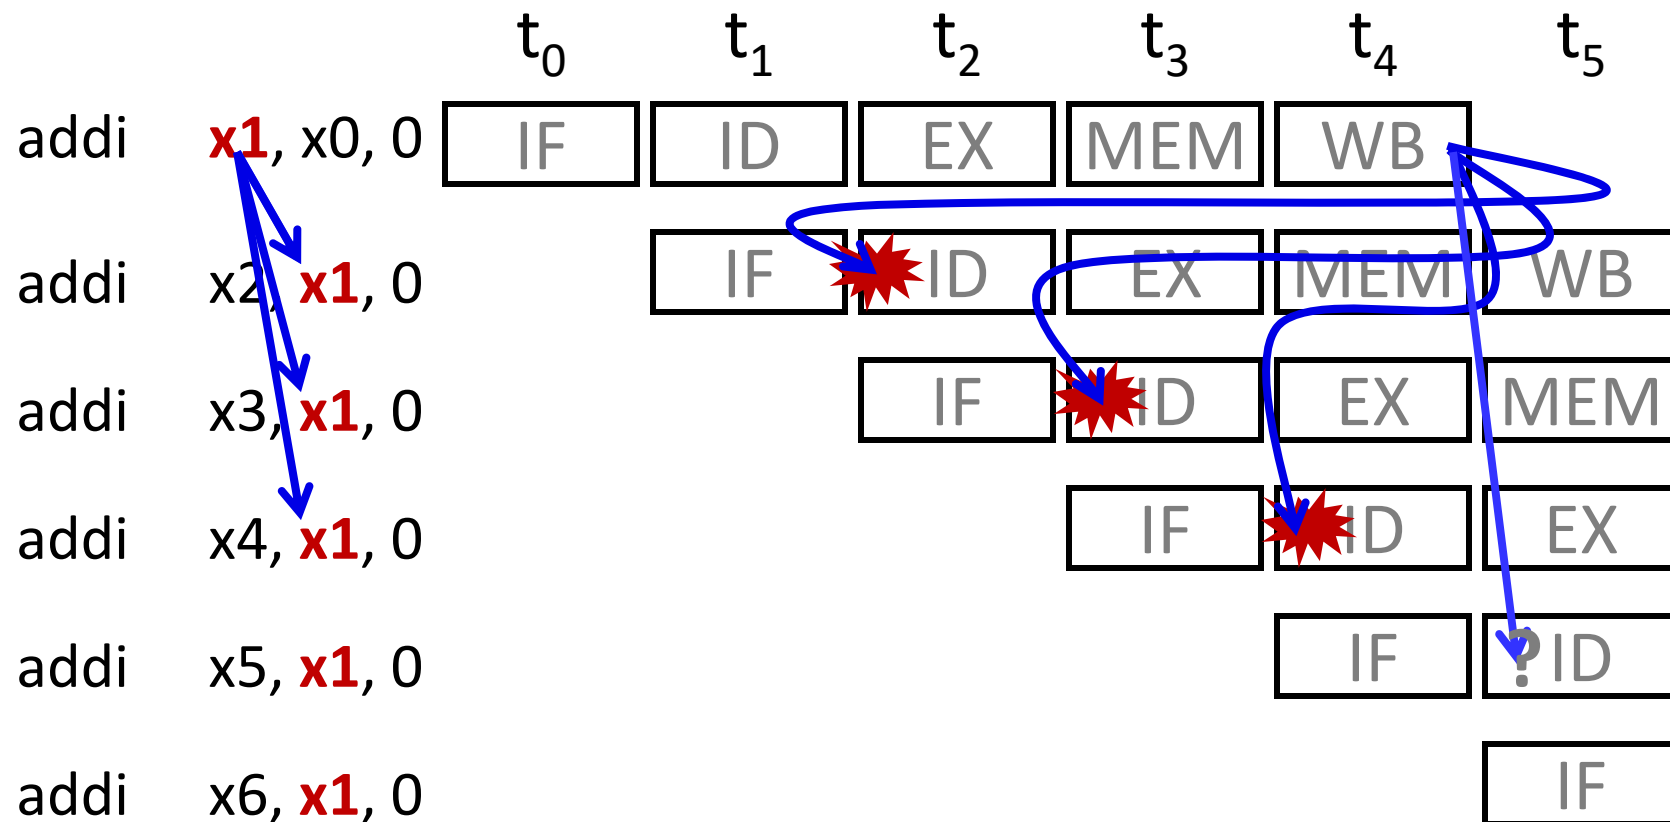
Resource View

	t_0	t_1	t_2	t_3	t_4	t_5	t_6	t_7	t_8	t_9	t_{10}
IF	I_0	I_1	I_2	I_3	I_4	I_5	I_6	I_7	I_8	I_9	I_{10}
ID		I_0	I_1	I_2	I_3	I_4	I_5	I_6	I_7	I_8	I_9
EX			I_0	I_1	I_2	I_3	I_4	I_5	I_6	I_7	I_8
MEM				I_0	I_1	I_2	I_3	I_4	I_5	I_6	I_7
WB					I_0	I_1	I_2	I_3	I_4	I_5	I_6

Illustrating Pipeline Operation: Operation View




Example: Read-after-Write Hazard



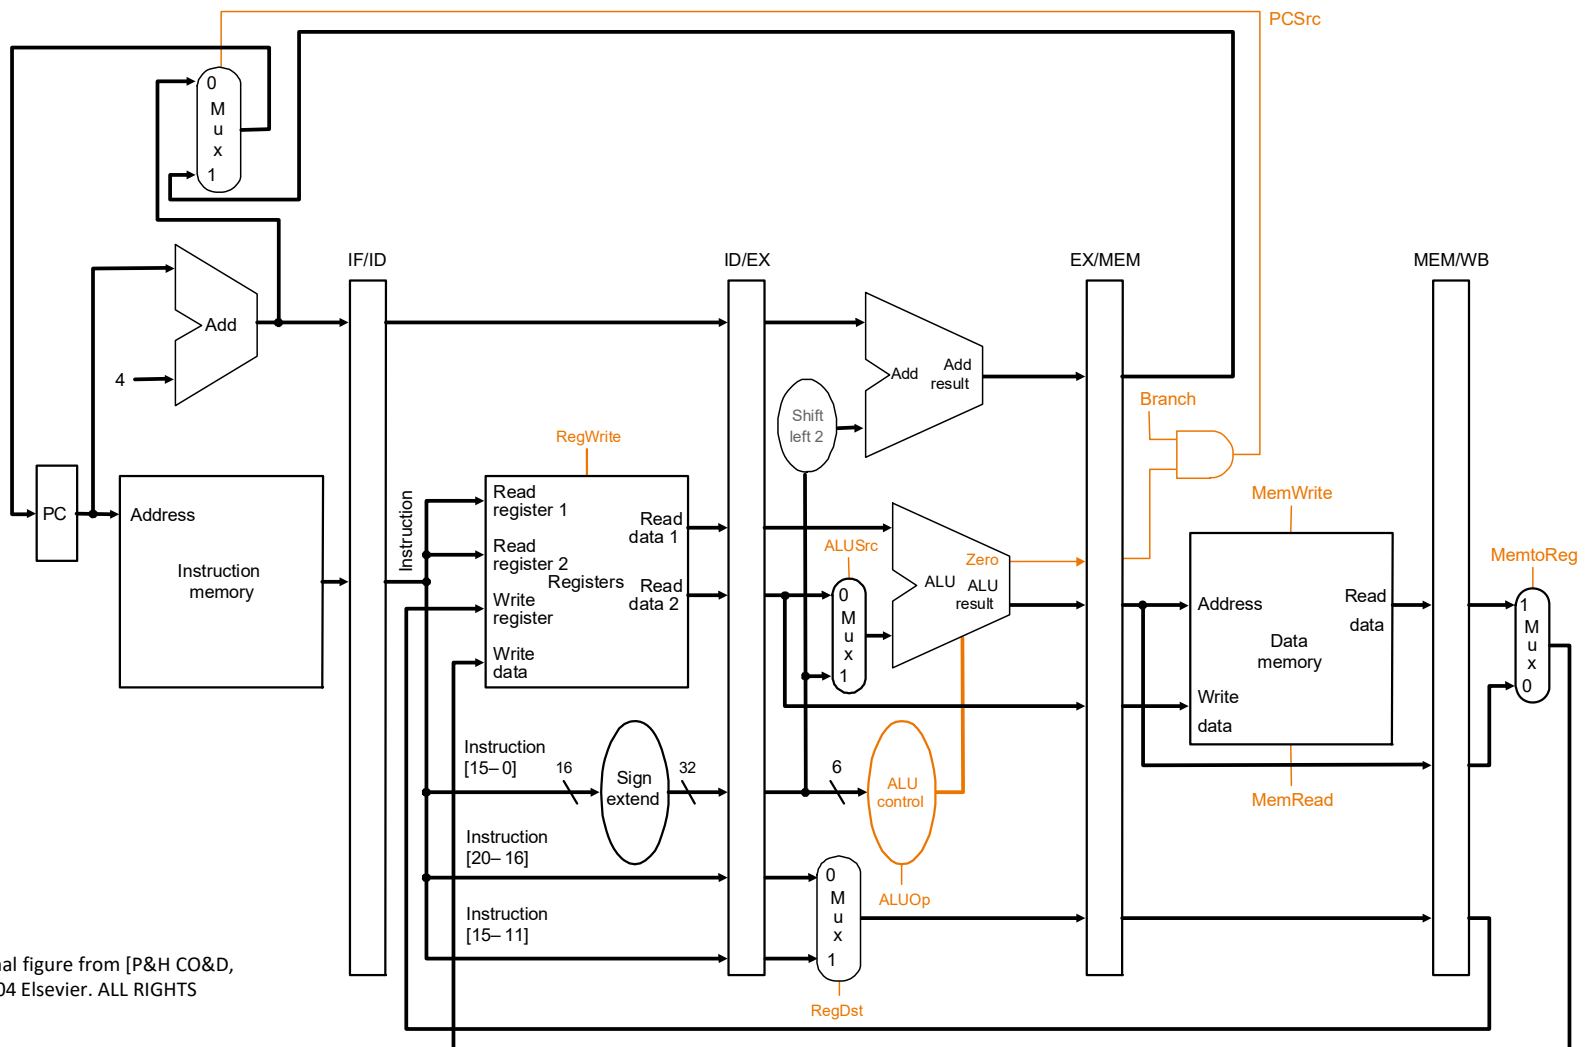
Example: Pipeline Stalls

	t_0	t_1	t_2	t_3	t_4	t_5	t_6	t_7	t_8	t_9	t_{10}
IF	I_0	I_1	I_2	I_3	I_4	I_4	I_4	I_4	I_5	I_6	I_7
ID		I_0	I_1	I_2	I_3	I_3	I_3	I_3	I_4	I_5	I_6
EX			I_0	I_1	I_2	\emptyset	\emptyset	\emptyset	I_3	I_4	I_5
MEM				I_0	I_1	I_2	\emptyset	\emptyset	\emptyset	I_3	I_4
WB					I_0	I_1	I_2	\emptyset	\emptyset	\emptyset	I_3



$I_2 = \text{addi } x1, x0, 0$; $I_3 = \text{addi } x2, x1, 0$;

Control Points




Based on original figure from [P&H CO&D,
COPYRIGHT 2004 Elsevier. ALL RIGHTS
RESERVED.]

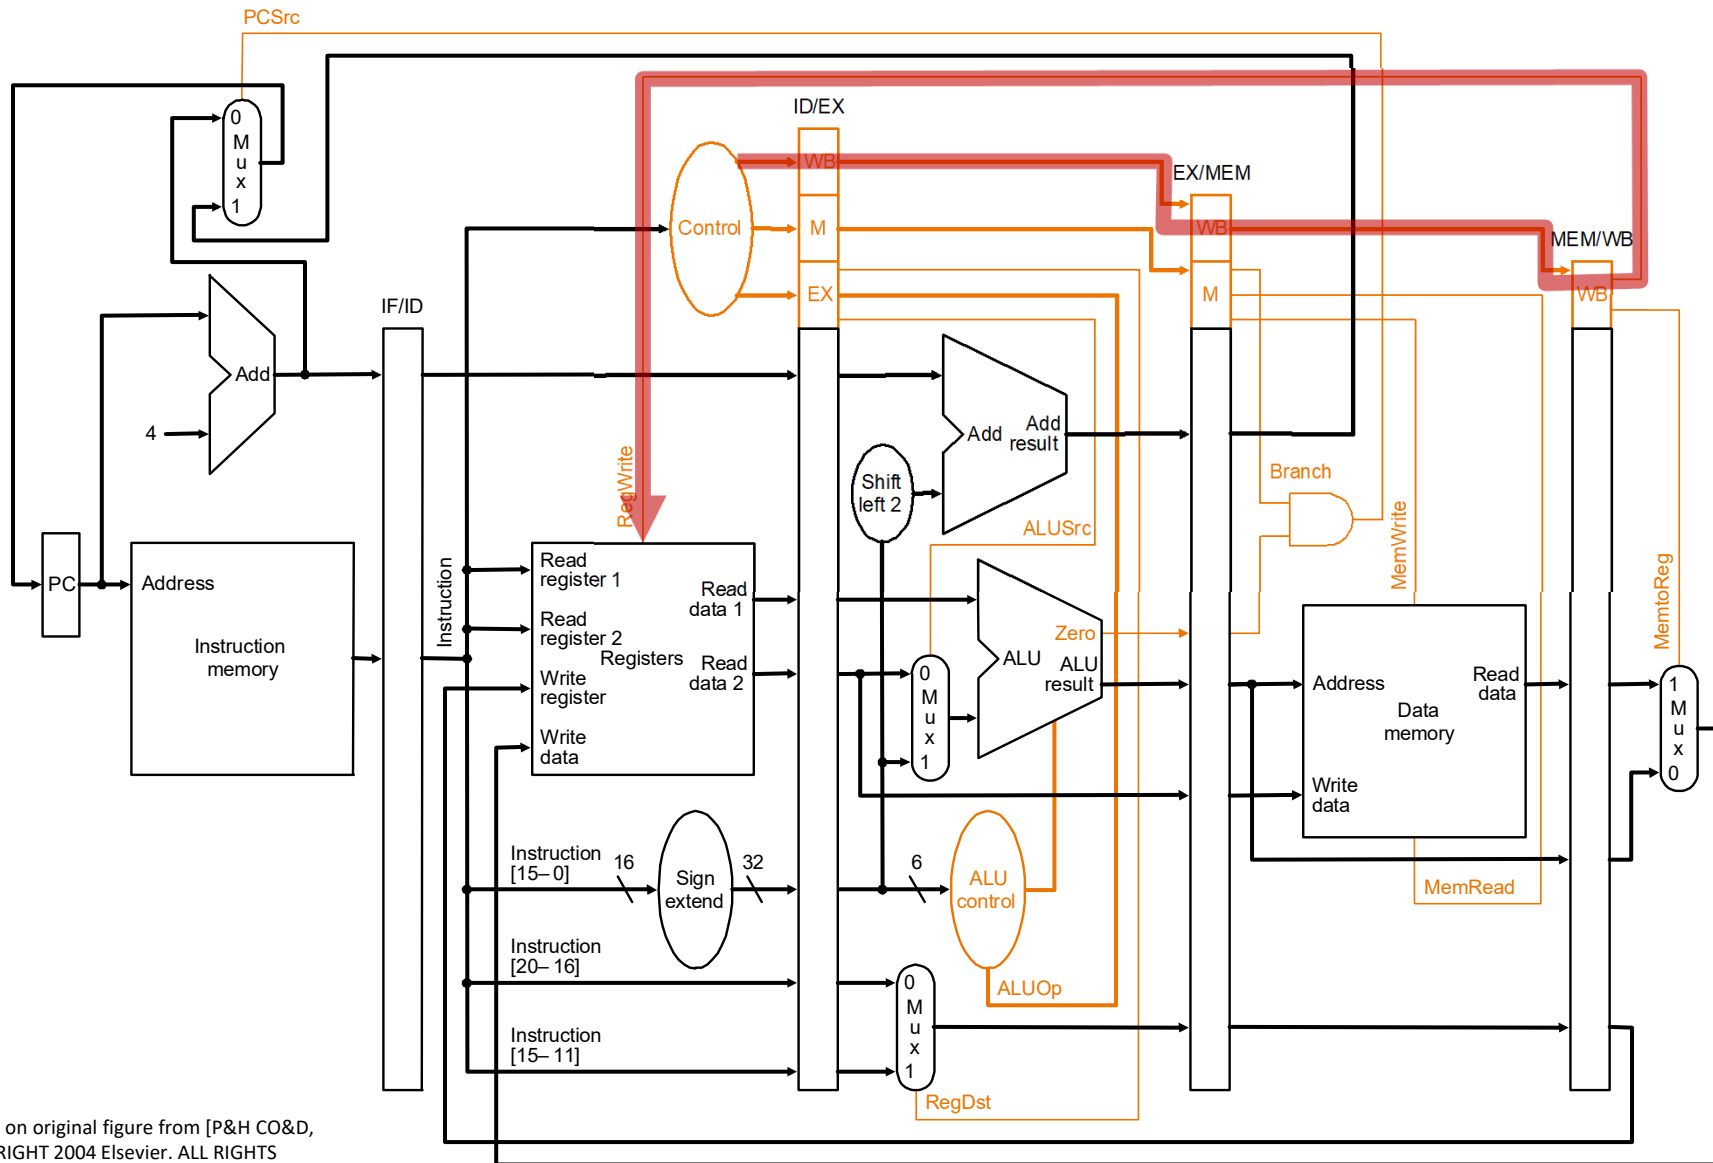
Identical set of control points as the single-cycle datapath!!

Sequential Control: Special Case

- For a given instruction
 - same control settings as single-cycle, but
 - control signals required at different cycles, depending on stage
 - decode once using the same logic as single-cycle and buffer control signals until consumed



Pipelined Control




Based on original figure from [P&H CO&D,
COPYRIGHT 2004 Elsevier. ALL RIGHTS
RESERVED.]

This is all there is to it (without hazards)!!

Instruction Pipeline Reality

- **Not identical tasks**
 - coalescing instruction types into one “multi-function” pipe
 - external fragmentation (some idle stages)
- **Not uniform suboperations**
 - group or sub-divide steps into stages to minimize variance
 - internal fragmentation (some too-fast stages)
- **Not independent tasks**
 - dependency detection and resolution
 - next lecture(s)

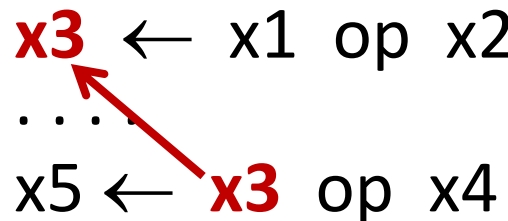


Even more messy if not RISC

Data Dependence

Data dependence

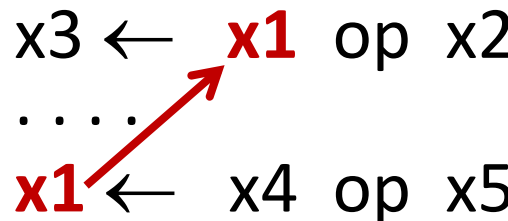
x3 ← x1 op x2
...
x5 ← **x3** op x4



Read-after-Write (RAW)

Anti-dependence

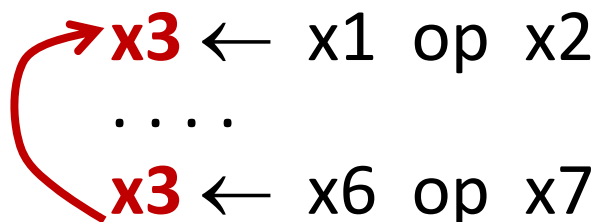
x3 ← **x1** op x2
...
x1 ← x4 op x5



Write-after-Read (WAR)

Output-dependence

x3 ← x1 op x2
...
x3 ← x6 op x7



Write-after-Write (WAW)

Don't forget memory instructions


Control Dependence

- C-Code

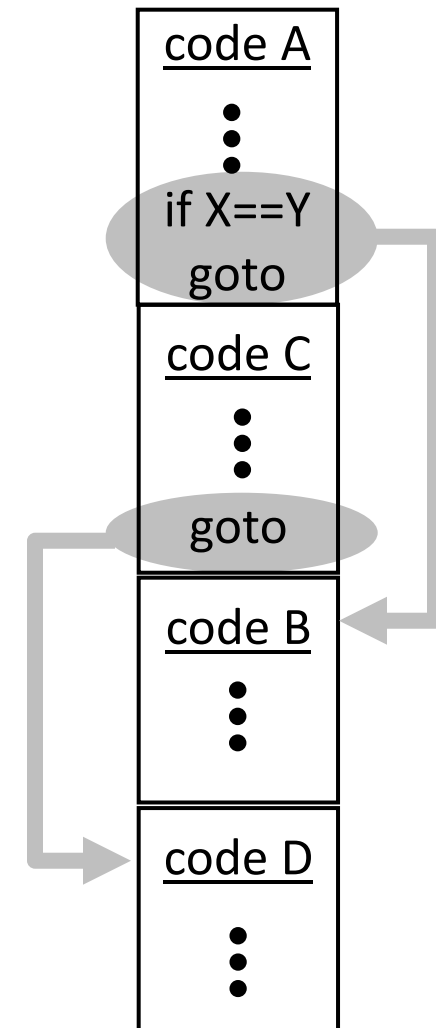
```

{ code A }
if X==Y then
    { code B }
else
    { code C }
{ code D }
  
```

Control Flow Graph



Assembly Code
(linearized)



Does B or C come after A?