

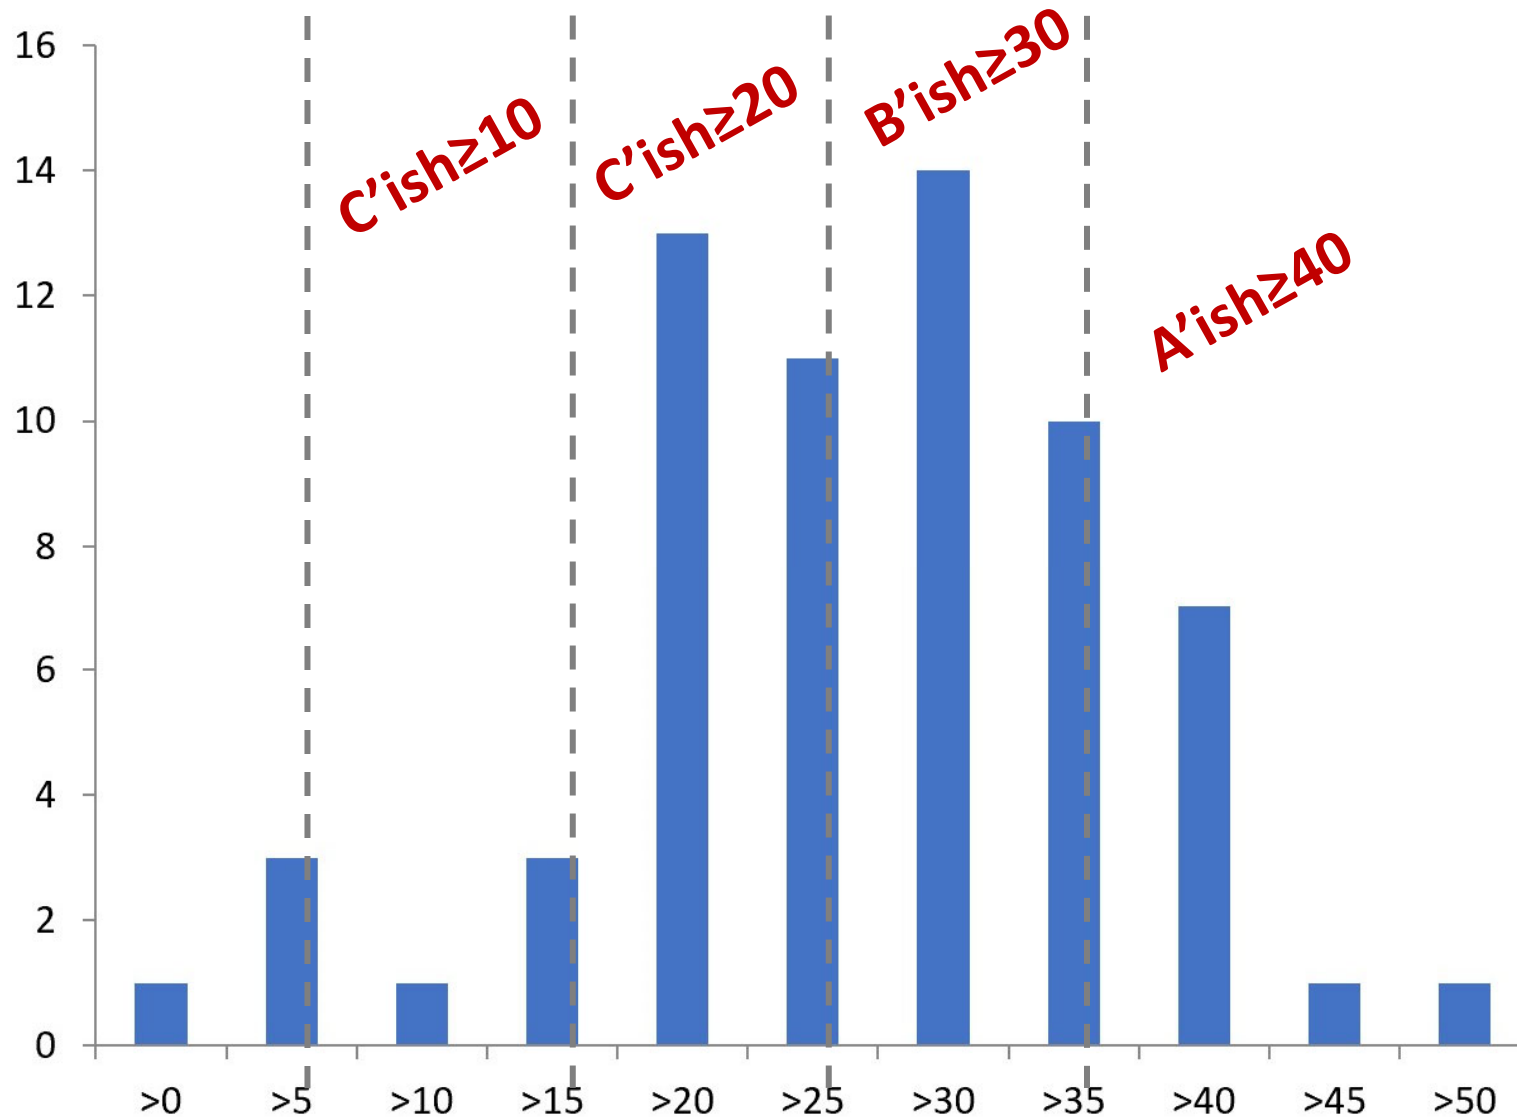
# **18-447 Lecture 12: Energy and Power**

James C. Hoe

Department of ECE

Carnegie Mellon University

# 2021 Midterm 1 Histogram



# Housekeeping

- Your goal today
  - a working understanding of energy and power
  - appreciate their significance in comp arch today
- Notices
  - Lab 2, due Friday 3/19 noon (not extended)
  - HW 3, due Monday 3/22 noon
- Readings
  - Design challenges of technology scaling, Borkar, 1999.
  - Synthesis Lectures (advanced optional):
    - Comp Arch Techniques for Power-Efficiency, 2008
    - Power-Efficient Comp Arch: Recent Advances, 2014

# First some intuitions


# Energy and Power

- CMOS logic transitions involve charging and discharging of parasitic capacitances
- Energy (Joule) dissipated as resistive heat when “charges” flow from VDD to GND
  - take a certain amount of energy per operation (e.g., addition, reg read/write, (dis)charge a node)
  - to the first order, energy  $\propto$  amount of compute
- Power (Watt=Joule/s) is rate of energy dissipation
  - more op/sec then more Joules/sec
  - to the first order, power  $\propto$  performance

**Power concerns usually more about heat removal**


# Heat and Thermal Resistance

- Resistive heat in the circuit must be removed in steadystate ([www.youtube.com/watch?v=BSGcnRanYMM](http://www.youtube.com/watch?v=BSGcnRanYMM))
- Can summarize everything between **circuit** and **ambient** by characteristic  $R_{\text{thermal}} = K/W$ 
  - convey power  $W$  in heat across temperature difference  $K = T_{\text{circuit}} - T_{\text{ambient}}$
- To dissipate more power in circuit
  1. let  $T_{\text{circuit}}$  get hotter (to a point)
  2. turn-up AC  $\Rightarrow$  lower  $T_{\text{ambient}}$
  3. better cooling  $\Rightarrow$  lower  $R$
- Economics/market driven choices



# Work and Perf. from Joules and Watt

- Fastest without energy/power awareness won't be fastest once constrained
  - power bounds performance directly
  - energy bounds work directly; want for lower J/op bounds performance indirectly



- Consider in context

recall that  $\text{power} \propto \text{perf}^{(\alpha > 1)}$

- mobile device: limited energy source, hard-to-cool form factor
- desktop: cooler size and noise
- data-center: electric bill, cooling capacity and cost

Ultimately driven by size/weight/\$\$\$

# Some (first-order) nitty-gritty



# Work and Runtime

- **Work**
  - scalar quantity for “amount of work” associated with a task
  - e.g., number of instructions to compute a SHA256 hash
- **$T = \text{Work} / k_{\text{perf}}$** 
  - runtime to perform a task
  - **$k_{\text{perf}}$**  is a scalar constant for the rate in which work is performed, e.g., “instructions per second”

# Energy and Power

- $E_{switch} = k_{switch} \cdot \text{Work}$ 
  - “switching” energy associated with task
  - $k_{switch}$  is a scalar constant for “energy per unit work”
- $E_{static} = k_{static} \cdot T = k_{static} \cdot \text{Work} / k_{perf}$ 
  - “leakage” energy just to keep the chip powered on
  - $k_{static}$  is the so called “leakage power”

*Faster execution means lower leakage energy???*

- $E_{total} = E_{switch} + E_{static} = k_{switch} \cdot \text{Work} + k_{static} \cdot \text{Work} / k_{perf}$
- $P_{total} = E_{total} / T = k_{switch} \cdot k_{perf} + k_{static}$

Static power can be 50% in high-perf processors

# In Short

- $T = \text{Work} / k_{perf}$

less work finishes faster

- $E = E_{switch} + E_{static} = (k_{switch} + k_{static}/k_{perf}) \cdot \text{Work}$

less work use less energy

- $P = P_{switch} + P_{static} = k_{switch} \cdot k_{perf} + k_{static}$

power independent of amount of work

- Reality check

- **Work** not a simple scalar, inst mix, dependencies ...


- **k's** are neither scalar nor constant

$k_{perf}$ : inst/sec

$k_{switch}$ : J/inst

$k_{static}$ : J/sec

# $k_{switch}$ , $k_{static}$ , $k_{perf}$ not independent




To increase  $k_{perf}$ : increase die area increases  $k_{switch}$  and  $k_{static}$   
faster transistors increases  $k_{static}$

# Why so important now?

# Technology Scaling for Dummies

- Planned scaling occurs in discrete “nodes” where each is  $\sim 0.7x$  of the previous in linear dimension
- Take the same design, reducing linear dimensions by  $0.7x$  (aka “gate shrink”) leads to **\*\*ideally\*\***
  - die area =  $0.5x$
  - delay =  $0.7x$ ; frequency =  $1.43x$
  - capacitance =  $0.7x$
  - $V_{dd}$  =  $0.7x$  (constant field) or  $1x$  (constant voltage)
  - power =  $0.5x$  (const. field) or  $1x$  (const. voltage)
- Take the same area, then
  - transistor count =  $2x$ , transistor speed =  $1.43x$
  - power =  $1x$  (const field) or  $2x$  (const voltage)

# The Other Moore's Law



# Moore's Law → Performance

- According to scaling theory

@constant complexity (“gate-shrink”):

1x transistors at 1.43x frequency

⇒ 1.43x performance at 0.5x power

@max complexity (“reticle limited”):

2x transistors at 1.43x frequency

⇒ 2.8x performance at constant power

- Historical (until 2005’ish), for high-perf CPUs

expected – ~2x transistors

higher – ~2x frequency (note: faster than scaling predicts)

– all together, ~2x performance at ~2x power

lower

higher


*Why so far off?*



# Performance (In)efficiency

- To hit “expected” performance target
  - push frequency harder by deepening pipelines
  - used the 2x transistors to build more complicated microarchitectures so fast/deep pipelines don’t stall (i.e., caches, BP, superscalar, out-of-order)
- The consequence of performance inefficiency is

limit of  
economical  
cooling [ITRS]




2005, Intel  
P4 Tehas 150W

Figure 8. Power dissipation projections.

[Borkar, IEEE Micro, July 1999]

# Moore's Law without Dennard Scaling



**Under fixed power ceiling, more ops/second  
only achievable if less Joules/op?**

**Frequency and Voltage Scaling:  
run slower at lower energy-per-op**

# Frequency and Voltage Scaling

- Switching energy per transition is  
 $\frac{1}{2}CV^2$  (modeling parasitic capacitance)
- Switching power at  $f$  transitions-per-sec is  
 $\frac{1}{2}CV^2f$
- To reduce power, slow down the clock
- If clock is slower ( $f'$ ), reduce supply voltage ( $V'$ ) too since transistors don't need to be as fast
  - reduced switching energy,  $\frac{1}{2}CV^2 \rightarrow \frac{1}{2}CV'^2$
  - lower  $V'$  also reduced leakage current/power

# Frequency Scaling (by itself)

- If  $Work / k_{perf} < T_{bound}$ , we can derate performance by frequency scaling by a factor  $S_{freq}$

$$(Work/k_{perf})/T_{bound} < S_{freq} < 1$$


s.t.  $k_{perf}' = k_{perf} S_{freq}$

- $T' = Work / (k_{perf} S_{freq})$ 
  - $1/S_{freq}$  longer runtime
- $E' = (k_{switch} + k_{static} / (k_{perf} S_{freq})) \cdot Work$ 
  - higher (leakage) energy due to longer runtime
- $P' = k_{switch} \cdot k_{perf} S_{freq} + k_{static}$ 
  - lower (switching) power due to longer runtime

Not such a good idea


# Intel P4 660 Frequency Scaling: FFT<sub>64K</sub>

circa 2005, 90nm



$$k_{perf} = 145 \text{ FFT}_{64K}/\text{sec}; \quad k_{switching} = 0.24 \text{ J/FFT}_{64K}; \quad k_{static} = 49.4 \text{ J/sec}$$

# Intel P4 660 Frequency Scaling: FFT<sub>64K</sub>




$$k_{perf}=145 \text{ FFT64K/sec}; k_{switching}=0.24 \text{ J/FFT64K}; k_{static}=49.4 \text{ J/sec}$$

# Frequency + Voltage Scaling

- Frequency scaling by  $s_{freq}$  allows supply voltage to be scaled by a corresponding factor  $s_{voltage}$
- $E \propto V^2$  thus  $k_{switch}'' = k_{switch} \cdot s_{voltage}^2$
- $k_{static}'' = k_{static} \cdot s_{voltage}^{2 \sim 3} \Leftarrow$  very gross approximation of something complicated
- $T'' = Work / (k_{perf} \cdot s_{freq})$   
 –  $1/s_{freq}$  longer runtime
- $E'' = (k_{switch} \cdot s_{voltage}^2 + k_{static} \cdot s_{voltage}^3 / k_{perf} \cdot s_{freq}) \cdot Work$
- $P'' = k_{switch} \cdot s_{voltage}^2 / k_{perf} \cdot s_{freq} + k_{static} \cdot s_{voltage}^3$   
 – superlinear reduction in power and energy to performance degradation

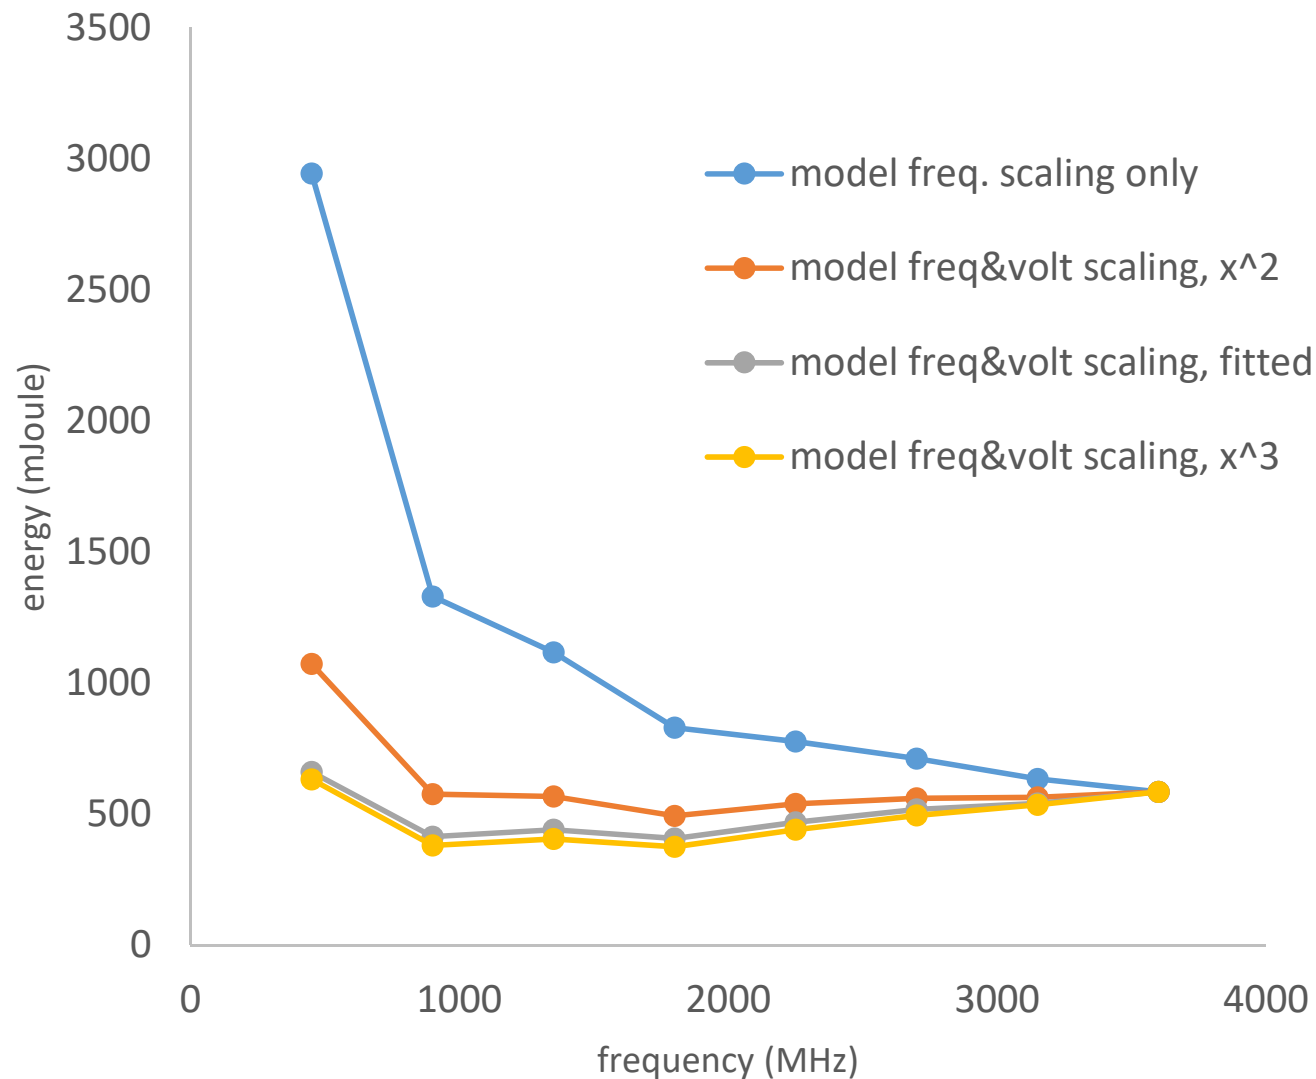


# Intel P4 660 F+V Scaling: FFT<sub>64K</sub>



circa 2005, 90nm

# Intel P4 660 F+V Scaling: FFT<sub>64K</sub>




circa 2005, 90nm

# Parallelization:

run faster at lower energy-per-op  
by  
running slower at lower energy-per-op

# Cost of Performance in Power



[Energy per Instruction Trends in Intel®  
Microprocessors, Grochowski et al., 2006]

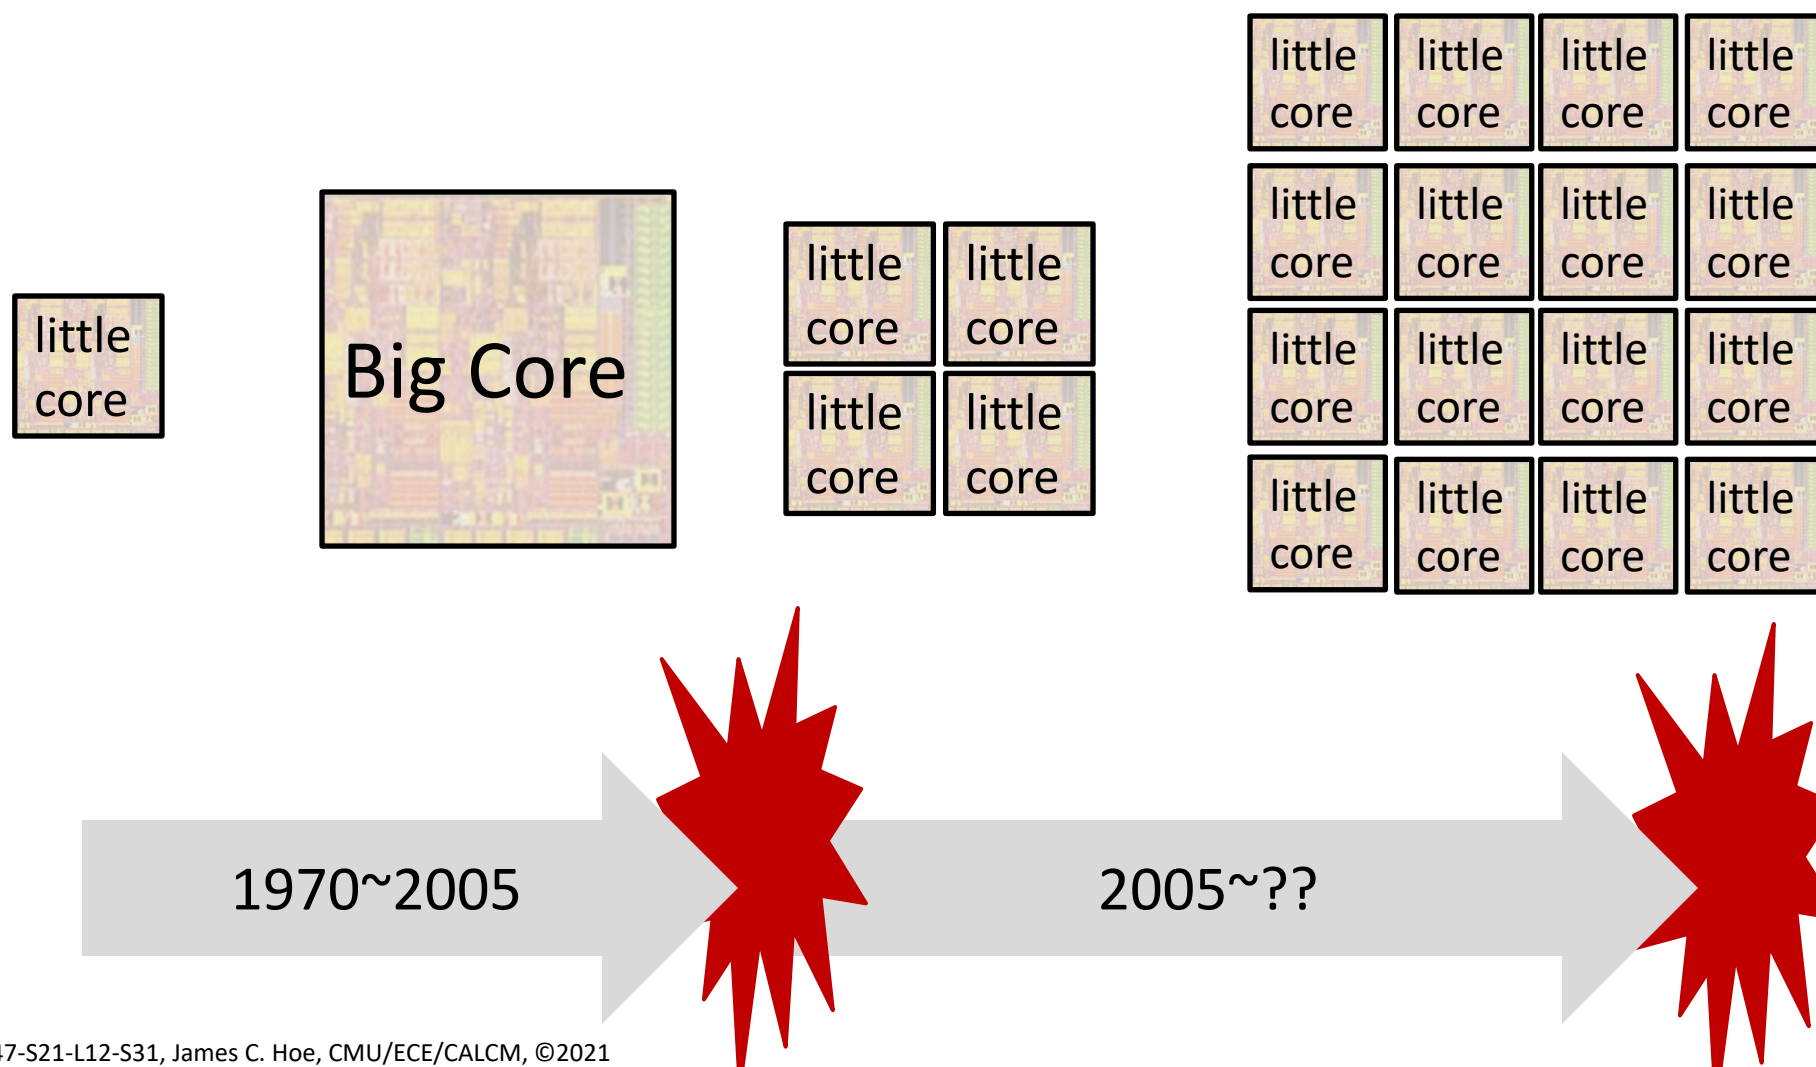
# Parallelization

- Ideal parallelization over  $N$  CPUs (to go fast)
  - $T = \text{Work} / (k_{\text{perf}} \cdot N)$
  - $E = (k_{\text{switch}} + k_{\text{static}} / k_{\text{perf}}) \cdot \text{Work}$   
 $N$ -times static power, but  $N$ -times faster runtime
  - $P = N (k_{\text{switch}} \cdot k_{\text{perf}} + k_{\text{static}})$
- Alternatively, forfeit speedup for power and energy reduction by  $s_{\text{freq}} = 1/N$  (assume  $s_{\text{voltage}} \approx s_{\text{freq}}$  below)
  - $T = \text{Work} / k_{\text{perf}}$
  - $E'' = (k_{\text{switch}} / N^2 + k_{\text{static}} / (k_{\text{perf}} N)) \cdot \text{Work}$
  - $P'' = k_{\text{switch}} \cdot k_{\text{perf}} / N^2 + k_{\text{static}} / N$
- Also works with using  $N$  slower-simpler CPUs

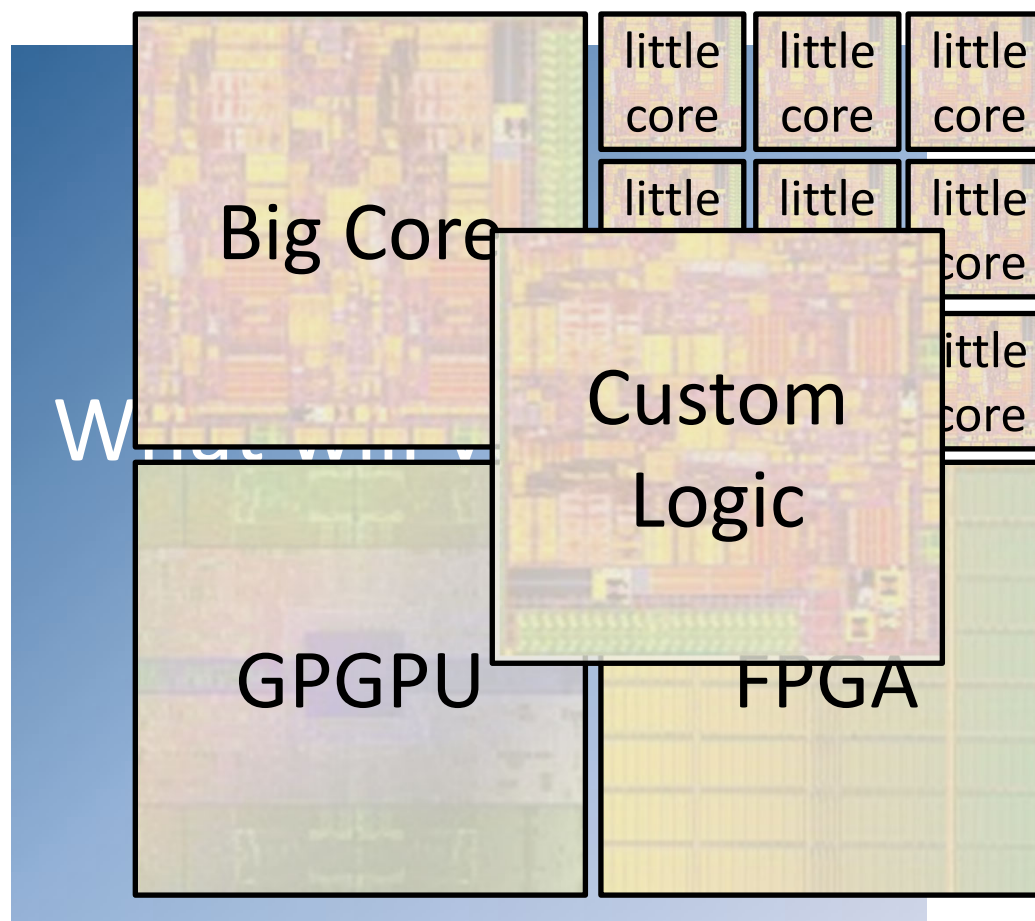
# So what is the problem?

- “Easy” to pack more cores on a die to stay on Moore’s law for “aggregate” or “throughput” performance
  - How to use them?
    - life is good if your **N** units of work is **N** independent programs  $\Rightarrow$  just run them
    - what if your **N** units of work is **N** operations of the same program?  $\Rightarrow$  rewrite as parallel program
    - what if your **N** units of work is **N** sequentially dependent operations of the same program?  $\Rightarrow$  ??
- How many cores can you use up meaningfully?**

# Moore's Law Scaling with Cores



# Remember: it is all about **Perf/Watt** and **Ops/Joules**



We talk about HW specialization in a later lecture