

# **18-447 Lecture 14: Memory Hierarchy**

James C. Hoe

Department of ECE

Carnegie Mellon University

# Housekeeping

- Your goal today
  - understand memory system and memory hierarchy design in big pictures
- Notices
  - Handout #10: Lab 3, **due Friday 4/9 noon**
  - Handout #11: HW 3 solutions
  - Handout #12: HW 4, **due Monday 4/12 noon**
- Readings
  - P&H Ch5 for the next many lectures

# Wishful Memory

- So far we imagined
  - a program owns contiguous 4GB private memory
  - 16 ExaByte if RV64I
  - a program can access anywhere in 1 proc. cycle
- We are in good company

4.1. Ideally one would desire an indefinitely large memory capacity such that any particular aggregate of 40 binary digits, *word* (cf. 2.3), would be immediately available—i.e. in a tin

---- Burks, Goldstein, von Neumann, 1946

# The Reality

- Can't afford/don't need as much memory as size of address space
  - RV32I said 4GB addr “space” not 4GB memory
- Can't find memory technology that is affordable in GByte and also cycle in GHz
- Most systems multi-task several programs
- But, “magic” memory is nevertheless a useful approximation of reality due to
  - memory hierarchy: appear large and fast
  - virtual memory: appear contiguous and private

cover this  
part first

cover this  
part later

# Memory Hierarchy: The Principles at Work

# The Law of Storage

- Bigger is slower
  - SRAM                      512 Bytes                      @ sub-nsec
  - SRAM                      KByte~MByte                      @ nsec
  - DRAM                      GByte                      @ ~50 nsec
  - SSD                      TByte                      @ msec
  - Hard Disk              TByte                      @ ~10 msec
- Faster is more expensive (dollars and chip area)
  - SRAM                      ~\$10K per GByte
  - DRAM                      ~\$10 per GByte
  - “Drives”                      ~\$0.1 per GByte

Note: order-of-magnitude only & changes with time

**How to make memory bigger, faster and cheaper?**

# Memory Locality

- “Typical” programs have strong locality in memory references—instruction and data  
we put them there ... loops, arrays, and structs ...
- Temporal: after accessing **A**, how many other distinct addresses before accessing **A** again
- Spatial: after accessing **A**, how many other distinct addresses before accessing a “near-by” **B**
- **Corollary:** a program with strong temporal and spatial locality must be accessing only a compact “**working set**” at a time

# Memoization

- If something is costly to compute, save the result to be reused
- With strong reuse
  - storing just a small number of frequently used results can avoid most recomputations
- With poor reuse
  - storing a large number of different results that are rarely or never reused
  - locating the needed result from a large number of stored ones can itself become as expensive as computing

being size  
effective



# Cost Amortization

- **overhead**: one-time cost to set up
- **unit-cost**: cost for each unit of work
- total cost = overhead + unit-cost x N
- average cost = total cost / N

$$= ( \text{overhead} / N ) + \text{unit-cost}$$

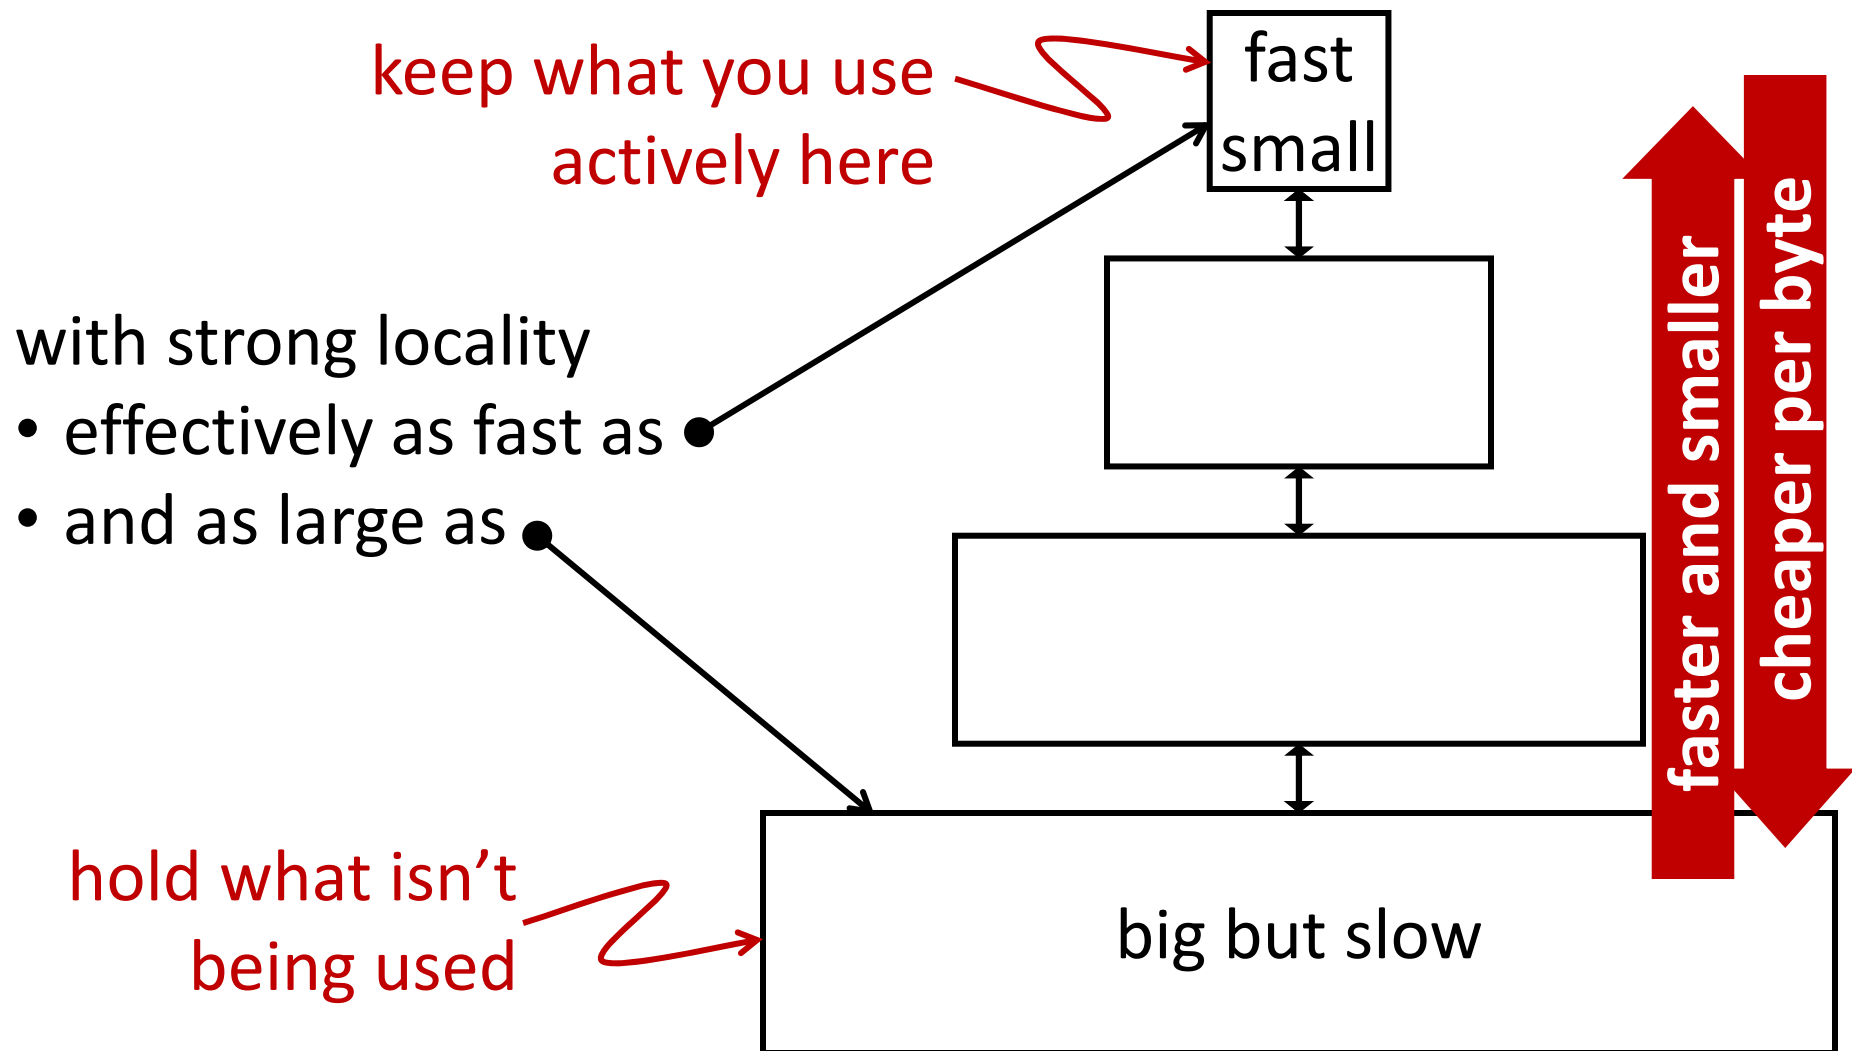
the essence of amortization

being time  
effective

In memoization, high up-front cost to compute once is no problem if results reused many times

# Putting the principles to work

# Memory Hierarchy



# Managing Memory Hierarchy

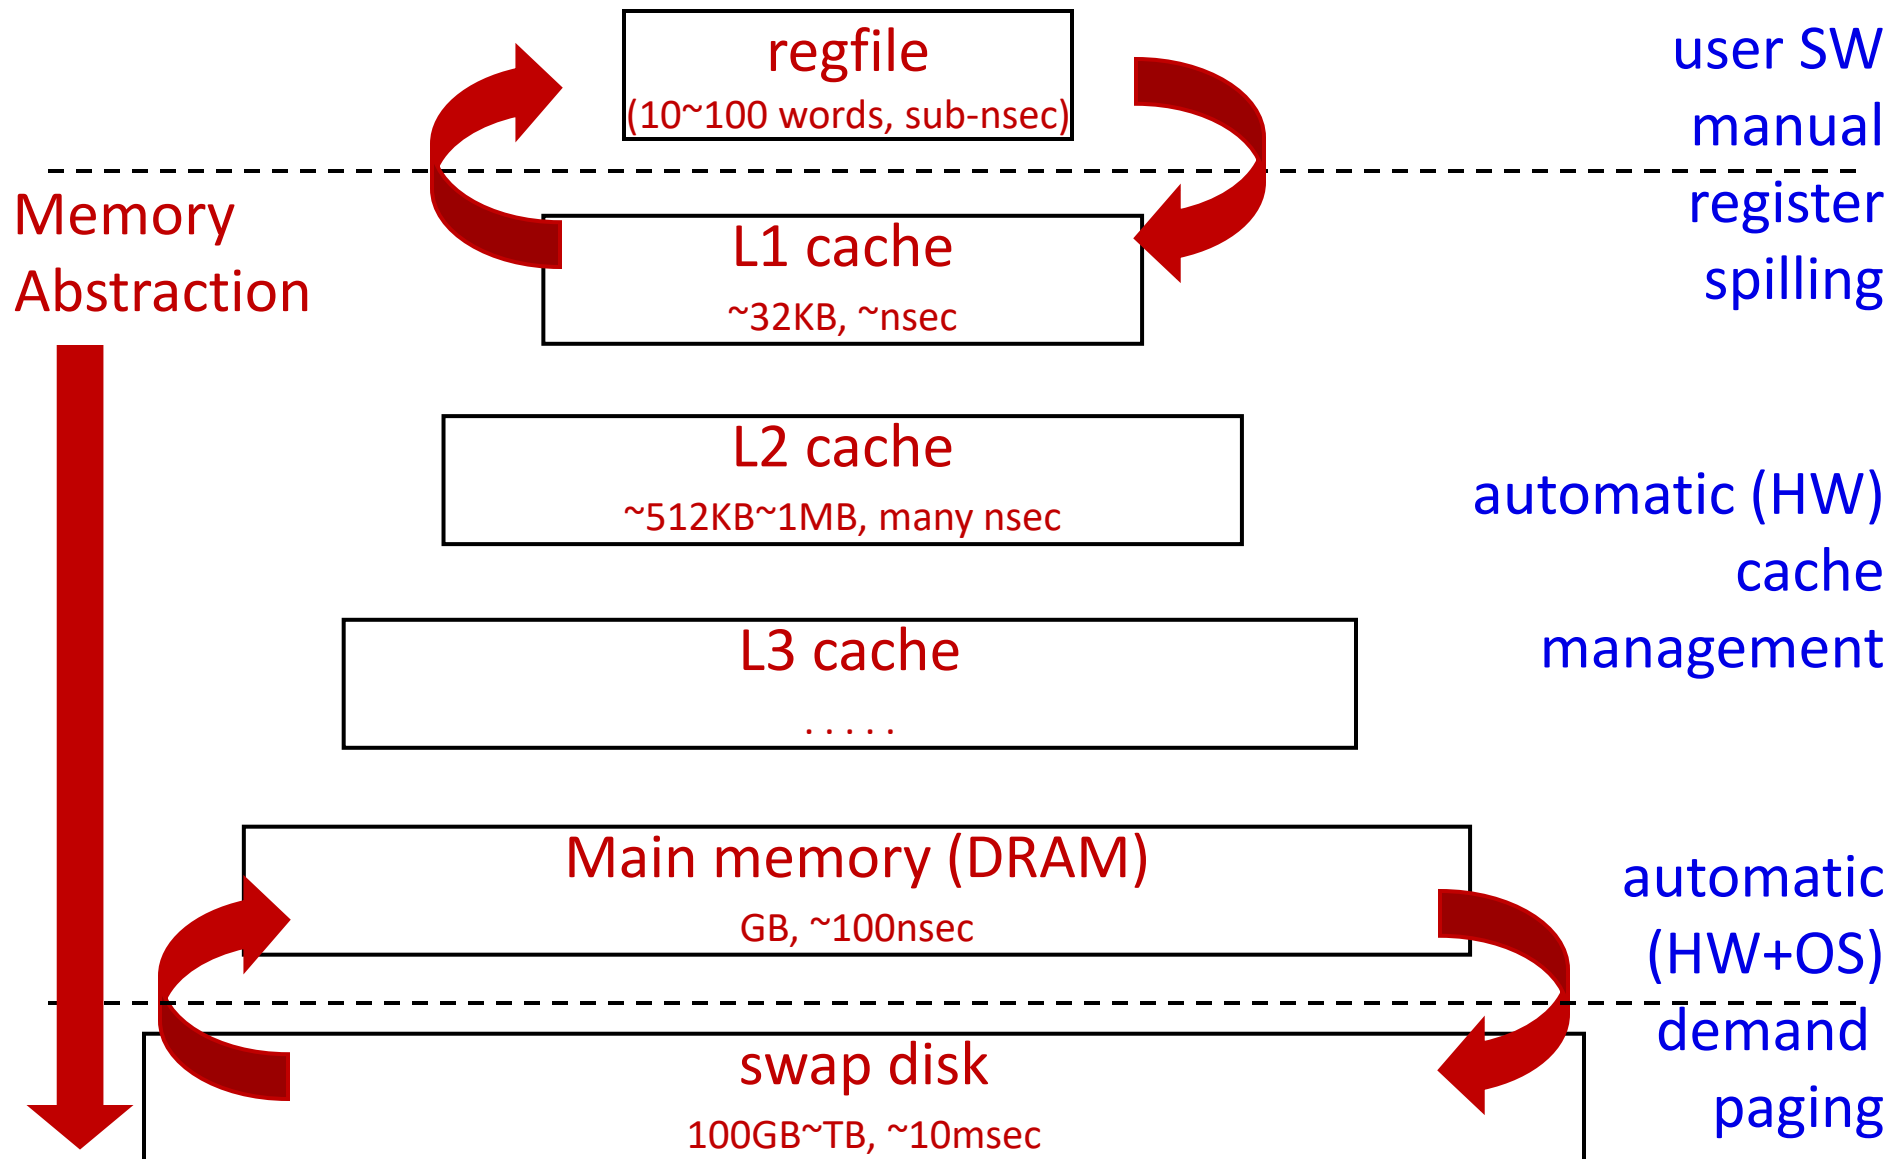
- Copy data between levels explicitly and manually
  - vacuum tubes vs Selectron (von Neumann paper)
  - “core” vs “drum” memory in the 50’s
  - “scratchpad” SRAM used on modern embedded and DSP

Register file is a level of storage hierarchy

- Single address space, automatic management
  - as early as ATLAS, 1962
  - common in today’s fast processor with slow DRAM
  - programmers don’t need to know about it for typical programs to be both fast and correct

What about atypical programs?

# Modern Storage Hierarchy



# Average Memory Access Time

- Memory hierarchy level  $L_1$  has raw access time of  $t_1$
- Average access time  $T_1$  is longer than  $t_1$ 
  - a chance (hit-rate  $h_1$ ) you find what you want  $\Rightarrow t_1$
  - a chance (miss-rate  $m_1$ ) you don't find it  $\Rightarrow t_1 + T_2$
  - $T_1 = h_1 \cdot t_1 + m_1 \cdot (t_1 + T_2)$  and  $h_1 + m_1 = 1.0$
- In general

$$T_i = h_i \cdot t_i + m_i \cdot (t_i + T_{i+1})$$

$$T_i = t_i + m_i \cdot T_{i+1}$$

think of this as  
“miss penalty”

Note:  $h_i$  and  $m_i$  are of references missed at  $L_{i-1}$

$$h_{\text{bottom-most}} = 1.0$$

$$T_i = t_i + m_i \cdot T_{i+1}$$

- Goal: achieve desired  $T_1$  within allowed cost

$T_i \approx t_i$  is not a goal:

- Keep  $m_i$  low
  - increase capacity  $C_i$  lowers  $m_i$ , but increases  $t_i$
  - lower  $m_i$  by smarter management, e.g.,
    - replacement: anticipate what you don't need
    - prefetching: anticipate what you will need
- Keep  $T_{i+1}$  low
  - reduce  $t_{i+1}$  with faster next level memory leads to increased cost and/or reduced capacity
  - better solved by adding intermediate levels

# Memory Hierarchy Design

- DRAM
  - optimized for capacity-per-dollar (cost)
  - $T_{\text{DRAM}}$  is essentially same regardless of capacity
- SRAM
  - optimized for latency at given capacity
  - tunable tradeoff between capacity and latency
  - possible,  $t = O(\sqrt{\text{capacity}})$
- Memory hierarchy bridges the difference between CPU speed and DRAM speed
  - $T_{\text{pclk}} \approx T_{\text{DRAM}} \Rightarrow$  no hierarchy needed
  - $T_{\text{pclk}} \ll T_{\text{DRAM}} \Rightarrow$  one or more levels of increasingly larger but slower SRAMs to minimize  $T_1$



## Aside: Why is DRAM slow?

- DRAM fabrication at forefront of VLSI, but scaled with Moore's law in capacity and cost not speed
- Between 1980 ~ 2004
  - 64K bit → 1024M bit (exponential ~55% annual)
  - 250ns → 50ns (linear)
- A deliberate engineering choice
  - memory capacity needs to grow linearly with processing speed in a balanced system – Amdahl's Other Law
  - DRAM/processor speed difference reconcilable by SRAM cache hierarchies (L1, L2, L3, .....)

Pareto-optimal faster/smaller/more-costly DRAM do exist

# Intel P4 Example

## (very fast, very deep pipeline)

- 90nm, 3.6 GHz
- 16KB L1 D-cache
  - $t_1 = 4$  cyc int (9 cycle fp)
- 1024KB L2 D-cache
  - $t_2 = 18$  cyc int (18 cyc fp)
- Main memory
  - $t_3 = \sim 50\text{ns}$  or 180 cyc
- Notice:
  - best case latency is not 1 cycle
  - worst case access latency is 300+ cycles depending on exactly what happens


if  $m_1=0.1, m_2=0.1$   
 $T_1=7.6, T_2=36$

if  $m_1=0.01, m_2=0.01$   
 $T_1=4.2, T_2=19.8$

if  $m_1=0.05, m_2=0.01$   
 $T_1=5.00, T_2=19.8$

if  $m_1=0.01, m_2=0.50$   
 $T_1=5.08, T_2=108$

# Working Set/Locality/Miss Rate



# Don't Forget Bandwidth and Energy

- Assume RISC pipeline 1GHz and IPC=1
  - 4GB/sec of instruction fetch bandwidth
  - 1GB/sec load and 0.6GB/sec store (if 25% LW and 15% SW, Agerwala&Cocke)
  - multiply by number of cores if multicore
- DDR4 ~20GB/sec/channel (under best-case access pattern) and ~10 Watt at full blast
- With memory hierarchy

$$BW_{i+1} = BW_1 \cdot \prod_1^i m_j$$

Critical for multicore and GPU


# Now we can talk about caches . . .

Generically in computing, any structure that “memoizes” frequently repeated computation results to save on the cost of reproducing the results from scratch, e.g. a web cache

# Cache in Computer Architecture

- An invisible, automatically-managed memory hierarchy
- Program expects reading  $M[A]$  to return most-recently written value, with or without cache
- Cache keeps “copies” of frequently accessed DRAM memory locations in a small fast memory
  - service load/store using fast memory copies if found
  - transparent to program if memory idempotent (L13)
  - funny things happen if mmap'ed or if memory can change (e.g., by other cores or DMA)

# Cache Interface for Dummies




- Like the magic memory
  - present address, R/W command, etc
  - result or update valid after a short/fixed latency
- Except occasionally, cache needs more time
  - will become valid/ready eventually
  - what to do with pipeline until then? Stall!!

# The Basic Problem

- Potentially  $M=2^m$  bytes of memory, how to keep “copies” of most frequently used locations in  $C$  bytes of fast storage where  $C \ll M$
- Basic issues (intertwined)
  - (1) when to cache a “copy” of a memory location
  - (2) where in fast storage to keep the “copy”
  - (3) how to find the “copy” later on (*LW and SW only give indices into  $M$* )




# Basic Operation (demand-driven version)

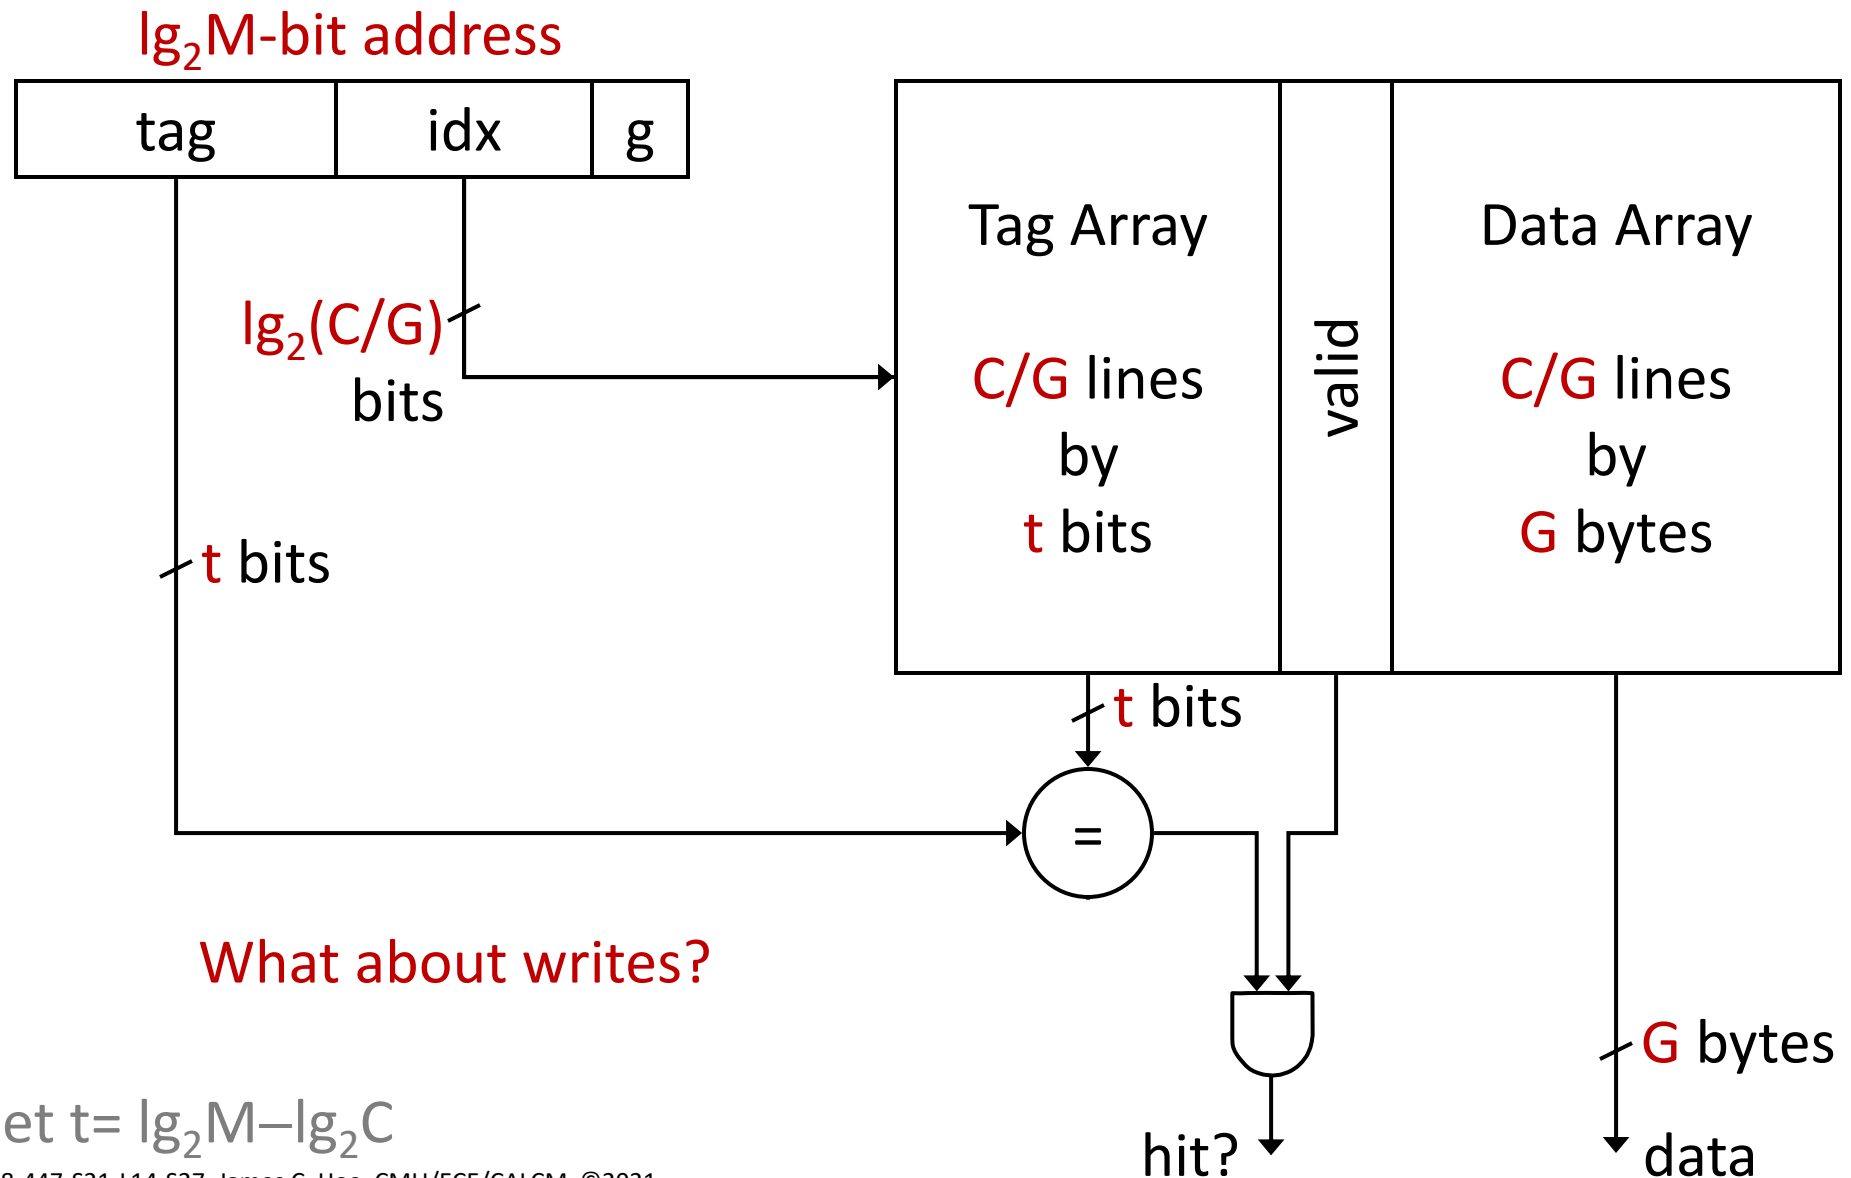


# Basic Cache Parameters

- **$M = 2^m$**  : size of address space in bytes  
sample values:  $2^{32}$ ,  $2^{64}$
  - **$G = 2^g$**  : cache access granularity in bytes  
sample values: 4, 8
- 
- **$C$**  : “capacity” of cache in bytes  
sample values: 16 KByte (L1), 1 MByte (L2)



# Direct-Mapped Placement (v1)



$$\text{let } t = \lg_2 M - \lg_2 C$$

# Storage Overhead and **B**lock Size

- For each cache block of **G** bytes, also storing “**t+1**” bits of tag (where **t** =  $\lg_2 M - \lg_2 C$ )
  - if **M** =  $2^{32}$ , **G** = 4, **C** = 16K =  $2^{14}$
  - $\Rightarrow$  **t** = 18 bits for each 4-byte block

60% overhead; 16KB cache actually 25.5KB SRAM

- Solution: “amortize” tag over larger **B**-byte block
  - manage **B/G** consecutive words as indivisible unit
  - if **M** =  $2^{32}$ , **B** = 16, **G** = 4, **C** = 16K
  - $\Rightarrow$  **t** = 18 bits for each 16-byte block

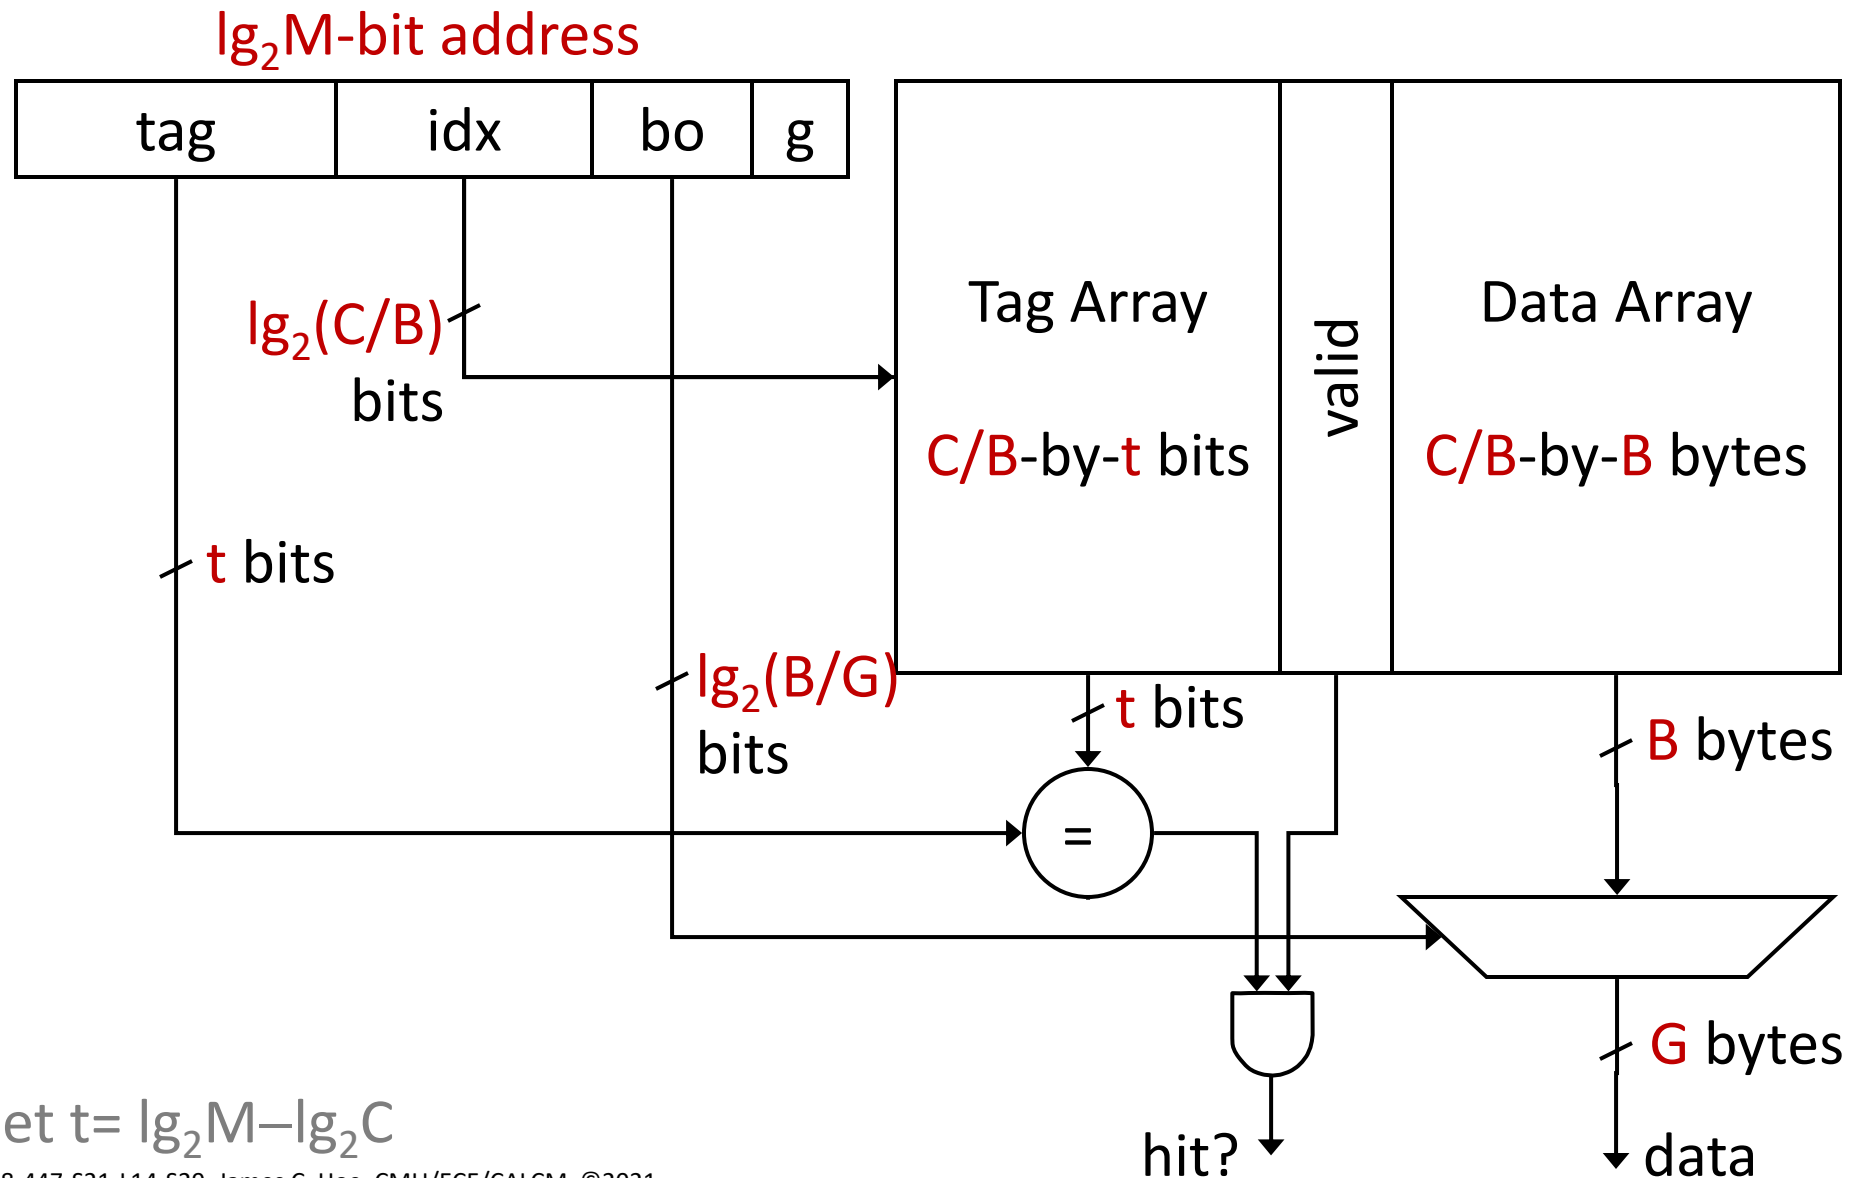
15% overhead; 16KB cache actually 18.4KB SRAM

- spatial locality also says this is good (*Q1: when*)

**B**

- Larger caches want even bigger blocks

# Direct-Mapped Placement (final)



# Basic Cache Parameters

ISA

- **$M = 2^m$**  : size of address space in bytes  
sample values:  $2^{32}$ ,  $2^{64}$
  - **$G = 2^g$**  : cache access granularity in bytes  
sample values: 4, 8
- 

Implementation

- **$C$**  : “capacity” of cache in bytes  
sample values: 16 KByte (L1), 1 MByte (L2)
- **$B = 2^b$**  : “block size” in bytes  
sample values: 16 (L1), >64 (L2)
- **$a$**  : “associativity” of the cache  
sample values: 1, 2, 4, 5(?), ... “ $C/B$ ”

to be continued

C/a should be a 2-power