

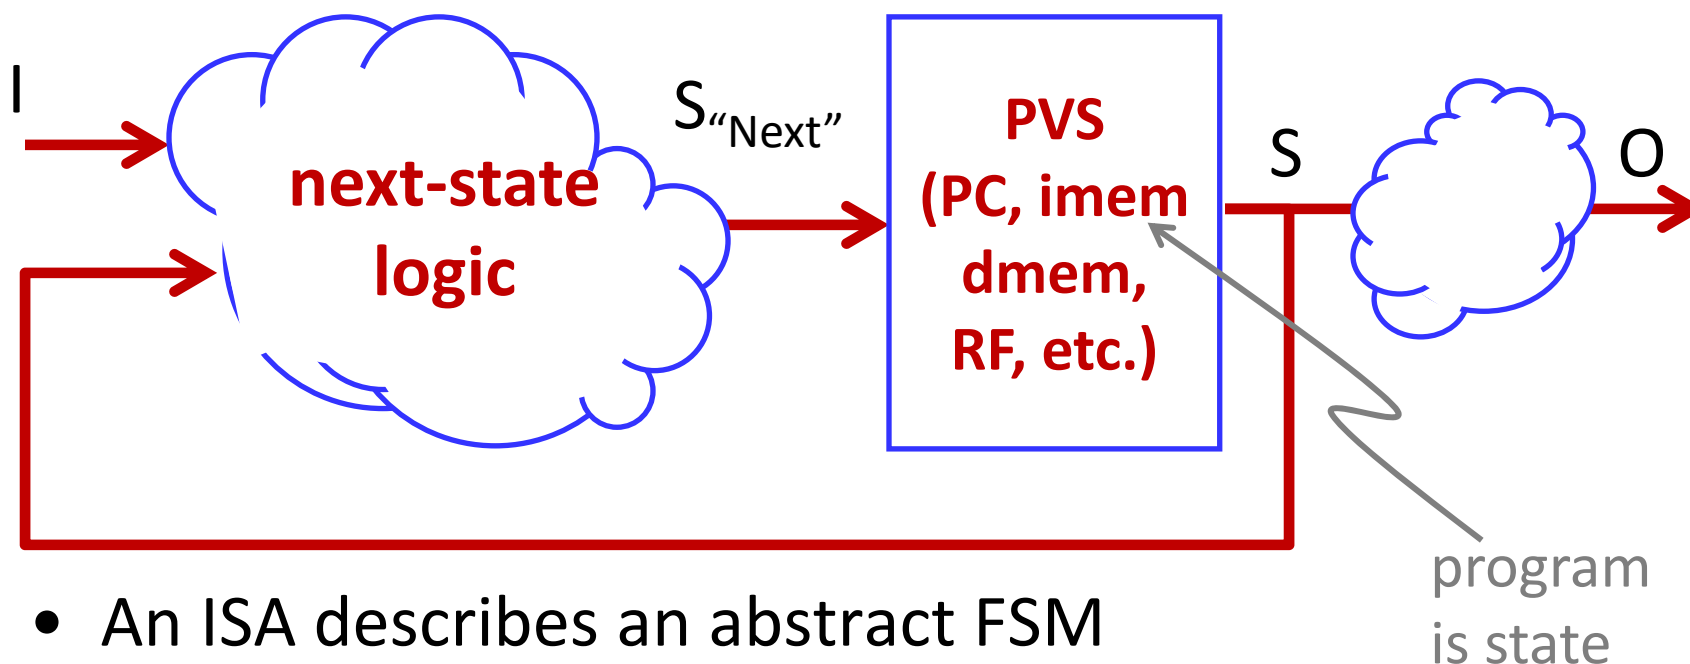
18-447 Lecture 3: Single-Cycle Microarchitecture

James C. Hoe
Department of ECE
Carnegie Mellon University

Housekeeping

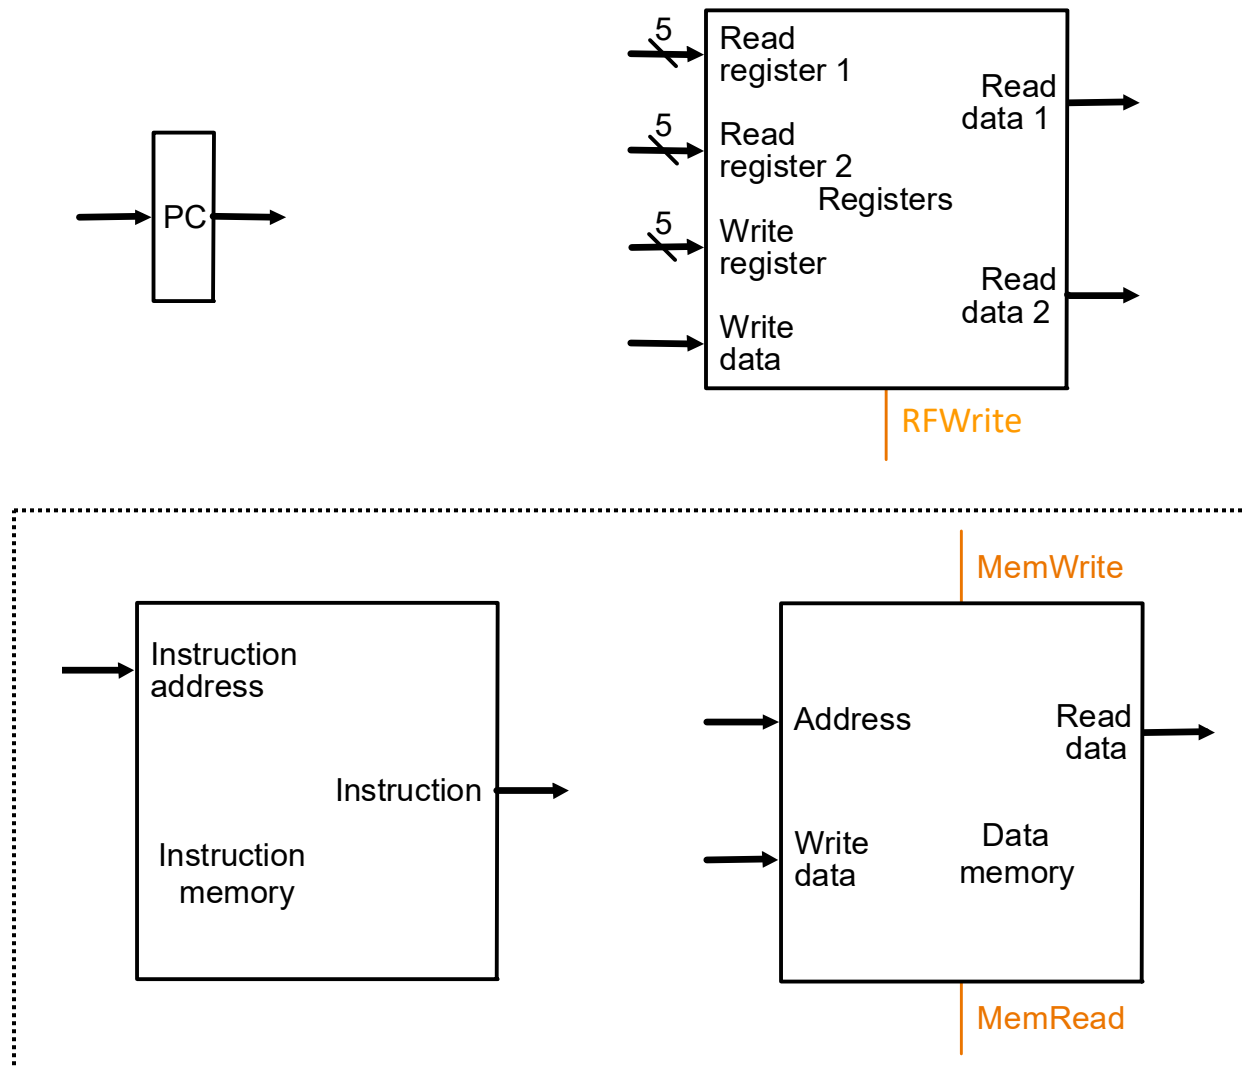
- Your goal today
 - first try at implementing the RV32I ISA
- Notices
 - Student survey on Canvas, **due Wednesday**
 - Handout #4: HW1, **due noon 2/22**
 - Lab 1, Part A, **due week of 2/19**
 - Lab 1, Part B, **due week of 2/26**
- Readings
 - P&H Ch 4.1~4.4
 - rest of P&H Ch 2 for next time

Instruction Processing FSM



- An ISA describes an abstract FSM
 - state = program visible state
 - state transition = instruction execution
- Nice ISAs have atomic instruction semantics
 - one state transition per instruction in abstract FSM
- The implementation FSM can look wildly different

Program Visible State (aka Architectural State)



**Based on original figure from [P&H CO&D, COPYRIGHT 2004 Elsevier. ALL RIGHTS RESERVED.]

“Magic” Memory and Register File

- Combinational Read
 - output of the read data port is a combinational function of the register file contents and the corresponding read select port
 - Synchronous write
 - the selected register (or memory location) is updated on the posedge clock transition when write enable is asserted
- Cannot affect read output in between clock edges

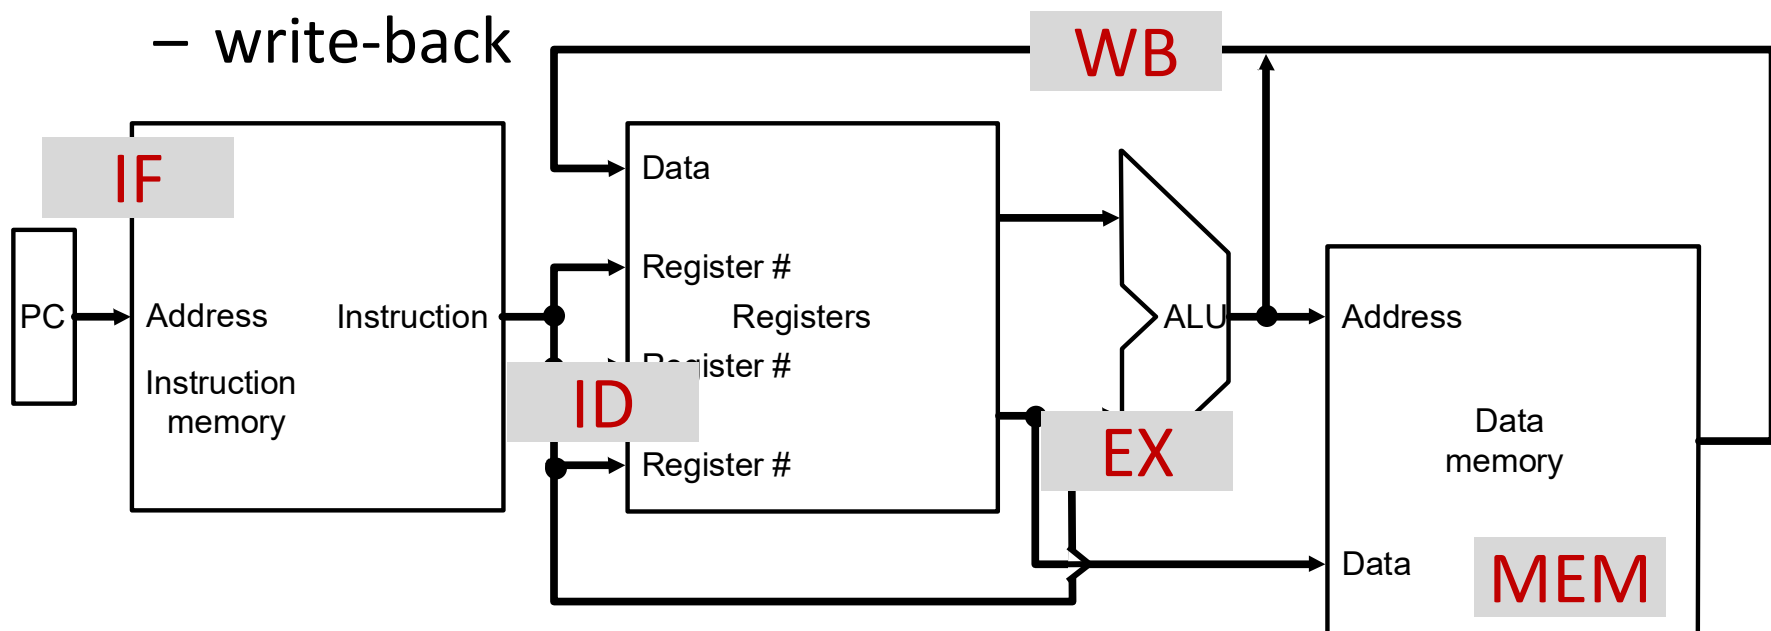
Simplifying Characteristics of “RISC”

- Simple operations
 - 2-input, 1-output arithmetic and logical operations
 - few alternatives for accomplishing the same thing
- Simple data movements
 - ALU ops are register-to-register, never memory
 - “load-store” architecture, 1 addressing mode
- Simple branches
 - limited varieties of branch conditions and targets
 - PC-offset
- Simple instruction encoding
 - all instructions encoded in the same number of bits
 - simple, fixed formats

(RISC=Reduced Instruction Set Computer)

RISC Instruction Processing

- 5 generic steps
 - instruction fetch
 - instruction decode and operand fetch
 - ALU/execute
 - memory access (not required by non-mem instructions)
 - write-back



Single-Cycle Datapath for RV32I ALU Instructions

Register-Register ALU Instructions

- Assembly (e.g., register-register addition)

ADD rd, rs1, rs2

- Machine encoding

0000000	rs2	rs1	000	rd	0110011
7-bit	5-bit	5-bit	3-bit	5-bit	7-bit

- Semantics

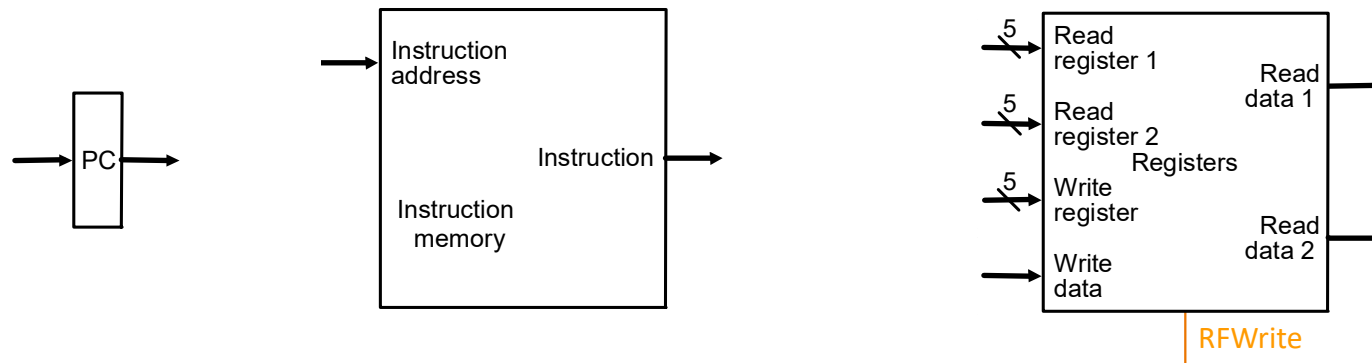
- $\text{GPR}[\text{rd}] \leftarrow \text{GPR}[\text{rs1}] + \text{GPR}[\text{rs2}]$
- $\text{PC} \leftarrow \text{PC} + 4$

- Exceptions: none (ignore carry and overflow)

- Variations

- Arithmetic: {ADD, SUB}
- Compare: {signed, unsigned} x {Set if Less Than}
- Logical: {AND, OR, XOR}
- Shift: {Left, Right-Logical, Right-Arithmetic}

ADD rd rs1 rs2



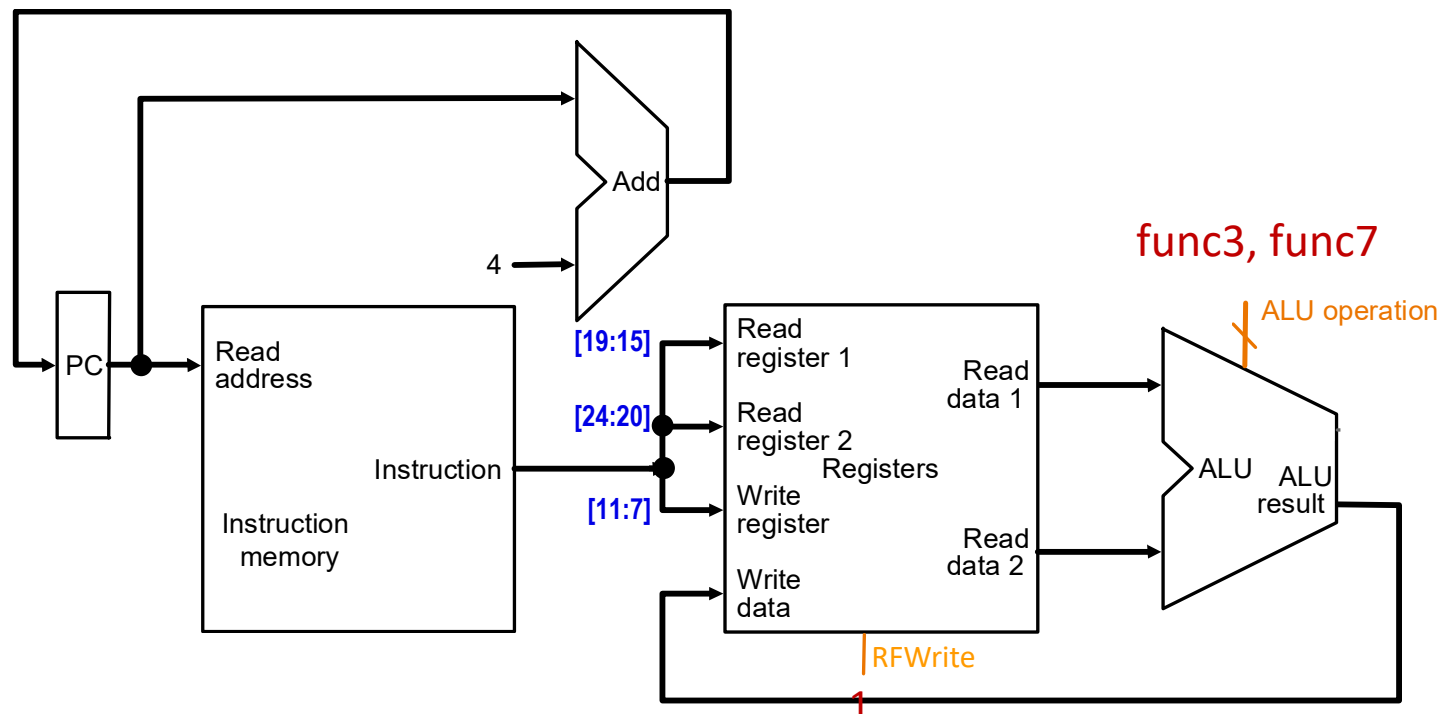
0000000	rs2	rs1	000	rd	0110011
7-bit	5-bit	5-bit	3-bit	5-bit	7-bit

if MEM[PC] == ADD rd rs1 rs2
 $GPR[rd] \leftarrow GPR[rs1] + GPR[rs2]$
 $PC \leftarrow PC + 4$



Combinational
state update logic

R-Type ALU Datapath



**Based on original figure from [P&H CO&D, COPYRIGHT 2004 Elsevier. ALL RIGHTS RESERVED.]

Reg-Immediate ALU Instructions

- Assembly (e.g., reg-immediate additions)

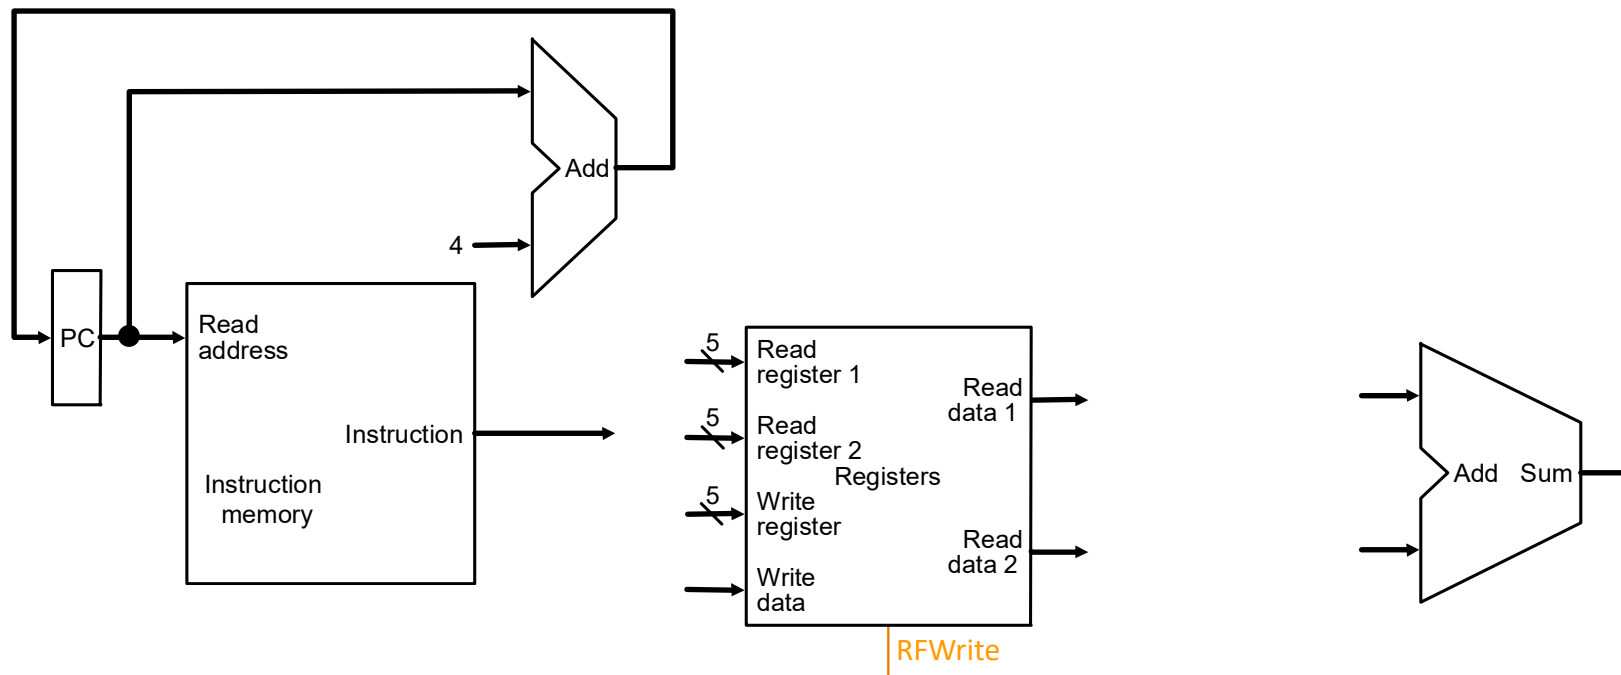
ADDI rd, rs1, imm₁₂

- Machine encoding

imm[11:0]	rs1	000	rd	0010011
12-bit	5-bit	3-bit	5-bit	7-bit

- Semantics
 - $\text{GPR}[\text{rd}] \leftarrow \text{GPR}[\text{rs1}] + \text{sign-extend}(\text{imm})$
 - $\text{PC} \leftarrow \text{PC} + 4$
- Exceptions: none (ignore carry and overflow)
- Variations
 - Arithmetic: {ADDI, ~~SUBI~~}
 - Compare: {signed, unsigned} x {Set if Less Than Imm}
 - Logical: {ANDI, ORI, XORI}
 - **Shifts by unsigned imm[4:0]: {SLLI, SRLI, SRAI}

ADDI rd rs1 immediate₁₂




imm[11:0]	rs1	000	rd	0010011
12-bit	5-bit	3-bit	5-bit	7-bit

if MEM[PC] == ADDI rd rs1 immediate
 GPR[rd] \leftarrow GPR[rs1] + sign-extend (immediate)
 PC \leftarrow PC + 4

IF	ID	EX	MEM	WB
----	----	----	-----	----

Combinational
state update logic

Datapath for R and I-type ALU Inst's



**Based on original figure from [P&H CO&D, COPYRIGHT 2004 Elsevier. ALL RIGHTS RESERVED.]

Single-Cycle Datapath for Data Movement Instructions


Load Instructions

- Assembly (e.g., load 4-byte word)

LW **rd**, **offset₁₂**(**base**)

rs1

- Machine encoding



- Semantics

- $\text{byte_address}_{32} = \text{sign-extend}(\text{offset}_{12}) + \text{GPR}[\text{base}]$
- $\text{GPR}[\text{rd}] \leftarrow \text{MEM}_{32}[\text{byte_address}]$
- $\text{PC} \leftarrow \text{PC} + 4$

- Exceptions: none for now

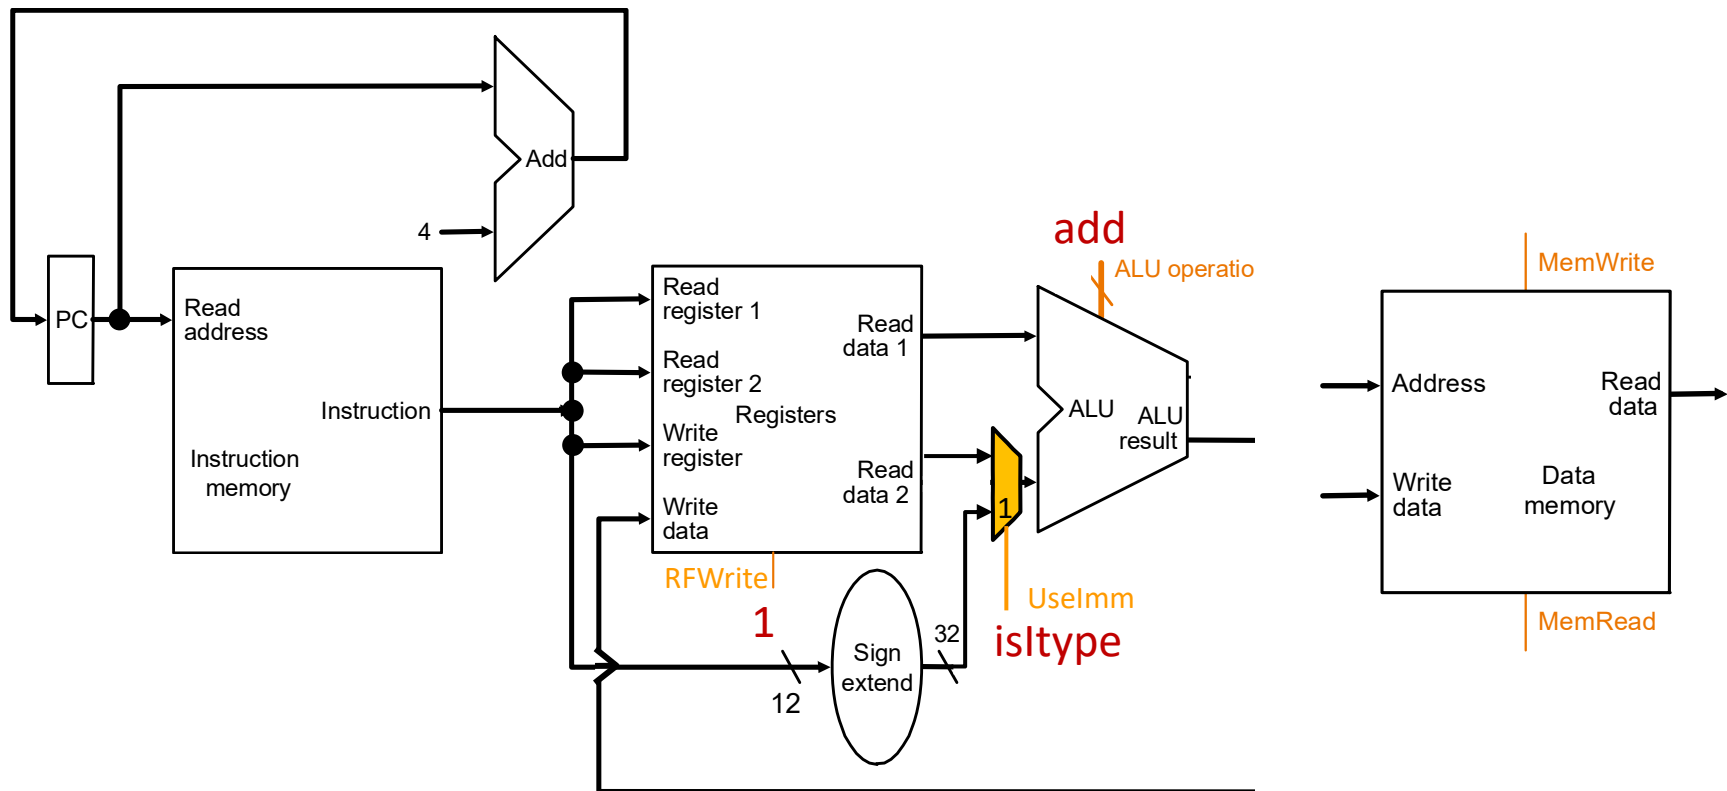
- Variations: LW, LH, LHU, LB, LBU

e.g., LB :: $\text{GPR}[\text{rd}] \leftarrow \text{sign-extend}(\text{MEM}_8[\text{byte_address}])$

LBU :: $\text{GPR}[\text{rd}] \leftarrow \text{zero-extend}(\text{MEM}_8[\text{byte_address}])$

Note: RV32I memory is byte-addressable, little-endian

LW Datapath



```

if MEM[PC]==LW rd offset12(base)
    EA = sign-extend(offset) + GPR[base]
    GPR[rd] ← MEM[ EA ]
    PC ← PC + 4

```



Combinational state update logic

Store Instructions

- Assembly (e.g., store 4-byte word)

SW $rs2, offset_{12}(base)$

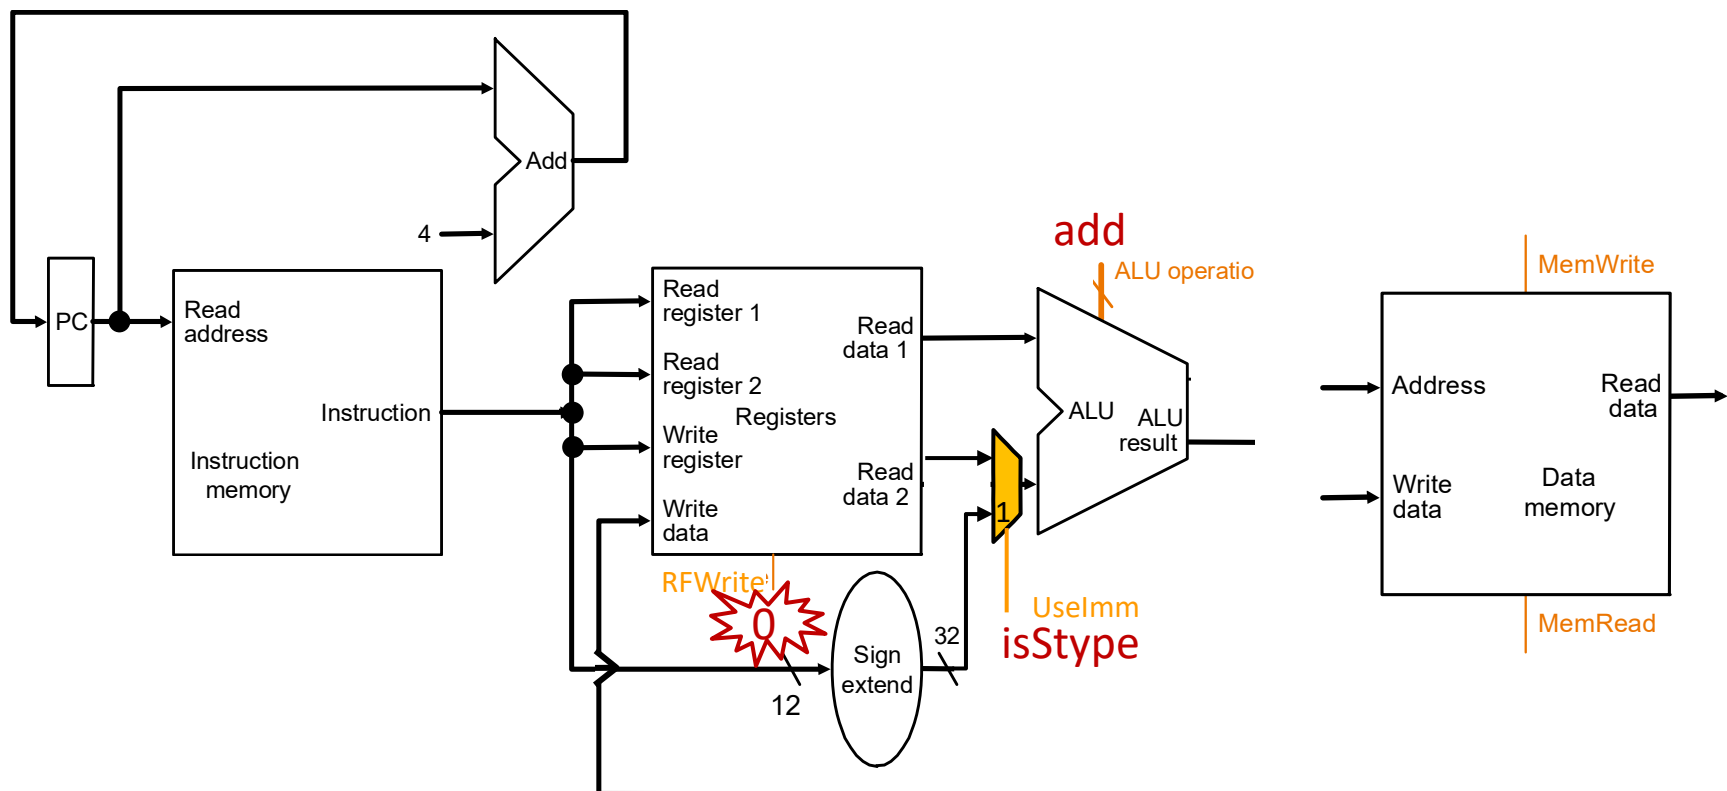
- Machine encoding

offset[11:5]	rs2	base	010	ofst[4:0]	0100011
7-bit	5-bit	5-bit	3-bit	5-bit	7-bit

- Semantics
 - $byte_address_{32} = \text{sign-extend}(offset_{12}) + GPR[base]$
 - $MEM_{32}[byte_address] \leftarrow GPR[rs2]$
 - $PC \leftarrow PC + 4$
- Exceptions: none for now
- Variations: SW, SH, SB

e.g., SB:: $MEM_8[byte_address] \leftarrow (GPR[rs2])[7:0]$

SW Datapath



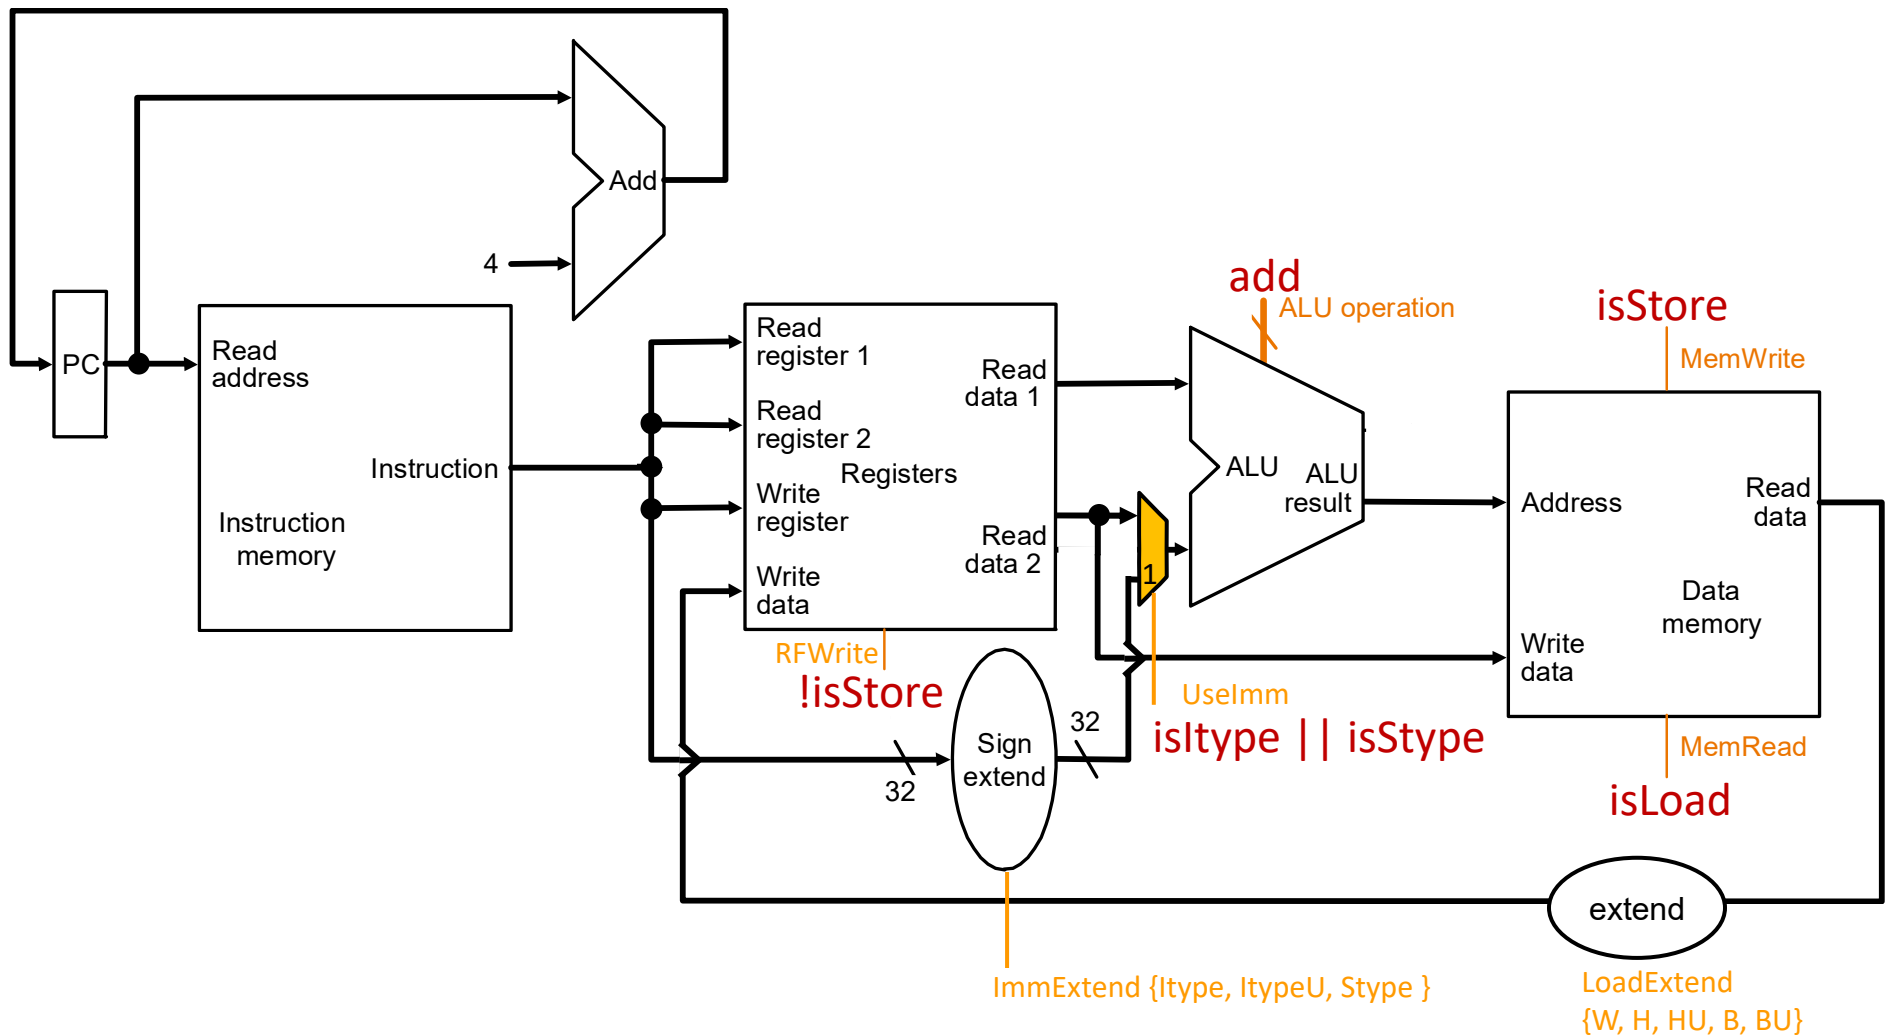
```

if MEM[PC]==SW rs2 offset12(base)
  EA = sign-extend(offset) + GPR[base]
  MEM[ EA ] ← GPR[rs2]
  PC ← PC + 4
  
```

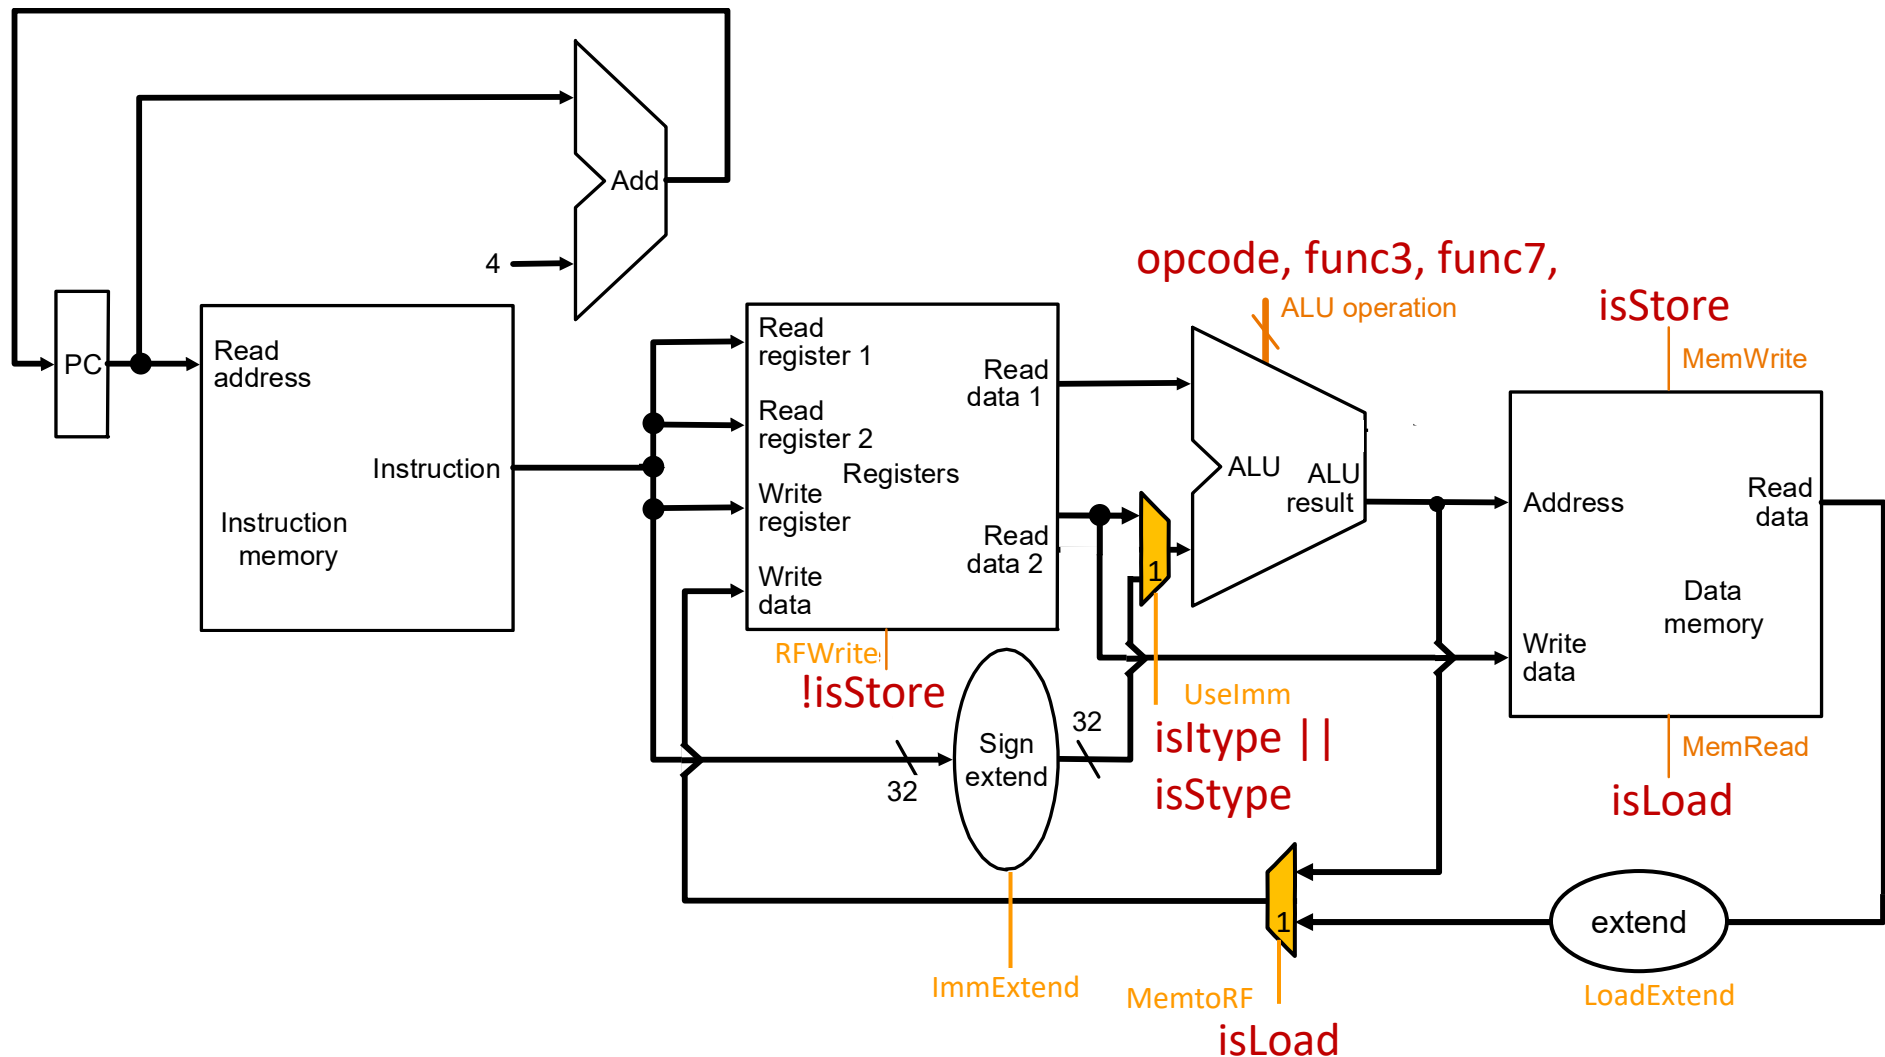
IF	ID	EX	MEM	WB
----	----	----	-----	----

Combinational
state update logic

Load-Store Datapath



Datapath for Non-Control Flow Inst's



Single-Cycle Datapath for Control Flow Instructions


Jump and Link Instruction

- Assembly

JAL *rd* *imm*₂₁

Note: implicit *imm*[0]=0

- Machine encoding



- Semantics

– target = PC + sign-extend(*imm*₂₁)

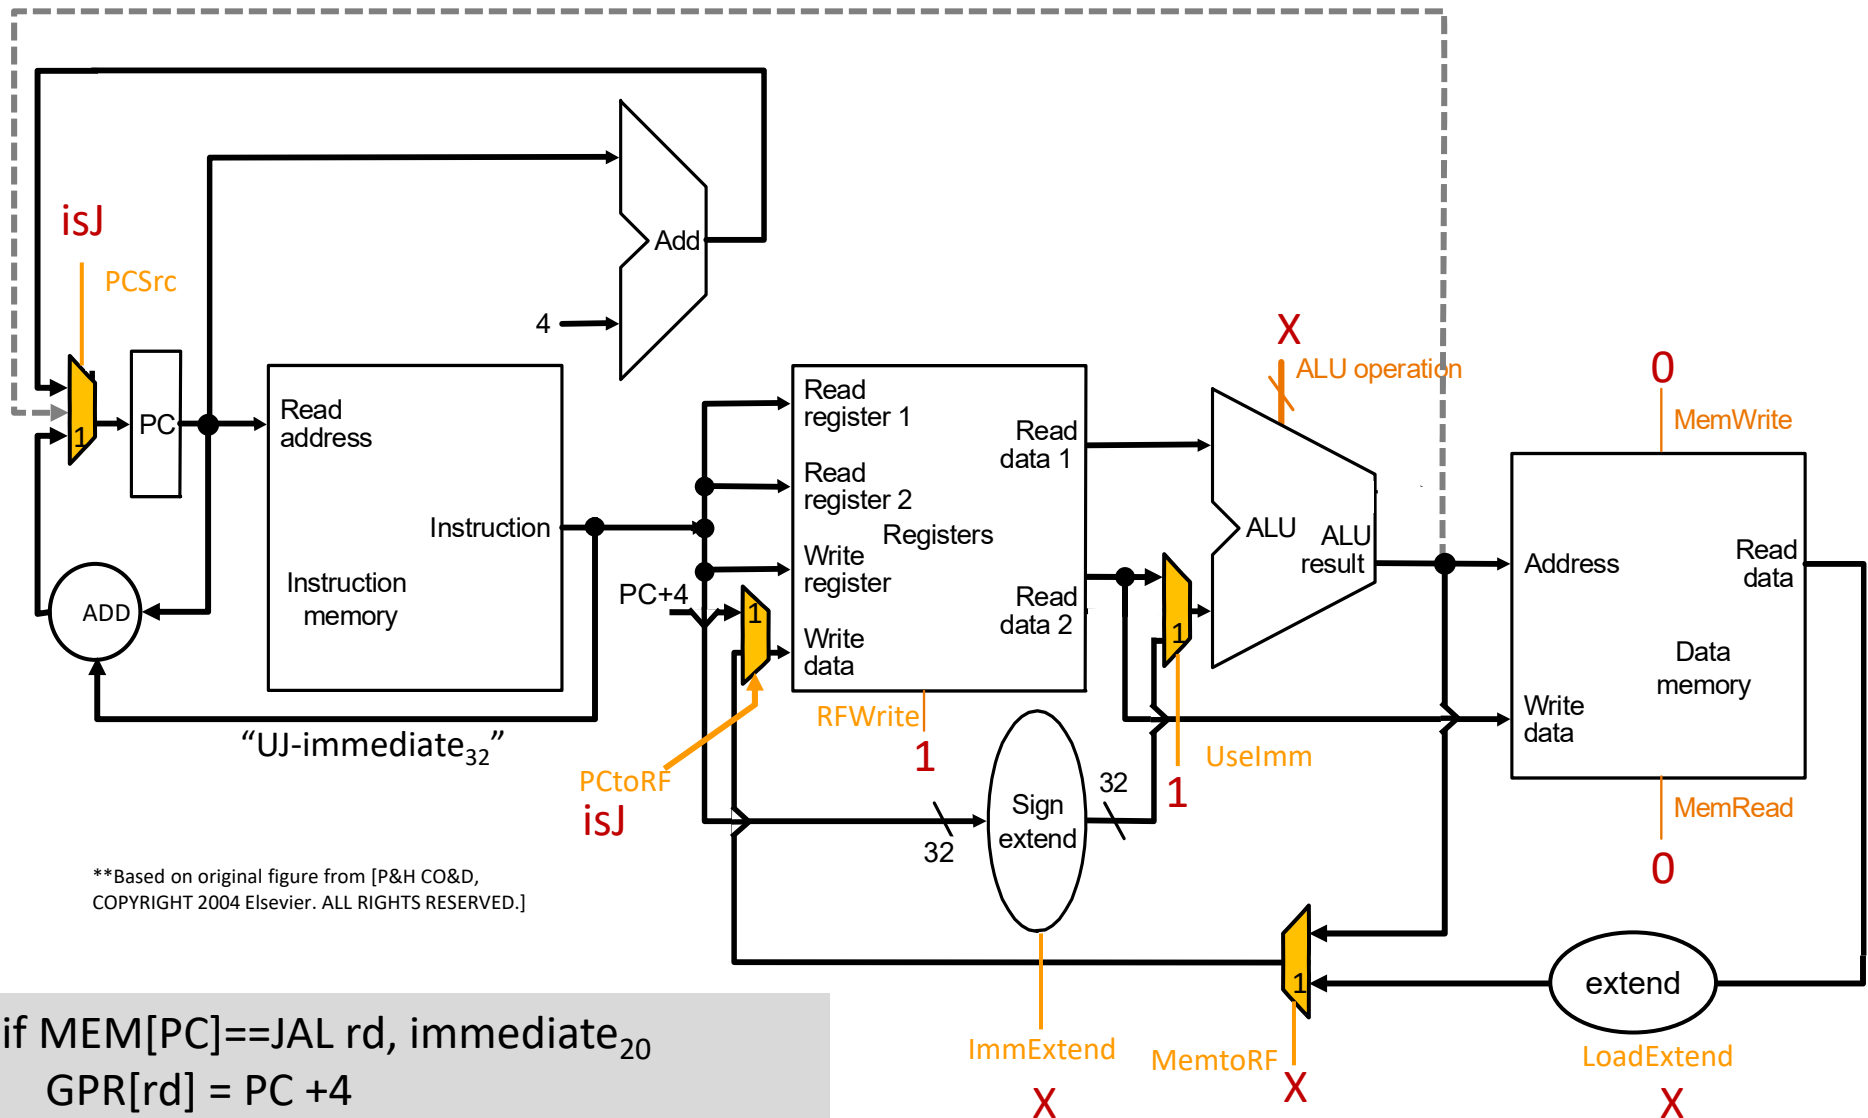
– GPR[*rd*] ← PC + 4

– PC ← target

How far can you jump?

- Exceptions: misaligned target (4-byte)

Unconditional Jump Datapath



**Based on original figure from [P&H CO&D,
COPYRIGHT 2004 Elsevier. ALL RIGHTS RESERVED.]

What about JALR?


(Conditional) Branch Instructions

- Assembly (e.g., branch if equal)

BEQ *rs1*, *rs2*, *imm*₁₃

Note: implicit *imm*[0]=0

- Machine encoding



- Semantics

– target = PC + sign-extend(*imm*₁₃)

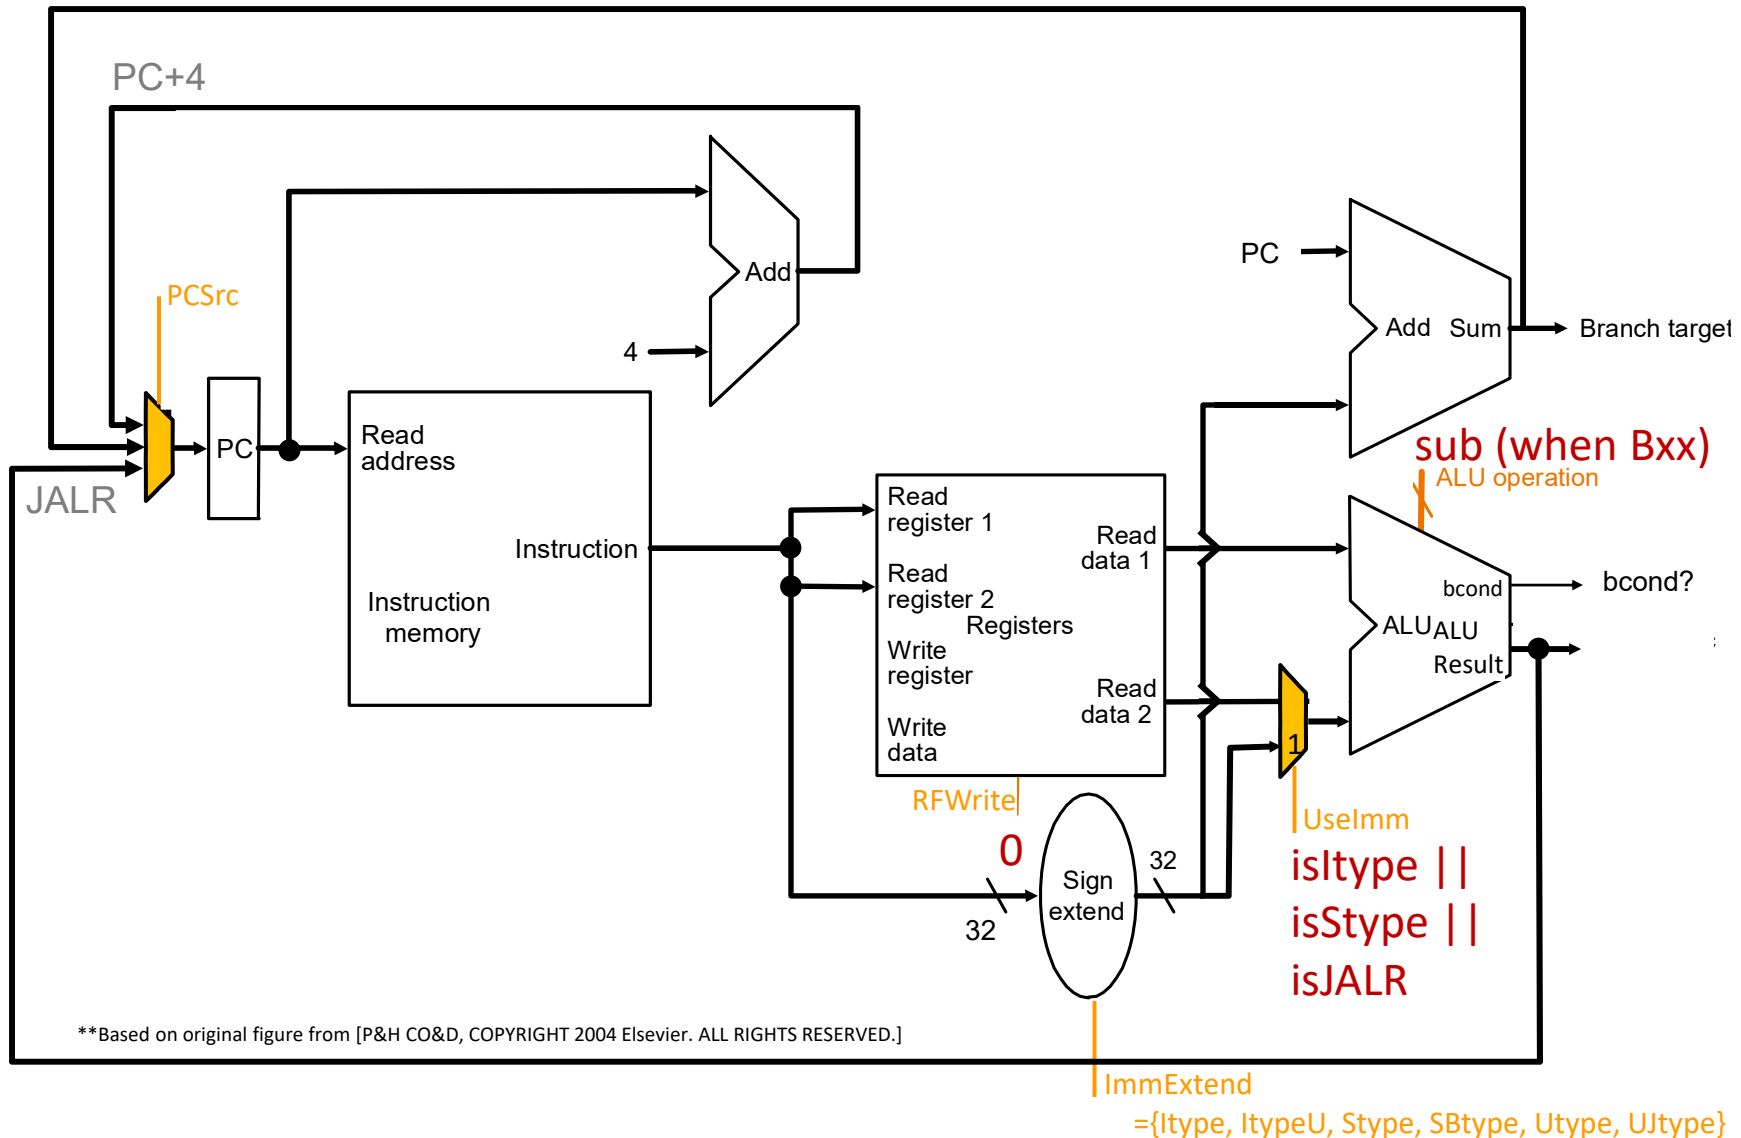
– if GPR[*rs1*]==GPR[*rs2*] then PC ← target
 else PC ← PC + 4

How far can you jump?

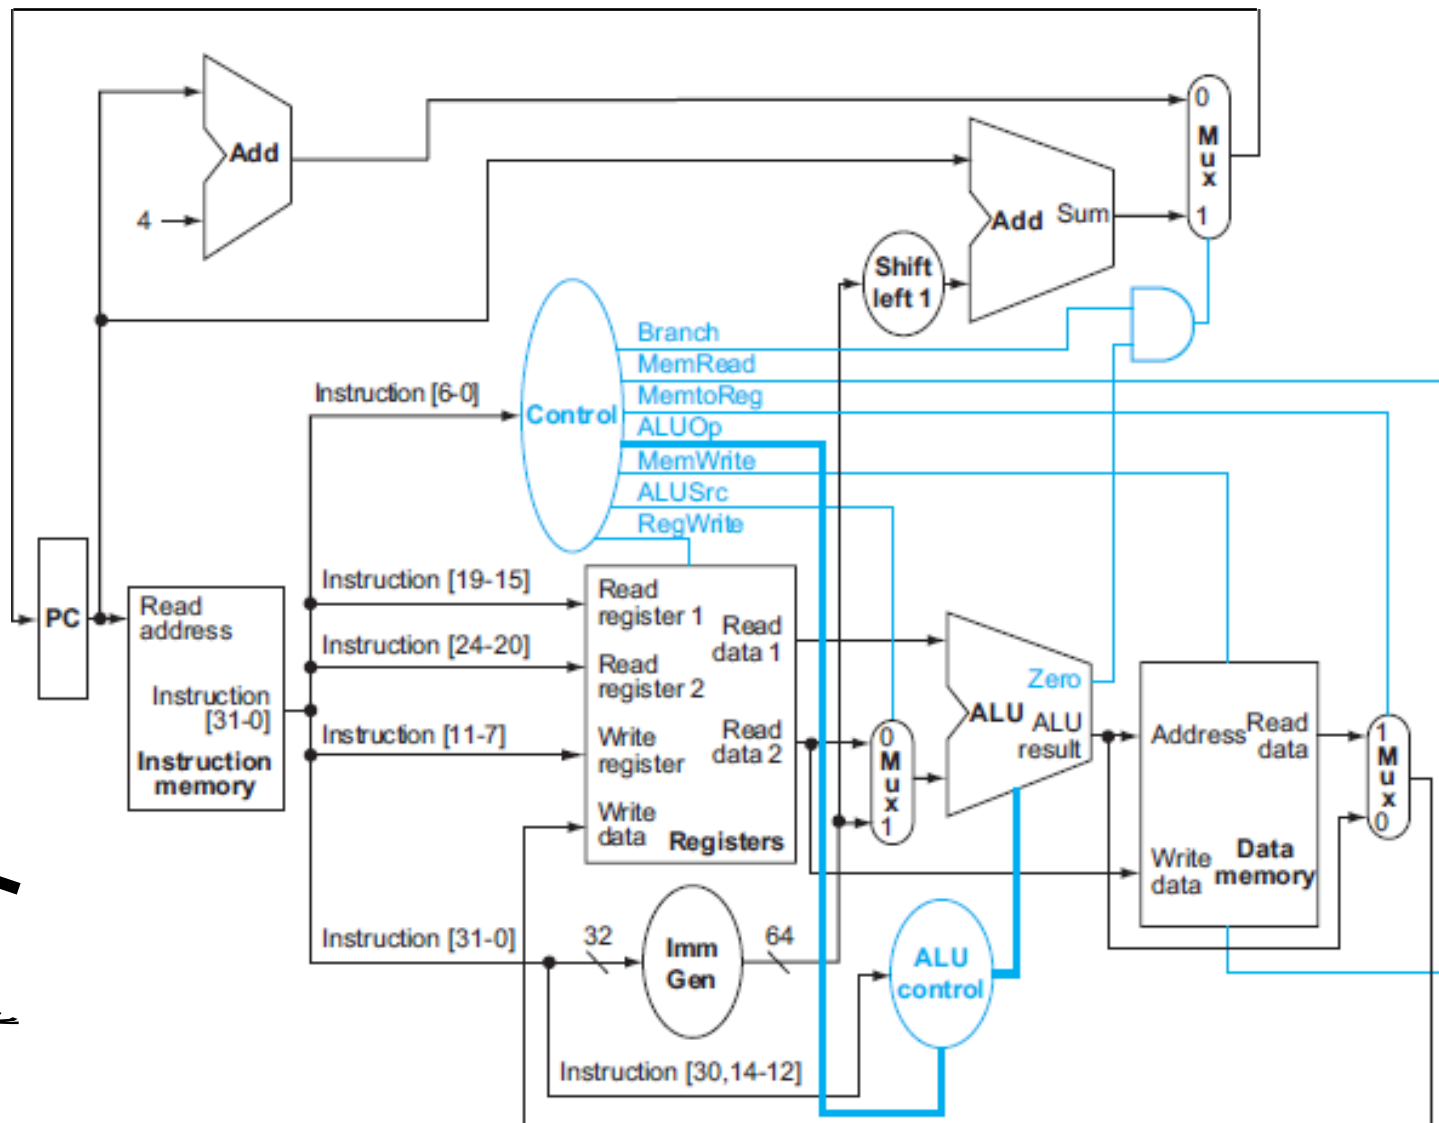
- Exceptions: misaligned target (4-byte) if taken
- Variations
 - BEQ, BNE, BLT, BGE, BLTU, BGEU

Conditional Branch Datapath

JAL and taken Branch

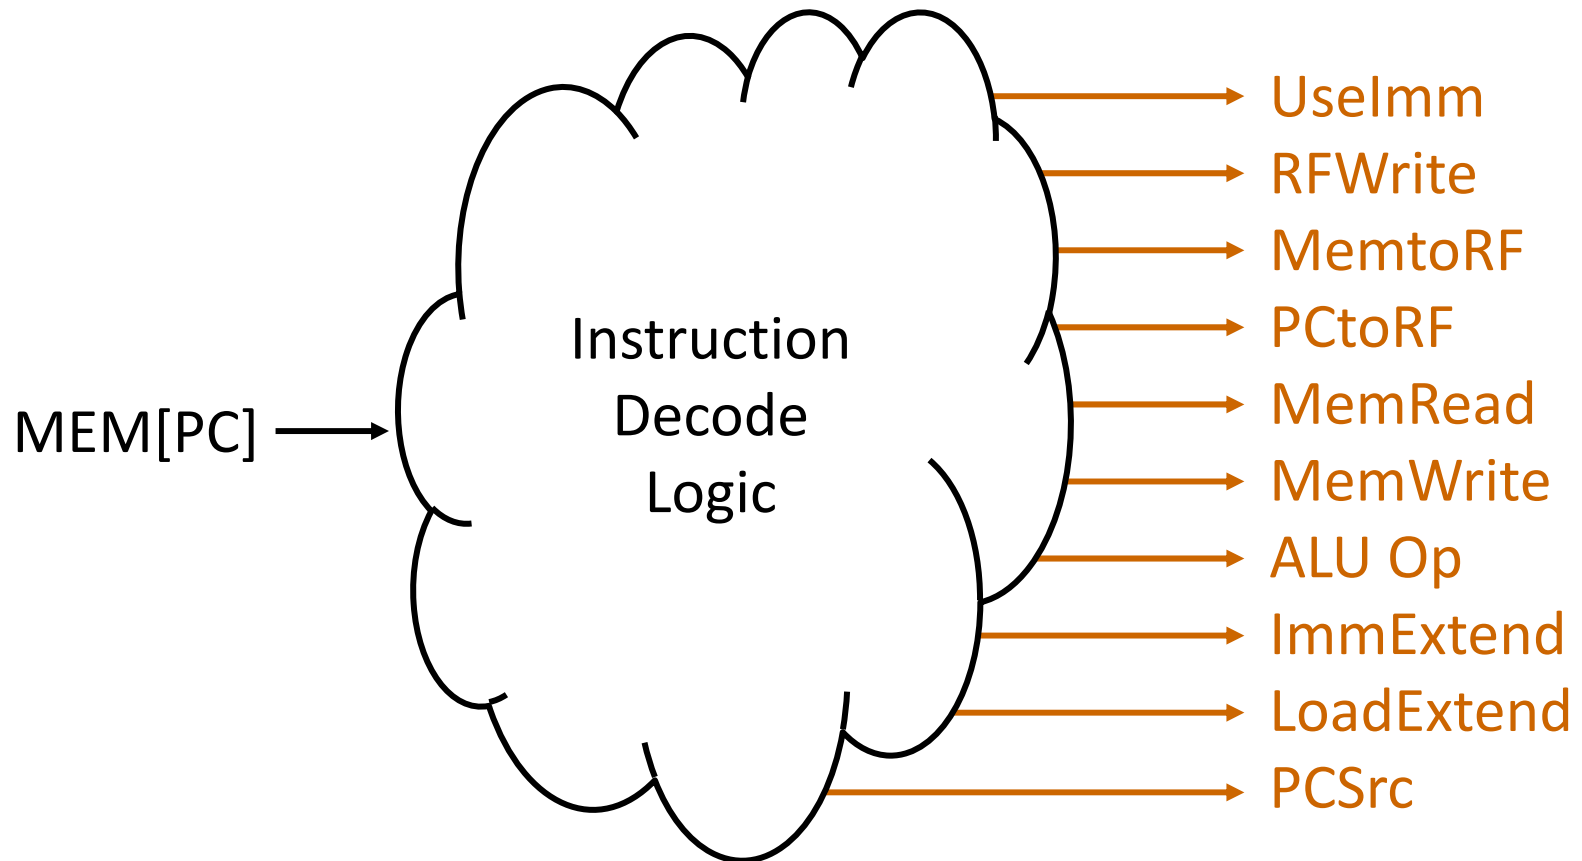


Adding Control to Datapath

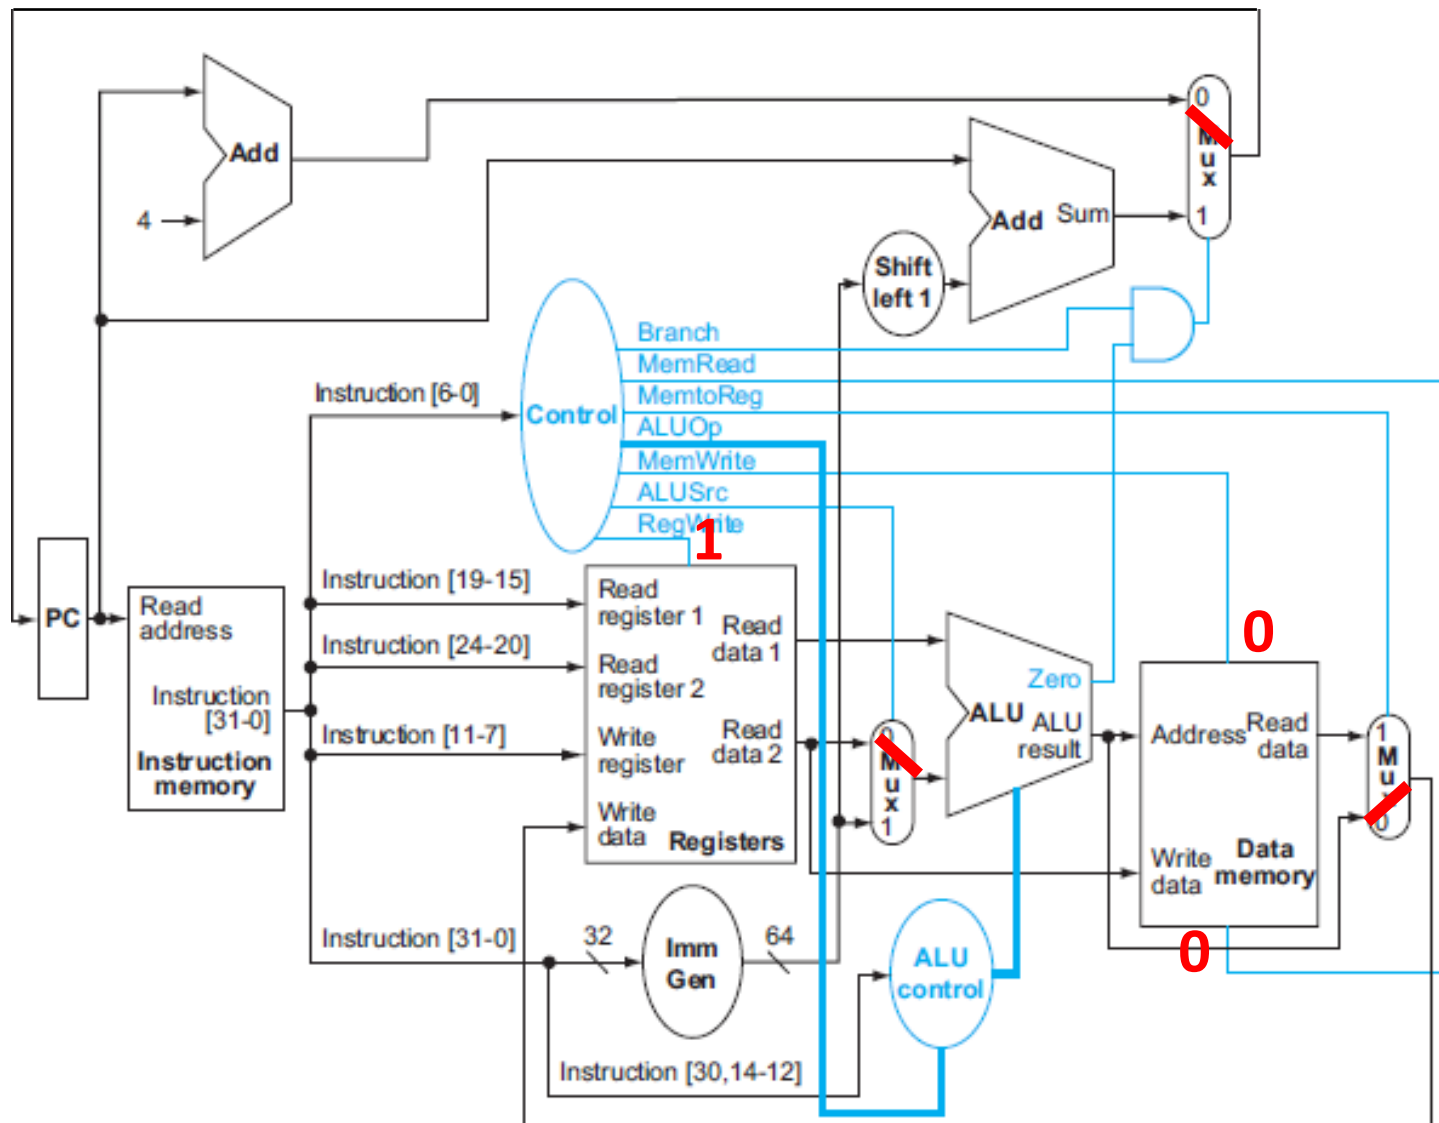


[Figure 4.17 from book, Copyright © 2018 Elsevier Inc. All rights reserved.]


Datapath Control Generation



R-Type ALU Worksheet:




I-Type ALU Worksheet:




[Figure 4.17 from book, Copyright © 2018 Elsevier Inc. All rights reserved.]

LW Worksheet:




[Figure 4.17 from book, Copyright © 2018 Elsevier Inc. All rights reserved.]

SW Worksheet:



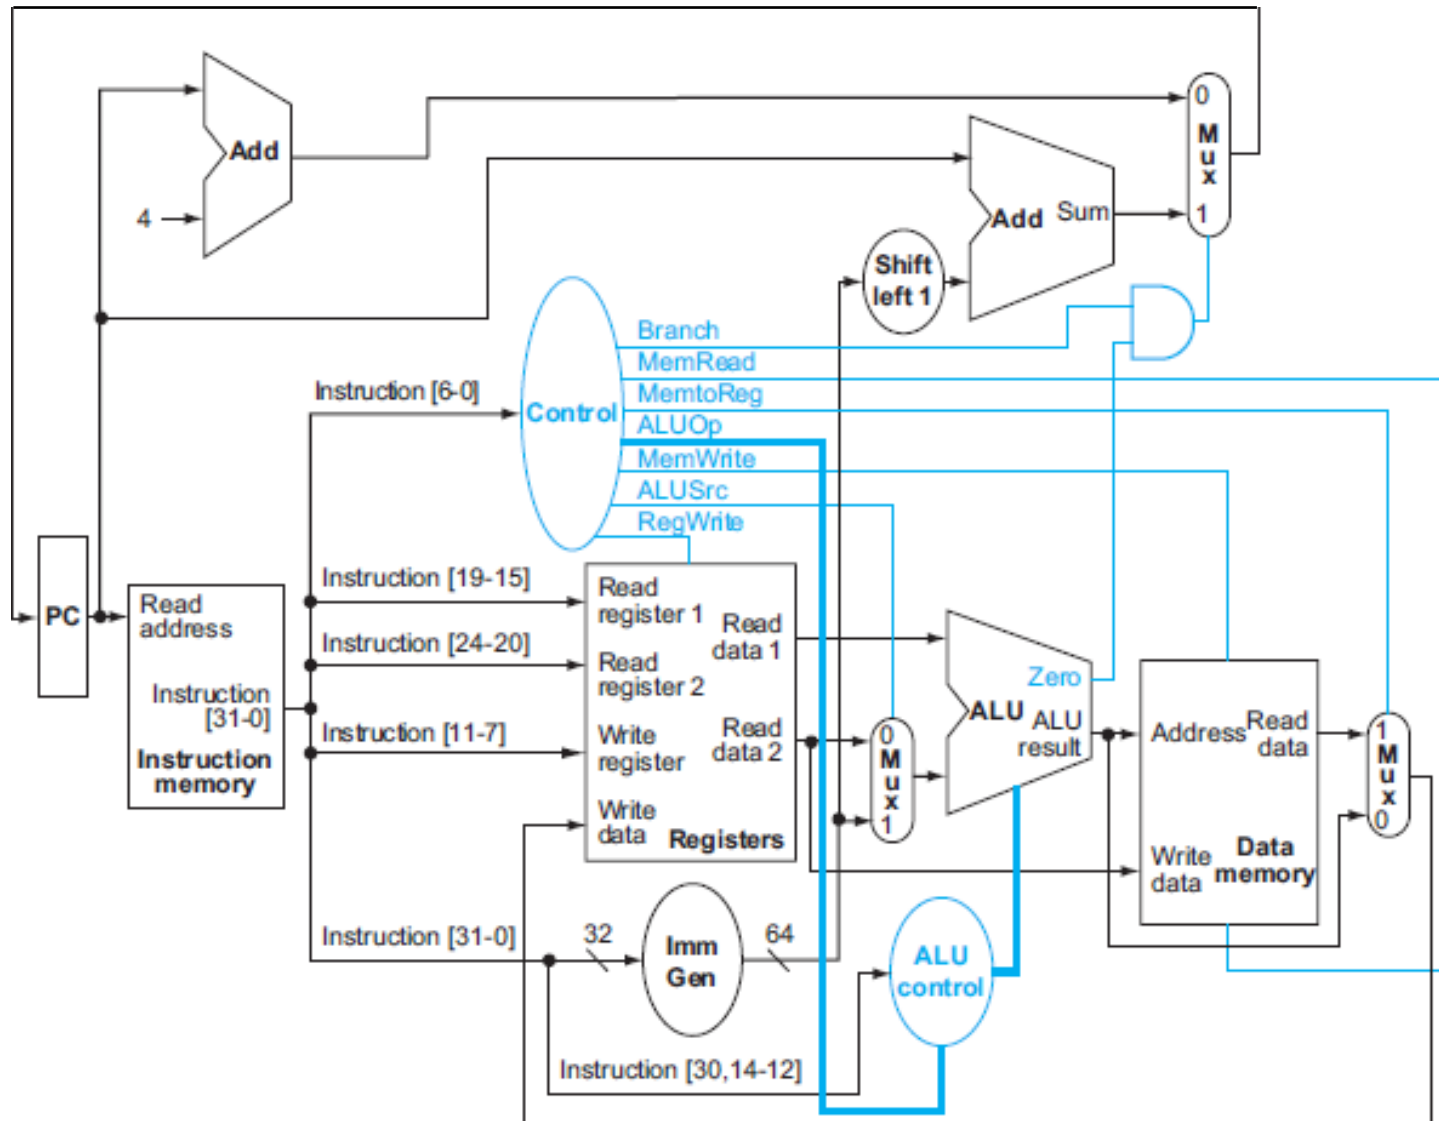
[Figure 4.17 from book, Copyright © 2018 Elsevier Inc. All rights reserved.]

Branch Taken Worksheet:



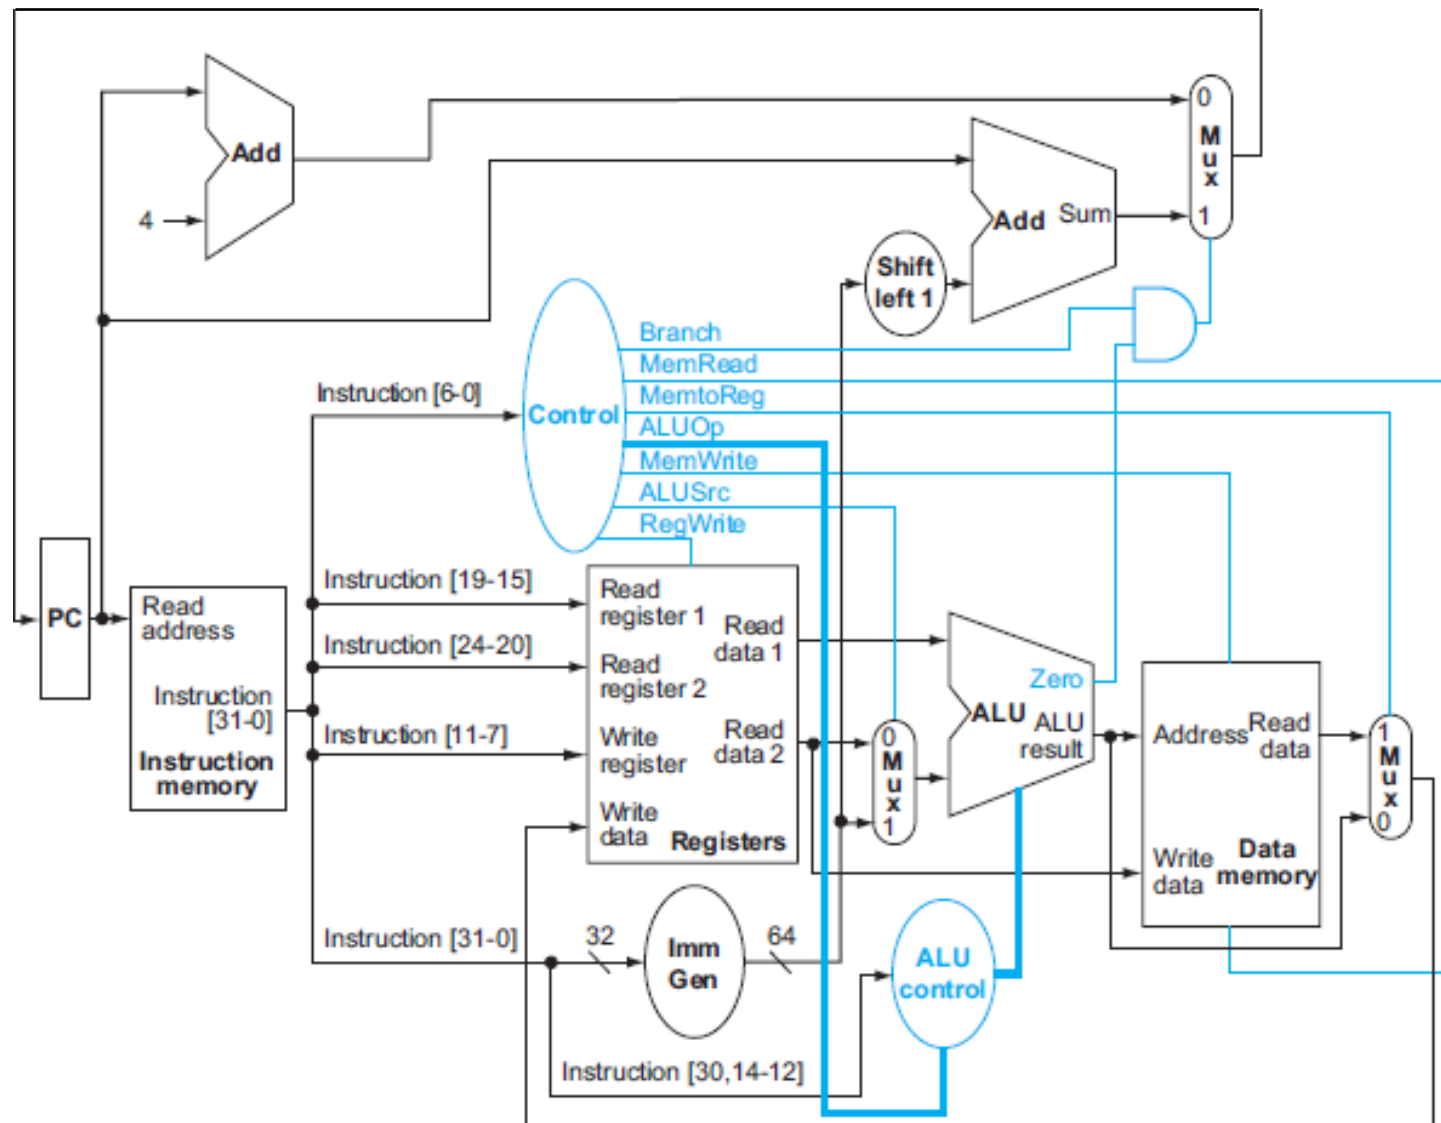
[Figure 4.17 from book, Copyright © 2018 Elsevier Inc. All rights reserved.]

Branch Not-Taken Worksheet:



[Figure 4.17 from book, Copyright © 2018 Elsevier Inc. All rights reserved.]

Jump (and Link?) ALU Worksheet:



[Figure 4.17 from book, Copyright © 2018 Elsevier Inc. All rights reserved.]

Single-Bit Control Signals

	When De-asserted	When asserted	Equation
UseImm	2 nd ALU input from 2 nd GPR read port	2 nd ALU input from immediate	$(\text{opcode} \neq \text{IsRtype}) \ \&\& \ (\text{opcode} \neq \text{isBtype})$
RFWrite	GPR write disabled	GPR write enabled	$(\text{opcode} \neq \text{SW}) \ \&\& \ (\text{opcode} \neq \text{Bxx})$
MemtoRF	Steer ALU result to GPR write port	steer memory load to GPR write port	$\text{opcode} == \text{LW/H/B}$
PCtoRF	Steer above result to GPR write port	Steer PC+4 to GPR write port	$(\text{opcode} == \text{JAL}) \ \parallel \ (\text{opcode} == \text{JALR})$
MemRead	Memory read disabled	Memory read port return load value	$\text{opcode} == \text{LW/H/B}$
MemWrite	Memory write disabled	Memory write enabled	$\text{opcode} == \text{SW/H/B}$

Multi-Bit Control Signals

	Options	Equation
ALU Op	<ul style="list-style-type: none"> ADD, SUB, AND, OR, XOR, NOR, LT, and Shift bcond: EQ, NE, GE, LT 	case opcode RTypeALU: according to funct3, funct7[5] ITypeALU : according to funct3 only (except shift) LW/SW/JALR : ADD Bxx : SUB and select bcond function — : ??
ImmExtend	Itype, ItypeU, Stype, SBtype, Utype, UJtype	<ul style="list-style-type: none"> select based on instruction format type (may want to have separate extension units for primary ALU and PC-offset adder)
PCSrc	PC+4, PCadder, ALU	case opcode JAL : PC + immediate JALR : GPR + immediate Bxx : taken?(PC + immediate):(PC + 4) — : PC+4
LoadExtend	W,H,HU,B,BU	case func3

Architecture vs Microarchitecture

Architecture

- Architectural Level



a clock has a hour hand and a minute hand,



a computer does?????....

You can read a clock without knowing how it works

conceptual

Microarchitecture

- Microarchitecture Level



a particular clockwork has a certain set of gears arranged in a certain configuration



a particular computer design has a certain datapath and a certain control logic

- Realization Level



machined alloy gears vs stamped sheet metal



CMOS vs ECL vs vacuum tubes

physical